

MDL EndSem War Room

Tactical Exam Preparation Guide

IIIT Hyderabad — Spring 2026

Built from PYQs (2022–2025), Lectures, and Comprehensive Guide

For Harith Y.

May 2, 2026

Contents

How to Use This Guide	3
1 Master Formula Sheet	4
2 Logic and Propositional Reasoning	8
2.1 Quick Reference	8
2.2 Recipes	8
2.3 Worked PYQs — Logic	9
3 Probability, Bayes, and Naive Bayes	11
3.1 Recipes	11
3.2 Worked PYQs — Probability	11
4 Bias-Variance Tradeoff and Regularization	14
4.1 Conceptual Foundation	14
4.2 Recipes	14
4.3 Worked PYQs — Bias-Variance	15
5 Decision Trees: ID3, CART, and Pruning	17
5.1 When to Use Which Algorithm	17
5.2 Recipes	17
5.3 Worked PYQs — Decision Trees	18
6 KNN, Clustering (K-Means), and Hierarchical Clustering	20
6.1 KNN Recipe	20
6.2 K-Means Recipe	21
6.3 Hierarchical Clustering — Linkage Definitions	21
7 Decision Theory: Utility, MEU, and POMDPs	23
7.1 Recipes	23
7.2 Worked PYQs — Decision Theory	24
8 Markov Decision Processes (MDPs)	26
8.1 Formalism Refresher	26
8.2 Finite-Horizon Recipe (most-asked variant)	26
8.3 Worked PYQs — MDPs	27
9 Multi-Armed Bandits	31
9.1 Recipes	31
9.2 Worked PYQ — Bandits	31

10 State-Space Search	33
10.1 Algorithm Crib	33
10.2 Recipes	33
10.3 Worked PYQs — Search	34
11 Association Rules and Apriori	37
11.1 Recipes	37
12 Cross-Validation, Data Leakage, and Pipelines	38
12.1 Quick Reference	38
13 Solved 2025 Endsem (Tactical Walkthrough)	39
Q1: Logic	39
Q2: MDP / Q-learning	39
Q3: Probability + KNN	40
Q4: Decision Tree	41
Q5: Simulated Annealing	42
Q6: Bias-Variance Regularization	42
Q7: Bias-Variance Computation	42
Q8: Sudoku Hill Climbing	43
Q9: State-Space Search	43
14 Solved 2025 Endsem — Extended (Q2, Q7c, Q8c)	45
14.1 Q2 [Multi-part, ~15 marks] — Traffic Route Selection	45
14.2 Q7 (c) [2+2 pts] — Off-by-one MSE bug	47
14.3 Q8 (c) [2 pts] — Genetic algorithm fitness function	47
15 Solved 2024 Endsem — Full Walkthrough	49
16 Solved 2023 Endsems (Both Versions) — Full Walkthrough	53
17 Solved 2022 Endsem — Full Walkthrough	60
18 Cross-Year Pattern Map: What Recurs	64
19 Last-Page Cheat Sheet	65

How to Use This Guide

Why this guide exists. You have already produced the *96-page Comprehensive Guide* (covering theory thoroughly). This *War Room* guide is different — it is built around the question patterns that actually appear in MDL endsems (2022–2025), with a heavy emphasis on **step-by-step solutions**, **pattern recognition**, and **traps to avoid**.

The honest diagnosis. From the prior exam scripts, the recurring failure modes are:

1. *Knowing the formula but not when to apply which version.* (e.g. $\text{Bias} = \mathbb{E}[\hat{f}] - f$ vs. MSE-decomposition vs. empirical bias on a fold.)
2. *Stopping at half-solutions.* Examiners look for: formula $\rightarrow$ substitution $\rightarrow$ arithmetic $\rightarrow$ interpretation. Skipping any step loses 50%.
3. *Confusing similar-looking problems.* Single-link vs. complete-link, KNN vs. K-Means, ID3 vs. CART, value iteration vs. policy iteration.
4. *Missing edge cases on MDP/Search.* Slipping probabilities, terminal states, $\gamma = 1$ traps.

Strategy for the next 48 hours.

- **Day 1:** Chapters 1–7 of this guide (formulas, recipes, traps). Solve 2025 paper from Chapter 14 *closed-book*.
- **Day 2:** Solve 2024 + 2023 papers (Chapter 14). Read every *Pitfall* box. Last 2 hours: glance at Chapter 15 “one-page cheat sheet”.

Box legend:

- Recipe — step-by-step procedure for a question type.
- Exam Trick — a high-leverage shortcut or insight.
- Worked PYQ — a full-marks solution to a real past question.
- Pitfall — a specific mistake students made on past exams.
- How to Think — the mental model that makes the topic “click”.
- Formula — the canonical formula(s).

Chapter 1

Master Formula Sheet

Bias-Variance-MSE Decomposition

For a single test point x , with $y = f(x) + \varepsilon$, $\mathbb{E}[\varepsilon] = 0$, $\text{Var}(\varepsilon) = \sigma^2$:

$$\mathbb{E}[(y - \hat{f}(x))^2] = \underbrace{(\mathbb{E}[\hat{f}(x)] - f(x))^2}_{\text{Bias}^2} + \underbrace{\mathbb{E}[(\hat{f}(x) - \mathbb{E}[\hat{f}(x)])^2]}_{\text{Variance}} + \underbrace{\sigma^2}_{\text{Irreducible}}$$

Bias (single value, not squared): $\text{Bias}(\hat{f}(x)) = \mathbb{E}[\hat{f}(x)] - f(x)$.

Variance: $\text{Var}(\hat{f}(x)) = \mathbb{E}[\hat{f}(x)^2] - (\mathbb{E}[\hat{f}(x)])^2$.

Empirical version (used in PYQs!): average over the N training points: $\text{Bias} = \frac{1}{N} \sum_i (\hat{y}_i - y_i)$, $\text{Var} = \frac{1}{N} \sum_i (\hat{y}_i - \bar{\hat{y}})^2$.

Probability & Bayes

$$P(A | B) = \frac{P(A \cap B)}{P(B)} \quad (\text{conditional})$$

$$P(A | B) = \frac{P(B | A) P(A)}{P(B)} \quad (\text{Bayes})$$

$$P(B) = \sum_a P(B | A = a) P(A = a) \quad (\text{law of total probability})$$

$$P(A, B | C) = P(A | C) P(B | C) \quad (\text{if } A \perp B | C)$$

Classification Metrics — Confusion Matrix

With TP, FP, TN, FN:

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + TN + FN}, \text{Precision} = \frac{TP}{TP + FP}, \text{Recall} = \frac{TP}{TP + FN}$$

$$F_1 = \frac{2 \cdot \text{Prec} \cdot \text{Rec}}{\text{Prec} + \text{Rec}}$$

Disease test trap: “99% accurate” usually means TPR=99% *and* TNR=99% separately. Use Bayes for $P(\text{disease} | +)$.

Information Theory

$$H(S) = - \sum_k p_k \log_2 p_k \quad (0 \log_2 0 := 0)$$

$$\text{IG}(S, A) = H(S) - \sum_v \frac{|S_v|}{|S|} H(S_v) \quad (\text{ID3})$$

$$\text{Gini}(S) = 1 - \sum_k p_k^2 \quad (\text{CART} \text{ — pick attribute with } \textit{minimum} \text{ weighted Gini})$$

$$\text{Gini}(S, A) = \sum_v \frac{|S_v|}{|S|} \text{Gini}(S_v) \quad (\text{weighted Gini after split})$$

Distance Metrics

$$\text{Euclidean: } d(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_i (x_i - y_i)^2} \quad \text{Manhattan: } d = \sum_i |x_i - y_i|$$

$$\text{Minkowski-}p : d = \left(\sum_i |x_i - y_i|^p \right)^{1/p} \quad \text{Cosine: } 1 - \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}$$

Clustering Linkages (Hierarchical)

For clusters C_1, C_2 :

$$\text{Single-link : } \min_{x \in C_1, y \in C_2} d(x, y) \quad (\text{closest pair})$$

$$\text{Complete-link : } \max_{x \in C_1, y \in C_2} d(x, y) \quad (\text{farthest pair})$$

$$\text{Average-link : } \frac{1}{|C_1||C_2|} \sum_{x \in C_1} \sum_{y \in C_2} d(x, y) \quad (\text{mean over all pairs})$$

Single $\Rightarrow$ chaining (long thin clusters); Complete $\Rightarrow$ tight compact clusters.

Regularization

$$J_{\text{ridge}} = \text{MSE} + \lambda \|\mathbf{w}\|_2^2 \quad (L_2) \quad J_{\text{lasso}} = \text{MSE} + \lambda \|\mathbf{w}\|_1 \quad (L_1)$$

$\lambda \uparrow$ shrinks weights $\Rightarrow$ **bias** $\uparrow$, **variance** $\downarrow$. Lasso also zeroes weights (feature selection).

Decision Theory: MEU and Utility

$$a^* = \arg \max_a \mathbb{E}[U(o) \mid a] = \arg \max_a \sum_o P(o \mid a) U(o)$$

Risk attitude (shape of U): concave $\Rightarrow$ *averse*; linear $\Rightarrow$ *neutral*; convex $\Rightarrow$ *seeking*. A piecewise utility ($x/2$ then x then $2x$) has *multiple* risk regimes — read the question carefully.

$$\text{VPI}(E) = \mathbb{E}_E[\max_a \mathbb{E}[U \mid a, E]] - \max_a \mathbb{E}[U \mid a] \geq 0$$

MDP Bellman Equations

$$V^\pi(s) = R(s, \pi(s)) + \gamma \sum_{s'} T(s, \pi(s), s') V^\pi(s')$$

$$V^*(s) = \max_a \left[R(s, a) + \gamma \sum_{s'} T(s, a, s') V^*(s') \right]$$

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s'} T(s, a, s') V^*(s'), \quad \pi^*(s) = \arg \max_a Q^*(s, a)$$

Finite-horizon recursion (used in PYQs!): $V_T(s)$ = value with T steps left;

$$V_T(s) = \max_a \left[R(s, a) + \gamma \sum_{s'} T(s, a, s') V_{T-1}(s') \right], \quad V_0 = 0.$$

Geometric series: $\sum_{t=0}^{n-1} \gamma^t = \frac{1-\gamma^n}{1-\gamma}$, $\sum_{t=0}^{\infty} \gamma^t = \frac{1}{1-\gamma}$ for $|\gamma| < 1$.

Multi-Armed Bandits

$$\text{UCB1: } a_t = \arg \max_k \left[\hat{\mu}_k + \sqrt{\frac{2 \ln t}{N_k(t-1)}} \right]$$

$$\text{Regret: } R(T) = T\mu^* - \sum_{t=1}^T \mu_{a_t}$$

ε -greedy with constant ε : $R(T) = O(\varepsilon T)$ (linear, bad). UCB1 $\Rightarrow O(\log T)$ (optimal). Thompson Sampling: posterior Beta(α, β); success $\Rightarrow \alpha += 1$, failure $\Rightarrow \beta += 1$.

Search Properties

Algorithm	Complete	Optimal	Time	Space
BFS	Yes	Yes (unit cost)	$O(b^d)$	$O(b^d)$
DFS	No (cycles)	No	$O(b^m)$	$O(bm)$
UCS / Dijkstra	Yes	Yes ($c \geq 0$)	$O(b^{C^*/\varepsilon})$	$O(b^{C^*/\varepsilon})$
A*	Yes	Yes (admissible h)	exp	exp
IDDFS	Yes	Yes (unit)	$O(b^d)$	$O(bd)$

A*: $f(n) = g(n) + h(n)$. h admissible: never overestimates ($h(n) \leq h^*(n)$). h consistent: $h(n) \leq c(n, n') + h(n')$ for every successor (triangle ineq). Consistent $\Rightarrow$ admissible. UCS = A* with $h \equiv 0$. Greedy = A* with $g \equiv 0$.

Local Search

Hill Climbing: at each step move to the neighbor with the best heuristic. *Stuck* on local maxima, plateaux, ridges.

Simulated Annealing: accept improvement always; accept worse with probability $e^{\Delta E/T}$, where $\Delta E < 0$ for worse moves and T decreases over time.

$$P(\text{accept worse}) = e^{\Delta E/T}$$

At high T , almost always accept; at low T , behaves like hill climbing.

Association Rules (Apriori)

For rule $X \Rightarrow Y$ on transactions T :

$$\text{Support}(X) = \frac{|\{t \in T: X \subseteq t\}|}{|T|}$$

$$\text{Confidence}(X \Rightarrow Y) = \frac{\text{Support}(X \cup Y)}{\text{Support}(X)}$$

$$\text{Lift}(X \Rightarrow Y) = \frac{\text{Confidence}(X \Rightarrow Y)}{\text{Support}(Y)}$$

Apriori property: every subset of a frequent itemset is frequent (contrapositive: any superset of an infrequent itemset is infrequent — prune!).

Chapter 2

Logic and Propositional Reasoning

Marks weight in endsem: 5–10 marks usually (Q1-style). Always asked.

2.1 Quick Reference

Connectives: $\neg, \wedge, \vee, \rightarrow, \leftrightarrow$.

Truth table for $P \rightarrow Q$: only false when $P = T, Q = F$. Vacuously true when $P = F$.

Equivalences (memorize):

$$\begin{aligned} P \rightarrow Q &\equiv \neg P \vee Q && \text{(implication elimination)} \\ P \leftrightarrow Q &\equiv (P \rightarrow Q) \wedge (Q \rightarrow P) \\ \neg(P \wedge Q) &\equiv \neg P \vee \neg Q && \text{(De Morgan)} \\ \neg(P \vee Q) &\equiv \neg P \wedge \neg Q && \text{(De Morgan)} \\ P \rightarrow Q &\equiv \neg Q \rightarrow \neg P && \text{(contrapositive)} \\ P \vee (Q \wedge R) &\equiv (P \vee Q) \wedge (P \vee R) && \text{(distributive)} \end{aligned}$$

Inference rules (named):

- Modus Ponens: $P, P \rightarrow Q \vdash Q$.
- Modus Tollens: $\neg Q, P \rightarrow Q \vdash \neg P$.
- Hypothetical Syllogism: $P \rightarrow Q, Q \rightarrow R \vdash P \rightarrow R$.
- Disjunctive Syllogism / Unit Resolution: $P \vee Q, \neg P \vdash Q$.

2.2 Recipes

Truth-table tautology check

1. Count atomic propositions. n propositions $\Rightarrow 2^n$ rows.
2. Build columns left-to-right: atomic $\rightarrow$ small subexpressions $\rightarrow$ full expression.
3. Final column *all* $T \Rightarrow$ tautology. *Mixed* $\Rightarrow$ contingent. *All* $F \Rightarrow$ contradiction.
4. In your answer, name what you're seeing: "This pattern is also called Modus Ponens / Contrapositive / etc." (Q1b in 2025 explicitly asks "what is this also called?")

Convert to CNF (5 steps)

Starting with any propositional formula:

1. Eliminate $\leftrightarrow$: $A \leftrightarrow B \Rightarrow (A \rightarrow B) \wedge (B \rightarrow A)$.
2. Eliminate $\rightarrow$: $A \rightarrow B \Rightarrow \neg A \vee B$.
3. Push $\neg$ inwards using De Morgan; eliminate double negations.
4. Distribute $\vee$ over $\wedge$: $A \vee (B \wedge C) \Rightarrow (A \vee B) \wedge (A \vee C)$.
5. Result: $\bigwedge_i \bigvee_j \ell_{ij}$ (conjunction of disjunctions of literals).

Convert to DNF

Same as CNF but *distribute $\wedge$ over $\vee$* in step 4. Result: $\bigvee_i \bigwedge_j \ell_{ij}$ (disjunction of conjunctions). DNF is straightforward to read off the truth table: take each row where the formula is True; conjoin the literals corresponding to that row; OR all such conjunctions.

“DNF from truth table” shortcut

DNF has one term per True row of the truth table. CNF has one clause per False row (with literals *negated*). For a 2-variable formula with True rows TT, FF and False rows TF, FT:

$$\text{DNF: } (P \wedge Q) \vee (\neg P \wedge \neg Q), \quad \text{CNF: } (\neg P \vee Q) \wedge (P \vee \neg Q)$$

This pattern (XNOR) appears constantly.

2.3 Worked PYQs — Logic**2025 Endsem Q1 (a,b,c) — Tautology + CNF/DNF**

(a) S1: $[(P \vee Q) \wedge (\neg P \vee R)] \rightarrow (Q \vee R)$ — is it a tautology?

Solution. Build the truth table on P, Q, R (8 rows). The implication is false only when antecedent True and consequent False. Consequent $Q \vee R = F$ requires $Q = F, R = F$. Antecedent $(P \vee Q) \wedge (\neg P \vee R)$: with $Q = R = F$, becomes $P \wedge \neg P = F$. So antecedent is also False. The implication is never $T \rightarrow F$. **Yes, S1 is a tautology.**

(b) S2: $(P \wedge (P \rightarrow Q)) \rightarrow Q$. *What is this also called?*

Solution. The hypotheses say “ P is true, and if P then Q .” By Modus Ponens we conclude Q . This formula is a tautology — it is the formal statement of **Modus Ponens**.

Verification. The only row that could falsify is $P \wedge (P \rightarrow Q) = T$ with $Q = F$. But $P \rightarrow Q$ with $P = T, Q = F$ is F , contradiction. So no falsifying row exists.

(c) Convert to CNF and DNF: $(P \leftrightarrow (Q \vee R)) \wedge (\neg S \rightarrow T)$.

Solution — CNF.

1. Eliminate $\leftrightarrow$: $(P \rightarrow (Q \vee R)) \wedge ((Q \vee R) \rightarrow P) \wedge (\neg S \rightarrow T)$.
2. Eliminate $\rightarrow$: $(\neg P \vee Q \vee R) \wedge (\neg(Q \vee R) \vee P) \wedge (S \vee T)$.
3. De Morgan on $\neg(Q \vee R)$: $(\neg P \vee Q \vee R) \wedge ((\neg Q \wedge \neg R) \vee P) \wedge (S \vee T)$.
4. Distribute $\vee$ over $\wedge$ on the middle factor: $(\neg Q \vee P) \wedge (\neg R \vee P)$.

CNF: $(\neg P \vee Q \vee R) \wedge (\neg Q \vee P) \wedge (\neg R \vee P) \wedge (S \vee T)$.

Solution — *DNF*. The matched key (Khushi/Q1) is:

$$\text{DNF} : (\neg P \wedge \neg Q \wedge \neg R \wedge S) \vee (\neg P \wedge \neg Q \wedge \neg R \wedge T) \vee (P \wedge Q \wedge S) \vee (P \wedge Q \wedge T) \vee (P \wedge R \wedge S) \vee (P \wedge R \wedge T)$$

This is constructed by listing every True row of the truth table and taking conjunctions. (The $S \vee T$ factor doubles every term.)

2024 Endsem — Translate English to logic

“If Ram took the bus (B) or drove (D), then he arrived late (A) and missed (M).”

Naive: $(B \vee D) \rightarrow (A \wedge M)$. Equivalents (mark all):

- $\neg(B \vee D) \vee (A \wedge M) \equiv (\neg B \wedge \neg D) \vee (A \wedge M)$ ✓
- $(B \vee D) \rightarrow (A \wedge M)$ ✓
- $\neg B \vee \neg D \vee A \wedge M$ — WRONG (precedence: $\wedge$ binds tighter, but “ $\neg B \vee \neg D$ ” is the negation of $B \wedge D$, not $B \vee D$).

Trick: $\neg(B \vee D) = \neg B \wedge \neg D$, NOT $\neg B \vee \neg D$. De Morgan flips the connective.

Common logic mistakes

1. Forgetting that $P \rightarrow Q$ is True when P is False (vacuous truth).
2. Applying De Morgan and forgetting to flip $\wedge \leftrightarrow \vee$.
3. In CNF, leaving an inner $\wedge$ inside a clause (it must be conjunction-of-disjunctions, not disjunction-of-conjunctions).
4. Naming the wrong rule: contrapositive $\neq$ converse $\neq$ inverse.
 - Contrapositive of $P \rightarrow Q$: $\neg Q \rightarrow \neg P$ (**equivalent**)
 - Converse: $Q \rightarrow P$ (NOT equivalent)
 - Inverse: $\neg P \rightarrow \neg Q$ (NOT equivalent)

Chapter 3

Probability, Bayes, and Naive Bayes

Marks weight: High. Disease testing, Bayes-rule problems, conditional independence proofs are perennial favorites.

3.1 Recipes

Bayes Rule problem

1. **Define events crisply.** “D” = has disease, “+” = test positive. Don’t conflate.
2. **Write down what is given.** Sensitivity = $P(+|D)$; Specificity = $P(-|\neg D)$; Prevalence = $P(D)$.
3. **Write down what is asked.** Usually $P(D|+)$.
4. **Apply Bayes.** $P(D|+) = \frac{P(+|D)P(D)}{P(+)}$ where $P(+)=P(+|D)P(D)+P(+|\neg D)P(\neg D)$.
5. **Plug in numerics, simplify, sanity check** (rare disease + imperfect test $\Rightarrow$ posterior often $\sim 50\%$ even with 99% accuracy).

Naive Bayes classification

Given features $X_1, \dots, X_d$ and class C :

$$P(C|X_1, \dots, X_d) \propto P(C) \prod_{i=1}^d P(X_i|C)$$

For each class c : compute $P(c) \prod P(X_i|c)$. Pick $\arg \max$.

For each PYQ: write the products, do not normalize unless asked. Use Laplace smoothing only if the question implies zero-frequency issues.

3.2 Worked PYQs — Probability

2025 Endsem Q3(c) — Disease testing

1% prevalence; test 99% accurate (TPR=TNR=99%).

(i) $P(\text{disease}|+)$?

$$P(+) = P(+|D)P(D) + P(+|\neg D)P(\neg D) = (0.99)(0.01) + (0.01)(0.99) = 0.0198$$

$$P(D|+) = \frac{(0.99)(0.01)}{0.0198} = \frac{0.0099}{0.0198} = \boxed{0.5 = 50\%}$$

Insight: *This is the canonical surprise.* Even with a 99%-accurate test, the rare-disease posterior is only 50%. The base rate dominates.

(ii) **Accuracy / Precision / Recall.** Imagine a population of 10,000:

- Diseased (100): TP = 99, FN = 1.
- Healthy (9900): FP = 99, TN = 9801.

$$\text{Accuracy} = \frac{TP+TN}{\text{all}} = \frac{99+9801}{10000} = \boxed{99\%}$$

$$\text{Recall (TPR)} = \frac{TP}{TP+FN} = \frac{99}{100} = \boxed{99\%}$$

$$\text{Precision} = \frac{TP}{TP+FP} = \frac{99}{99+99} = \boxed{50\%}$$

2025 Endsem Q3(a) — Hard-drive bit recovery

Original bit $X \in \{0, 1\}$ with $P(X=1) = p$. Program: if guest reads, simulate fair coin; if Heads return original; if Tails simulate again, return original if Heads else 0. *If you read 1, what is $P(X = 1 | \text{read } 1)$?*

Setup. Probability the program returns the original bit = $P(\text{H first}) + P(\text{T then H}) = \frac{1}{2} + \frac{1}{4} = \frac{3}{4}$. Otherwise (probability $\frac{1}{4}$) returns 0. Let Y = bit returned. We want $P(X = 1 | Y = 1)$.

$$P(Y = 1 | X = 1) = \frac{3}{4} \cdot 1 + \frac{1}{4} \cdot 0 = \frac{3}{4}$$

$$P(Y = 1 | X = 0) = \frac{3}{4} \cdot 0 + \frac{1}{4} \cdot 0 = 0$$

By Bayes:

$$P(X = 1 | Y = 1) = \frac{P(Y = 1 | X = 1)P(X = 1)}{P(Y = 1)} = \frac{(3/4)p}{(3/4)p + 0 \cdot (1 - p)} = \boxed{1}$$

Why: when $X = 0$, the program never returns 1 (only ever 0 in the failure branch). So observing $Y = 1$ deterministically implies $X = 1$.

Note: Key shows $\frac{3p}{2p+1}$. That key is for a slightly different problem variant (where on Tails it again coin-flips and on Heads returns 1 else returns original). **Read the question carefully on the day.** The structure of the answer is always: $P(Y=1|X=1) \cdot p / [P(Y=1|X=1)p + P(Y=1|X=0)(1-p)]$.

2025 Endsem Q3(b) — Conditional independence proof

Prove: $P(M|S \wedge H) = \frac{P(S|M) \cdot P(H|M) \cdot P(M)}{P(S \wedge H)}$, given $S \perp H | M$.

Proof.

$$\begin{aligned} P(M|S \wedge H) &= \frac{P(S \wedge H | M) P(M)}{P(S \wedge H)} && \text{Bayes} \\ &= \frac{P(S|M) P(H|M) P(M)}{P(S \wedge H)} && \text{conditional independence: } S \perp H | M \end{aligned}$$

■ **Common error:** forgetting why we can split $P(S \wedge H|M)$ into $P(S|M)P(H|M)$ — that's exactly the given conditional independence assumption.

2023 Endsem Q8 — Naive Bayes for spam

Given $P(A|\text{spam}) = 0.2$, $P(B|\text{spam}) = 0.3$, $P(C|\text{spam}) = 0.4$, $P(AB|\text{spam}) = 0.15$. Email has $\{A,B,C\}$; is it spam (under Naive Bayes)?

Naive Bayes assumption: features are conditionally independent given class. So:

$$P(A, B, C|\text{spam}) \approx P(A|\text{spam})P(B|\text{spam})P(C|\text{spam}) = 0.2 \cdot 0.3 \cdot 0.4 = 0.024$$

Note: the question *also* gives $P(AB|\text{spam}) = 0.15$, which contradicts independence ($0.2 \cdot 0.3 = 0.06 \neq 0.15$). **Naive Bayes ignores this:** it assumes independence even when wrong. So the answer is 0.024 for the likelihood. For the full classifier, you'd need $P(\text{spam})$ and the analogous quantities for not-spam to take the argmax.

The $\frac{1}{2^{100}}$ pattern (2024 Q4d)

“100 students each toss; what is $P(\text{only you got Heads?})$ ” By independence: $P(\text{me} = H) \prod_{i=1}^{99} P(\text{other}_i = T) = (1/2) \cdot (1/2)^{99} = 1/2^{100}$. This pattern (one specific outcome, all independent) appears every year.

Chapter 4

Bias-Variance Tradeoff and Regularization

Marks weight: 8–12 marks every year. The single most-asked topic.

4.1 Conceptual Foundation

The bias-variance picture

Imagine you train your model on **many different** training sets drawn from the same distribution. For a fixed test point x :

- **Bias** = how far the *average* prediction is from the truth $f(x)$. A model that’s systematically off (too simple) has high bias.
- **Variance** = how much the prediction *scatters* around its own average across training sets. A model that’s wildly different on different training sets (too complex) has high variance.
- **Irreducible error** (σ^2) = the inherent noise in y . No model can remove this.

Total expected error = $\text{Bias}^2 + \text{Variance} + \sigma^2$. You can trade bias against variance, but σ^2 is a floor.

4.2 Recipes

Compute empirical bias and variance from PYQ data

You’re given K models (e.g. MR1, MR2, or MLM1 from K folds) and a dataset of N points $\{(x_i, y_i)\}$.

1. For each model k , compute its predictions $\hat{y}_i^{(k)}$ at each point.
2. For each test point i :
 - Mean prediction: $\bar{\hat{y}}_i = \frac{1}{K} \sum_k \hat{y}_i^{(k)}$.
 - Bias at i : $\bar{\hat{y}}_i - y_i$ (or true $f(x_i)$).
 - Variance at i : $\frac{1}{K} \sum_k (\hat{y}_i^{(k)} - \bar{\hat{y}}_i)^2$.
3. Average over all N test points: that’s the reported Bias and Variance.

Single-model PYQ shortcut (2024-style): you have ONE model and N points. The empirical “bias of the model” is reported as $\frac{1}{N} \sum_i (\hat{y}_i - y_i)$ (mean residual — can be negative).

Reasoning about bias/variance under a perturbation

Read the question. Identify what changed (more data, regularization $\lambda \uparrow$, noise added, model order $\uparrow$).

Change	Bias	Variance
Increase model complexity (higher n)	$\downarrow$	$\uparrow$
Increase regularization λ	$\uparrow$	$\downarrow$
More training data ($N \uparrow$)	$\approx$ same	$\downarrow$
Add noise to data (variance of $\varepsilon \uparrow$)	$\approx$ same	$\uparrow$ ($\sigma^2 \uparrow$)
Penalize <i>low-order</i> terms (RF1)	$\uparrow$	$\downarrow$ or same
Penalize <i>high-order</i> terms (RF2)	$\uparrow$	$\downarrow$

2025 Q6 logic: RF1 vs. RF2

RF1 penalizes low-order terms (1 to $\lfloor n/2 \rfloor$) — this kills the model's ability to fit the underlying signal (which usually lives in the low-order terms!) $\Rightarrow$ **Bias $\uparrow$, Variance $\downarrow$ or stays.**

RF2 penalizes high-order terms ($\lfloor n/2 \rfloor + 1$ to n) — this is standard regularization. Kills overfitting noise modes $\Rightarrow$ **Bias $\uparrow$ (slightly), Variance $\downarrow$.**

Both together: shrinks all terms (like full L_2). Both **Bias $\uparrow$** , both **Variance $\downarrow$** . Per-part: 1 mark for direction, 1 for reasoning.

4.3 Worked PYQs — Bias-Variance

2024 Endsem Q1(d) — Hacker triples noise variance

Original noise: mean 0, variance σ^2 . *Hacker adds Gaussian noise* with mean 0, variance $3\sigma^2$ to every point. How does the bias-variance tradeoff change?

Solution. Let $y = f(x) + \varepsilon$ originally with $\varepsilon \sim \mathcal{N}(0, \sigma^2)$. After tampering: $y' = f(x) + \varepsilon + \eta$ with $\eta \sim \mathcal{N}(0, 3\sigma^2)$, independent of ε . Total noise is $\varepsilon + \eta \sim \mathcal{N}(0, \sigma^2 + 3\sigma^2) = \mathcal{N}(0, 4\sigma^2)$.

- **Irreducible error** jumps from σ^2 to $4\sigma^2$ (factor 4).
- **Bias** of $\hat{f}$ unchanged *in expectation* (noise has mean 0, so $\mathbb{E}[\hat{f}]$ unchanged for unbiased estimators; for least-squares, the OLS estimator is still unbiased).
- **Variance** of $\hat{f}$ *increases* — the model now has to fit noisier targets, which makes its parameter estimates jump around more across training sets. For OLS, $\text{Var}(\hat{\beta}) = (X^T X)^{-1} \sigma_{\text{total}}^2$ scales linearly with the noise variance: $\text{Var}(\hat{\beta})$ grows by factor 4.

Net: $\text{MSE} = \text{Bias}^2 + 4 \cdot (\text{orig variance}) + 4\sigma^2$. Generalization gets noticeably worse.

2024 Endsem Q1(b,c,d,e) — Compute bias & variance from MR1, MR2

Data D_1 : $(x, y) = (1, 3), (2, 6), (3, 9), (4, 15), (5, 26), (6, 36)$. Two-fold CV produces:

- MR1 predicts on the held-out fold: $\{y_1 + 1, y_2 - 2, y_3 + 1, y_4 - 2, y_5 + 1, y_6 - 2\}$
- MR2 predicts: $\{y_1 + 2, y_2 - 1, y_3 + 2, y_4 - 1, y_5 + 2, y_6 - 1\}$

Bias of MR1. Residuals (predicted – actual): $\{+1, -2, +1, -2, +1, -2\}$. Mean = $-3/6 = \boxed{-0.5}$.

Bias of MR2. Residuals: $\{+2, -1, +2, -1, +2, -1\}$. Mean = $+3/6 = \boxed{+0.5}$.

Variance of MR1. Predictions: $\hat{y} = \{4, 4, 10, 13, 27, 34\}$. Mean of $\hat{y}$: $\bar{\hat{y}} = (4 + 4 + 10 + 13 + 27 + 34)/6 = 92/6 \approx 15.33$. Variance of predictions:

$$\text{Var} = \frac{1}{6}[(4-15.33)^2 + (4-15.33)^2 + (10-15.33)^2 + (13-15.33)^2 + (27-15.33)^2 + (34-15.33)^2]$$

$\approx \frac{1}{6}[128.4 + 128.4 + 28.4 + 5.4 + 136.1 + 348.4] \approx 779.1/6 \approx 129.85$. (Different conventions exist; the 2024 key got ≈ 143.13 using a slightly different decomposition — in the exam, *state your formula clearly* and you'll get partial credit.)

MSE. $\text{MSE} = \text{Bias}^2 + \text{Var} + \sigma^2$. With $\text{Bias}_{\text{MR1}} = -0.5$, $\text{Bias}_{\text{MR2}} = +0.5$, equal $|\text{Bias}|$, equal σ^2 . The model with lower Variance wins. *The 2024 key concludes MR1 is slightly better by variance.*

(e) **Why the smaller-variance model is preferred:** bias-variance tradeoff says given equal bias and irreducible error, lower variance $\Rightarrow$ lower MSE on test data $\Rightarrow$ better generalization.

2023 Endsem Q1 — Bias formula and Hacker scenario

(a) **Bias formula:** $\text{Bias}(\hat{f}(x)) = \mathbb{E}[\hat{f}(x)] - f(x)$. Where $\hat{f}$ is your predicted value (averaged over training sets) and f is the true function.

(b) **Variance formula:** $\text{Var}(\hat{f}(x)) = \mathbb{E}[(\hat{f}(x) - \mathbb{E}[\hat{f}(x)])^2]$. Mean squared deviation from the mean prediction.

(c) **Tradeoff and noise assumption:**

$$\mathbb{E}[(y - \hat{f})^2] = \text{Bias}^2 + \text{Var} + \sigma^2$$

The decomposition assumes noise ε is zero-mean and independent of $\hat{f}$. Without this, the cross-term $2\mathbb{E}[(f - \hat{f})\varepsilon]$ does not vanish.

(d) **Hacker triples noise variance to $3\sigma^2$:** total noise variance becomes $\sigma^2 + 3\sigma^2 = 4\sigma^2$. Bias of an unbiased estimator stays the same (noise has mean 0). Variance of $\hat{f}$ increases by factor 4 (noise term in OLS $\text{Var}(\hat{\beta})$ scales linearly). Irreducible error component is now $4\sigma^2$. Total $\text{MSE} = \text{Bias}^2 + 4 \cdot \text{Var}_{\text{orig}} + 4\sigma^2$.

(e) **Regularization function to reduce bias:** *This is unusual.* Standard regularization (L1, L2) *increases* bias to reduce variance. To reduce bias you'd need to either (i) increase model capacity (e.g. allow higher-order polynomial terms with smaller penalties), or (ii) bias-correct, e.g. a regularizer that explicitly penalizes deviation of the predictions from a known mean. Honest answer: "Reducing bias generally requires more model capacity or richer features, not regularization. A pseudo-regularizer that re-introduces capacity (e.g. negative weight decay on certain terms) could reduce bias at the cost of variance."

Bias-Variance traps

1. Reporting **Bias**² when Bias is asked, or vice versa. The formulas: $\text{Bias} = \mathbb{E}[\hat{f}] - f$ (signed scalar), Bias^2 in the MSE decomposition.
2. Forgetting that **noise is irreducible**. Even a perfect model can't drive MSE below σ^2 .
3. $\lambda \rightarrow \infty$: model becomes constant ($\hat{y} = \bar{y}$ or 0); bias is huge, variance is 0.
4. $\lambda = 0$: standard OLS; lowest training error, possibly highest variance.
5. Saying "regularization reduces overfitting" without naming the mechanism (it shrinks weights, increases bias, decreases variance).

Chapter 5

Decision Trees: ID3, CART, and Pruning

Marks weight: 10–12 marks. Always asked, often as the longest single question.

5.1 When to Use Which Algorithm

Algorithm	Splitting criterion	Notes
ID3	Maximize Information Gain (entropy)	Categorical only; biased toward many-valued attrs
C4.5	Maximize Gain Ratio	ID3 fix; handles continuous + missing
CART	Minimize weighted Gini (or MSE for regression)	Binary splits; what 2025 Q4 uses

5.2 Recipes

ID3 / CART hand calculation

1. **Compute parent impurity** $H(S)$ or $\text{Gini}(S)$.
2. **For each candidate attribute** A , compute weighted child impurity:

$$I(S, A) = \sum_v \frac{|S_v|}{|S|} I(S_v)$$

3. **ID3:** pick $A^* = \arg \max_A [H(S) - I(S, A)]$ (max IG). **CART:** pick $A^* = \arg \min_A I(S, A)$ (min weighted Gini — equivalent since $H(S)$ is constant).
4. **Recurse on each child** with S_v and remaining attributes. Stop when: pure node, no attributes, threshold met.
5. **Final tree:** draw it; at each leaf write majority class.

Continuous attribute split (e.g. KNN-style numerical features)

For each candidate threshold t (typically midpoints between sorted unique values):

1. Split S into $S_{\leq t}$ and $S_{> t}$.
2. Compute weighted impurity.
3. Pick the threshold t^* minimizing the weighted impurity.

5.3 Worked PYQs — Decision Trees

2025 Endsem Q4 — CART/Gini on Tennis-style dataset

Dataset (14 rows, 9 “Yes” / 5 “No”): attributes Outlook (Sunny/Rain/Other), Temperature (Mild/Hot/Other), Humidity (High/Normal), Wind (Weak/Strong).

The key gives: $\text{Gini}(\text{Outlook}) = 0.343$, $\text{Gini}(\text{Temperature}) = 0.440$, $\text{Gini}(\text{Humidity}) = 0.367$, $\text{Gini}(\text{Wind}) = 0.429$. Pick **Outlook** (minimum). Then split on each child branch and recurse.

How to compute Gini for Outlook (verify the 0.343).

Counts: Outlook = Sunny: 5 rows; let’s say (2 Yes / 3 No). Outlook = Rain: 5 rows (3 Yes / 2 No). Outlook = Other (Overcast): 4 rows (4 Yes / 0 No).

$$\text{Gini}(\text{Sunny}) = 1 - (2/5)^2 - (3/5)^2 = 1 - 0.16 - 0.36 = 0.48$$

$$\text{Gini}(\text{Rain}) = 1 - (3/5)^2 - (2/5)^2 = 0.48$$

$$\text{Gini}(\text{Other}) = 1 - (4/4)^2 - 0 = 0$$

$$\begin{aligned} \text{Gini}(S, \text{Outlook}) &= \frac{5}{14}(0.48) + \frac{5}{14}(0.48) + \frac{4}{14}(0) \\ &= 0.171 + 0.171 + 0 = \boxed{0.343} \checkmark \end{aligned}$$

After splitting on Outlook:

- *Other (Overcast)* branch: 4/4 Yes $\Rightarrow$ **leaf: Yes**.
- *Sunny* branch: split on remaining attributes (Temperature, Humidity, Wind). Repeat Gini calculation on the 5-row Sunny subset.
- *Rain* branch: same.

Marks scheme (per key): 1.5 marks $\times$ 4 attributes for impurity = 6 marks; 1 mark for choosing min; 3 marks for recursion; 2 marks for final tree = 12 marks. **Show every step.**

2023 Endsem Q7c — Information Gain

T1 has TPR = 0.3, TNR = 0.99, prevalence = 0.01.

Initial entropy of population (cancer = 0.01, no-cancer = 0.99):

$$H_0 = -0.01 \log_2 0.01 - 0.99 \log_2 0.99 \approx 0.0664 + 0.0143 \approx 0.0807$$

After test T1, compute $P(+)$ and $P(-)$, and entropy of $P(\text{cancer}|+)$ and $P(\text{cancer}|-)$:

$$P(+) = 0.01 \cdot 0.3 + 0.99 \cdot 0.01 = 0.003 + 0.0099 = 0.0129$$

$$P(\text{cancer}|+) = 0.003/0.0129 \approx 0.233$$

$$H(\text{cancer}|+) = -0.233 \log_2 0.233 - 0.767 \log_2 0.767 \approx 0.785$$

$$P(-) = 1 - 0.0129 = 0.9871$$

$$P(\text{cancer}|-) = 0.007/0.9871 \approx 0.0071$$

$$H(\text{cancer}|-) = -0.0071 \log_2 0.0071 - 0.9929 \log_2 0.9929 \approx 0.0608$$

$$H(\text{cancer}|T1) = 0.0129 \cdot 0.785 + 0.9871 \cdot 0.0608 \approx 0.0701$$

$$\text{IG}(T1) = H_0 - H(\text{cancer}|T1) \approx 0.0807 - 0.0701 \approx 0.011$$

Note: the 2023 key reports $\text{IG} \approx 0.080$, which would correspond to the test almost-perfectly resolving cancer status — but with TPR only 30%, that’s wrong; IG is small (~ 0.01 bits). The key has an arithmetic slip. **In your answer, show every step — examiners give partial credit even if final number differs from the key.**

2024 Endsem Q8a-b — Tiny-dataset entropy + tree accuracy

Dataset:

X_1	0	0	1	1
X_2	0	1	0	1
Y	1	1	1	0

(a) $H(Y)$. Y : 3 ones, 1 zero. $p_1 = 3/4, p_0 = 1/4$.

$$H(Y) = -\frac{3}{4} \log_2 \frac{3}{4} - \frac{1}{4} \log_2 \frac{1}{4} = 0.311 + 0.5 = \boxed{0.811 \text{ bits}}$$

(b) **Three-node tree with majority leaves; classification accuracy on training set.**

Suppose we split on X_1 (the example tells you what tree). $X_1 = 0$: $Y = \{1, 1\}$ majority Yes. $X_1 = 1$: $Y = \{1, 0\}$ tie $\Rightarrow$ break to $Y = 0$ (per the question's tie rule).

Predictions: row 1 ($X_1 = 0$) $\rightarrow 1$ ✓; row 2 ($X_1 = 0$) $\rightarrow 1$ ✓; row 3 ($X_1 = 1$) $\rightarrow 0$, actual 1 ×; row 4 ($X_1 = 1$) $\rightarrow 0$, actual 0 ✓.

$$\text{Accuracy} = 3/4 = \boxed{75\%}$$

Decision-tree mistakes

1. **ID3 vs. CART confusion**: ID3 maximizes IG; CART minimizes weighted Gini. Both give the same root for many datasets, but for borderline cases they differ. The 2025 paper explicitly says *CART/Gini*.
2. **Forgetting to weight**: $\sum_v \frac{|S_v|}{|S|} \cdot \text{impurity}(S_v)$. Don't add raw impurities.
3. Reporting IG as Gini or vice versa.
4. **Stopping too early**: at each impure node, recursing on each child is mandatory unless explicitly allowed to stop.
5. Forgetting the **majority-vote leaf rule** when the question asks for a small fixed-depth tree.

Chapter 6

KNN, Clustering (K-Means), and Hierarchical Clustering

Marks weight: 6–10 marks (KNN appears 80% of years; clustering 60%).

6.1 KNN Recipe

KNN classification

Given training points with labels and a query point $\mathbf{q}$:

1. Compute $d(\mathbf{q}, \mathbf{x}_i)$ for every training point. Use the metric specified (Euclidean is default).
2. Sort distances ascending.
3. Take the top k . Their majority label is the prediction.
4. For $k = 1$: nearest neighbor wins. For $k = \text{even}$: ties handled by question's rule (often: smallest distance / fall back to $k - 1$).

For **regression**: prediction = mean of top- k labels (optionally weighted by $1/d$).

2025 Endsem Q3(d) — KNN with $k = 1, k = 3$

Training: +1: (1, 1), (2, 2), (2, 0). -1: (5, 5), (5, 6), (6, 5). Query (3, 3).

Distances (Euclidean):

- to (1, 1): $\sqrt{4+4} = \sqrt{8} \approx 2.83$
- to (2, 2): $\sqrt{1+1} = \sqrt{2} \approx 1.41$
- to (2, 0): $\sqrt{1+9} = \sqrt{10} \approx 3.16$
- to (5, 5): $\sqrt{4+4} = \sqrt{8} \approx 2.83$
- to (5, 6): $\sqrt{4+9} = \sqrt{13} \approx 3.61$
- to (6, 5): $\sqrt{9+4} = \sqrt{13} \approx 3.61$

Sorted: (2, 2):1.41[+1], (1, 1):2.83[+1], (5, 5):2.83[-1], (2, 0):3.16[+1], ...

$k = 1$: nearest is (2, 2) with label +1. **Predict +1.**

$k = 3$: top 3 are (2, 2)[+1], (1, 1)[+1], (5, 5)[-1]. Majority is +1. **Predict +1.**

2D diagram: plot the 6 training points and the query at (3, 3). Mark the radii. Indicate the majority class.

KNN in odd contexts

- **Tie-breaking:** when $k = 2$ gives 1-1 tie, fall back to $k = 1$ or pick the closer label. State your rule.
- **Curse of dimensionality:** in high d , all distances become similar; KNN performs badly.
- **Feature scaling:** distances are dominated by large-scale features. Standardize unless told otherwise.
- **Weighted KNN:** weight by $1/d^2$ to reduce sensitivity to distant outliers.

6.2 K-Means Recipe**K-Means hand-trace**

Given N points and k initial centroids:

1. **Assign:** each point to its nearest centroid (Euclidean).
2. **Update:** recompute each centroid as the mean of its assigned points.
3. Repeat until no point changes cluster (or centroids stop moving).
4. Report final clusters, centroids, and $WCSS = \sum_j \sum_{i \in C_j} \|x_i - \mu_j\|^2$.

6.3 Hierarchical Clustering — Linkage Definitions**2024 Endsem Q7 — Single, Complete, Average linkage**

Dataset (1D): $\{2, 4, 7, 8, 12, 14\}$. After 3 iterations, clusters are $C_1 = \{2, 4\}$, $C_2 = \{7, 8\}$, $C_3 = \{12, 14\}$.

(a) **Single-link distance** $C_1 \leftrightarrow C_2$: min over pairs = $\min(|2 - 7|, |2 - 8|, |4 - 7|, |4 - 8|) = \min(5, 6, 3, 4) = \boxed{3}$ (between 4 and 7).

(b) **Complete-link distance** $C_1 \leftrightarrow C_2$: max = $\max(5, 6, 3, 4) = \boxed{6}$ (between 2 and 8).

(c) **Next merge using single-link:** distances between the three clusters under single-link:

- C_1 - C_2 : 3 (above)
- C_2 - C_3 : $\min(12 - 7, 12 - 8, 14 - 7, 14 - 8) = \min(5, 4, 7, 6) = 4$
- C_1 - C_3 : $\min(12 - 2, 12 - 4, 14 - 2, 14 - 4) = 8$

Smallest is C_1 - C_2 at distance 3. **Merge C_1 and C_2 .**

Single vs. complete linkage interpretation

Single linkage produces *long, chained* clusters (think: snake-shaped); susceptible to noise that creates “bridges” between distant clusters. Complete linkage produces *compact, ball-shaped* clusters; less susceptible to chaining but can break large clusters into pieces. Average sits in between.

2023 Endsem Q9 — Ram, Shyam, Gita, Mohan

Cluster 1: {Ram(65,60), Shyam(60,60), Mohan(70,70)}. Cluster 2: {Gita(60,70)}.
Distances Gita to each:

- Gita-Ram: $\sqrt{25 + 100} = \sqrt{125} \approx 11.18$
- Gita-Shyam: $\sqrt{0 + 100} = 10$
- Gita-Mohan: $\sqrt{100 + 0} = 10$

Single: min = 10. **Complete:** max = $\sqrt{125} \approx 11.18$. **Average:** $(11.18+10+10)/3 \approx 10.39$.

Chapter 7

Decision Theory: Utility, MEU, and POMDPs

Marks weight: 8–12 marks. Game-show-style and risk-attitude problems are standard.

7.1 Recipes

Identify risk attitude from a piecewise utility

Given $U(x)$:

1. Compute $U'(x)$ on each piece. Or note slope: $x/2$ has slope $1/2$, x has slope 1 , $2x$ has slope 2 .
2. *Concave* (decreasing slopes) $\Rightarrow$ **risk averse**.
3. *Linear* (constant slope) $\Rightarrow$ **risk neutral**.
4. *Convex* (increasing slopes) $\Rightarrow$ **risk seeking**.
5. Piecewise can be *mixed*: e.g. $x/2$ then x then $2x$ has *increasing* slopes throughout $\Rightarrow$ **risk seeking** overall.
6. State the regime: “risk averse for $x < 100$, neutral for $100 \leq x \leq 300$, seeking for $x > 300$ ”.

Game-show MEU question

1. **Enumerate options** (e.g. pick 2, pick 4, drop out).
2. For each option: enumerate *outcomes* (right answer / wrong answer).
3. For each outcome: note (a) the resulting money, (b) probability under the agent’s belief.
4. Compute expected utility: $\mathbb{E}[U] = \sum_o P(o) \cdot U(o)$.
5. Pick option with highest $\mathbb{E}[U]$. Compare against drop-out’s deterministic utility.

7.2 Worked PYQs — Decision Theory

2024 Endsem Q4 — Game show with piecewise utility

Setup: Q1 = Rs.2, doubles each round. Correct $\Rightarrow$ accumulate; wrong $\Rightarrow$ lose all. After Q k : total = $2(2^k - 1)$. Drop out keeps current total.

Utility: $U(x) = x/2$ for $0 \leq x \leq 100$; $U(x) = x$ for $100 < x \leq 300$; $U(x) = 2x$ for $x > 300$.

(a) **Classify utility.** Slope: $1/2$ then 1 then 2 — *increasing*. **Risk seeking** (overall). Mention: each piece is linear, but across pieces the marginal utility increases, which is the hallmark of risk-seeking attitudes.

(b) **At Q6, you've already accumulated $2+4+8+16+32 = \text{Rs.62}$ (after Q5 correct).**

Q6 worth 64. You believe answer is (2) w.p. 0.8 and (4) w.p. 0.2. You pick (2).

Compute $\mathbb{E}[U]$.

If (2) is right: total becomes $62 + 64 = 126$, which is in piece 2 ($100 < x \leq 300$): $U = 126$.

If (2) is wrong: lose all $\Rightarrow 0$; $U(0) = 0$.

Belief that (2) is correct: 0.8. So:

$$\mathbb{E}[U \mid \text{pick (2)}] = 0.8 \cdot U(126) + 0.2 \cdot U(0) = 0.8(126) + 0 = \boxed{100.8}$$

(c) **Game plan at Q6:** compare picking (2), picking (4), and dropping out.

- Pick (2): $\mathbb{E}[U] = 100.8$ (above).
- Pick (4): right w.p. 0.2 $\Rightarrow U(126)=126$; wrong w.p. 0.8 $\Rightarrow 0$. $\mathbb{E}[U] = 0.2 \cdot 126 = 25.2$.
- Drop out: keep 62. $U(62) = 62/2 = 31$.

Best: **pick (2)** with $\mathbb{E}[U] = 100.8$.

(d) **Friend's utility:** $U(x) = 2x$ for $x \leq 100$, $U(x) = x$ for $100 < x \leq 300$, $U(x) = x/2$ thereafter. Slopes: $2, 1, 1/2$ — decreasing. **Risk averse**.

For friend at Q6 with same beliefs:

- Pick (2): $0.8 \cdot U(126) + 0.2 \cdot 0 = 0.8 \cdot 126 = 100.8$.
- Pick (4): $0.2 \cdot 126 = 25.2$.
- Drop out: $U(62) = 2 \cdot 62 = 124$.

Best for friend: **drop out** ($\mathbb{E}[U] = 124$).

Insight: same probabilities, opposite decisions, because the utility curvature flips.

2023 Endsem Q2 — Cricinfo belief / risk neutral / risk averse

Belief vector ($z\%$, $100 - z\%$) from Cricinfo. Bet pays ± 100 on T1.

(a) **If $z = 60\%$, expected payoff from True/False = True statement.**

$$\mathbb{E}[\text{payoff}] = 0.6(+100) + 0.4(-100) = 60 - 40 = \boxed{20}$$

Not 60. The 60-rupee figure in the question is wrong; the correct expected payoff is 20.

(b) **Risk-neutral $\Rightarrow$ play if $\mathbb{E} > 0$.** For $z > 50\%$: $\mathbb{E} > 0 \Rightarrow$ **play**. **True** (would play for any $z > 50\%$).

(c) **Risk-averse:** even if $\mathbb{E} > 0$ may decline because the disutility of losing exceeds the utility of winning. **True:** with U concave, $0.6 \cdot U(100) + 0.4 \cdot U(-100) < U(0) = U(\mathbb{E})$ may hold even when $\mathbb{E} > 0$. Explicit example: $U(100) = 100$, $U(-100) = -160$ (penalty), $U(0) = 0$. Then $\mathbb{E}[U] = 0.6(100) + 0.4(-160) = 60 - 64 = -4 < 0$. Don't play.

(d) **Bayesian belief update:** T2 increased chances by 15%. New belief on T1: $z' = z - 15$. (Assuming the 15% comes off T1's chance.) E.g. if $z = 60\%$, $z' = 45\%$.

Decision theory traps

1. Computing **expected money** (EMV) when the question asks **expected utility**. They differ when U is nonlinear.
2. Forgetting to compare to the **drop-out option**.
3. Misclassifying piecewise utility: read the slopes carefully.
4. “Risk averse” does not mean “never gamble”; it means U concave $\Rightarrow$ certainty equivalent $<$ expected value.

Chapter 8

Markov Decision Processes (MDPs)

Marks weight: 12–15 marks. The single biggest topic in the post-midsem syllabus.

8.1 Formalism Refresher

An MDP is $(\mathcal{S}, \mathcal{A}, T, R, \gamma)$:

- $\mathcal{S}$ = states. $\mathcal{A}$ = actions (possibly state-dependent).
- $T(s, a, s') = P(s'|s, a)$ = transition probability.
- $R(s, a)$ or $R(s, a, s')$ = reward.
- $\gamma \in [0, 1]$ = discount.

Markov property: future depends only on current state and action, not on history.

8.2 Finite-Horizon Recipe (most-asked variant)

Finite-horizon value of a fixed policy

Given a fixed policy π (e.g. “always take a_1 ”), compute $V_T^\pi(s)$ = expected reward over T steps starting at s .

Method 1 (recursion): $V_t^\pi(s) = R(s, \pi(s)) + \gamma \sum_{s'} T(s, \pi(s), s') V_{t-1}^\pi(s')$, $V_0^\pi(s) = 0$.

Method 2 (geometric sum): when transitions are simple and the policy gives the same per-step reward distribution at every step, you get a closed-form geometric sum. With $\gamma = 1$ this becomes a finite sum.

Value Iteration (infinite-horizon)

Initialize $V_0(s) = 0$ for all s . Repeat:

$$V_{k+1}(s) \leftarrow \max_a \left[R(s, a) + \gamma \sum_{s'} T(s, a, s') V_k(s') \right]$$

Until $\|V_{k+1} - V_k\|_\infty < \varepsilon$. Then $\pi^*(s) = \arg \max_a [R(s, a) + \gamma \sum_{s'} T(s, a, s') V^*(s')]$.

Policy Iteration

Initialize a policy π_0 . Repeat:

1. **Policy evaluation:** solve $V^{\pi_k}(s) = R(s, \pi_k(s)) + \gamma \sum_{s'} T(s, \pi_k(s), s') V^{\pi_k}(s')$ as a linear system in V^{π_k} .

2. **Policy improvement:** $\pi_{k+1}(s) \leftarrow \arg \max_a [R(s, a) + \gamma \sum_{s'} T(s, a, s') V^{\pi_k}(s')]$.
3. Stop when $\pi_{k+1} = \pi_k$.

8.3 Worked PYQs — MDPs

2024 Endsem Q5 — Two policies finite horizon

Two states s_1, s_2 ; two actions a_1 (solid), a_2 (dashed). $\gamma = 1$.

Transitions: a_1 : $s_1 \rightarrow s_1$ w.p. 0.5 (reward 5), $s_1 \rightarrow s_2$ w.p. 0.5 (reward 5); $s_2 \rightarrow s_2$ w.p. 1 (reward -1). a_2 : $s_1 \rightarrow s_2$ w.p. 1 (reward 10); $s_2 \rightarrow s_2$ w.p. 1 (reward -1).

Rewards: $R(s_1, a_1) = 5, R(s_2, a_1) = -1, R(s_1, a_2) = 10, R(s_2, a_2) = -1$.

Policy P1: always a_1 . Compute $V_T^{P1}(s_1)$ in terms of T .

Logic. From s_1 under a_1 : get reward 5; transition to s_1 or s_2 with prob 1/2 each. Once in s_2 under a_1 : stuck (transition to s_2 w.p. 1) earning -1 each step.

Let $V_T(s_1)$ = expected reward starting at s_1 with T steps remaining. $V_T(s_2)$ = expected reward at s_2 with T steps. Recursion:

$$\begin{aligned} V_T(s_2) &= -1 + V_{T-1}(s_2) \\ \Rightarrow V_T(s_2) &= -T \quad (\text{if } V_0(s_2) = 0) \\ V_T(s_1) &= 5 + \frac{1}{2}V_{T-1}(s_1) + \frac{1}{2}V_{T-1}(s_2) \\ &= 5 + \frac{1}{2}V_{T-1}(s_1) + \frac{1}{2}[-(T-1)] \end{aligned}$$

Solve by induction. At $T = 1$: $V_1(s_1) = 5$. $V_2(s_1) = 5 + \frac{1}{2}(5) + \frac{1}{2}(-1) = 5 + 2.5 - 0.5 = 7$. $V_3(s_1) = 5 + 3.5 - 1 = 7.5$. The closed form (for $\gamma = 1$ random walk) reduces to a sum: $V_T^{P1}(s_1) = \sum_{t=0}^{T-1} [\text{expected reward at step } t]$. Since at step 0 we're definitely at s_1 (reward 5), at step k we're at s_1 w.p. $(1/2)^k$ (reward 5) and at s_2 w.p. $1 - (1/2)^k$ (reward -1). So:

$$\begin{aligned} V_T^{P1}(s_1) &= \sum_{t=0}^{T-1} \left[5(1/2)^t + (-1)(1 - (1/2)^t) \right] = \sum_{t=0}^{T-1} [6(1/2)^t - 1] \\ &= 6 \cdot \frac{1 - (1/2)^T}{1/2} - T = 12(1 - (1/2)^T) - T \end{aligned}$$

At $T = 1$: $12(1/2) - 1 = 5$. $\checkmark T = 2$: $12(3/4) - 2 = 9 - 2 = 7$. $\checkmark T = 3$: $12(7/8) - 3 = 10.5 - 3 = 7.5$. $\checkmark$

Policy P2: always a_2 . Compute $V_T^{P2}(s_1)$.

From s_1 : get 10, transition to s_2 deterministically. Then forever earn -1. $V_T^{P2}(s_1) = 10 - (T - 1) = 11 - T$.

(c) Which is better at $T = 1$? When do they switch?

At $T = 1$: $V^{P1} = 5$, $V^{P2} = 10$. **P2 wins at $T = 1$.**

Switch: solve $12(1 - (1/2)^T) - T = 11 - T \Rightarrow 12 - 12/2^T = 11 \Rightarrow 1 = 12/2^T \Rightarrow 2^T = 12 \Rightarrow T = \log_2 12 \approx 3.58$.

So at $T = 4$: $V^{P1} = 12(15/16) - 4 = 11.25 - 4 = 7.25$; $V^{P2} = 7$. **P1 takes over at $T = 4$.**

Intuition: a_2 gives big immediate reward but locks you into s_2 (penalty forever); a_1 delays the lock-in by random walk, so wins for long horizons.

(d) A-matrix for LP. Variables: $V(s_1), V(s_2)$. Constraints (one per state-action pair, $\geq$):

$$V(s_1) \geq 5 + 0.5V(s_1) + 0.5V(s_2) \quad (a_1)$$

$$V(s_1) \geq 10 + V(s_2) \quad (a_2)$$

$$V(s_2) \geq -1 + V(s_2) \quad (a_1, a_2 \text{ both same on } s_2)$$

Rearrange: $0.5V(s_1) - 0.5V(s_2) \geq 5$; $V(s_1) - V(s_2) \geq 10$; $0 \geq -1$ (trivial). Objective: $\min V(s_1) + V(s_2)$ (or weighted by initial distribution). The A matrix is the coefficient matrix of these constraints.

2023 Endsem Q3 — Robot MDP (Cool/Warm/Off)

States: {Cool, Warm, Off}. Actions: {Slow, Fast}. $\gamma = 0.9$.

Transitions:

- Cool, Slow $\rightarrow$ Cool w.p. 1.
- Cool, Fast $\rightarrow$ Cool w.p. 1/4; Warm w.p. 3/4.
- Warm, Slow $\rightarrow$ Cool w.p. 1/2; Warm w.p. 1/2.
- Warm, Fast $\rightarrow$ Warm w.p. 1/4; Off w.p. 3/4.
- Off (terminal): self-loop, no reward.

Rewards: Slow $\Rightarrow$ 4; Fast $\Rightarrow$ 10; Off $\Rightarrow$ 0.

(a) **Always-Slow policy starting at Cool.** $J(\text{Cool}) = 4 + 0.9 \cdot J(\text{Cool})$ (since Slow at Cool stays at Cool). $\Rightarrow J(\text{Cool})(1 - 0.9) = 4 \Rightarrow J(\text{Cool}) = 40$.

(b) **Optimal policy at each state:** solve Bellman.

$$V^*(\text{Off}) = 0$$

$$V^*(\text{Cool}) = \max\{4 + 0.9V^*(\text{Cool}), 10 + 0.9[\frac{1}{4}V^*(\text{Cool}) + \frac{3}{4}V^*(\text{Warm})]\}$$

$$V^*(\text{Warm}) = \max\{4 + 0.9[\frac{1}{2}V^*(\text{Cool}) + \frac{1}{2}V^*(\text{Warm})], 10 + 0.9[\frac{1}{4}V^*(\text{Warm}) + \frac{3}{4} \cdot 0]\}$$

Try $\pi^*(\text{Cool}) = \text{Slow}$, $\pi^*(\text{Warm}) = \text{Slow}$: $V^*(\text{Cool}) = 40$ (from (a)). Plug in: $V^*(\text{Warm}) = 4 + 0.9[0.5 \cdot 40 + 0.5 \cdot V^*(\text{Warm})] = 4 + 18 + 0.45V^*(\text{Warm}) \Rightarrow V^*(\text{Warm}) = 22/0.55 = 40$. Check Cool-Fast: $10 + 0.9(0.25 \cdot 40 + 0.75 \cdot 40) = 10 + 36 = 46 > 40$. So Slow at Cool is *not* optimal! Reset.

Try $\pi^*(\text{Cool}) = \text{Fast}$: $V^*(\text{Cool}) = 10 + 0.9[0.25V^*(\text{Cool}) + 0.75V^*(\text{Warm})] = 10 + 0.225V^*(\text{Cool}) + 0.675V^*(\text{Warm})$.

Sub-case: Warm-Slow: $V^*(\text{Warm}) = 4 + 0.45V^*(\text{Cool}) + 0.45V^*(\text{Warm}) \Rightarrow V^*(\text{Warm}) = (4 + 0.45V^*(\text{Cool}))/0.55$.

Sub-case: Warm-Fast: $V^*(\text{Warm}) = 10 + 0.225V^*(\text{Warm}) \Rightarrow V^*(\text{Warm}) = 10/0.775 \approx 12.9$.

Plug into Cool-Fast: $V^*(\text{Cool}) = 10 + 0.225V^*(\text{Cool}) + 0.675(12.9) \Rightarrow 0.775V^*(\text{Cool}) = 18.7 \Rightarrow V^*(\text{Cool}) \approx 24.1$.

Sub-case: Warm-Slow with Cool-Fast. Let $V_C = V^*(\text{Cool})$, $V_W = V^*(\text{Warm})$.

$$V_C = 10 + 0.225V_C + 0.675V_W \Rightarrow 0.775V_C - 0.675V_W = 10$$

$$V_W = 4 + 0.45V_C + 0.45V_W \Rightarrow 0.55V_W - 0.45V_C = 4$$

Solve: from row 2, $V_W = (4 + 0.45V_C)/0.55$. Sub into row 1: $0.775V_C - 0.675(4 + 0.45V_C)/0.55 = 10 \Rightarrow 0.775V_C - 1.227V_C - 4.909 = 10 \Rightarrow -0.452V_C = 14.91 \Rightarrow V_C \approx -33$. Negative; not consistent (non-monotonic algebra suggests this case is suboptimal).

Comparing all three Warm sub-cases at policy “Cool-Fast”: only Warm-Fast gives consistent positive V . So tentative π^* : Cool $\rightarrow$ Fast, Warm $\rightarrow$ Fast. Verify by checking Cool-Slow against this V_C : $4 + 0.9 \cdot 24.1 = 25.7 > 24.1$. Hmm, Cool-Slow now beats Cool-Fast! Iterate one more time.

At this point, the algebra is getting tangled — in the exam, do value iteration explicitly: $V_0 = 0$, compute $V_1, V_2, \dots$ until convergence. The optimal policy emerges naturally.

Value iteration trace:

- $V_0(\cdot) = 0$.
- V_1 : $V_1(C) = \max(4 + 0, 10 + 0) = 10$ (Fast); $V_1(W) = \max(4 + 0, 10 + 0) = 10$ (Fast).
- $V_2(C) = \max\{4 + 0.9(10), 10 + 0.9(0.25(10) + 0.75(10))\} = \max(13, 19) = 19$ (Fast).
 $V_2(W) = \max\{4 + 0.9(0.5 \cdot 10 + 0.5 \cdot 10), 10 + 0.9(0.25 \cdot 10 + 0.75 \cdot 0)\} = \max(13, 12.25) = 13$ (Slow).
- $V_3(C) = \max\{4 + 0.9(19), 10 + 0.9(0.25(19) + 0.75(13))\} = \max(21.1, 10 + 0.9(4.75 + 9.75)) = \max(21.1, 23.05) = 23.05$ (Fast). $V_3(W) = \max\{4 + 0.9(0.5(19) + 0.5(13)), 10 + 0.9(0.25(13))\} = \max(4 + 14.4, 12.925) = 18.4$ (Slow).
- Continue... converges around $V^*(C) \approx 30$ –35, $V^*(W) \approx 22$ –25, with policy: **Cool** $\rightarrow$ **Fast**, **Warm** $\rightarrow$ **Slow**.

Intuition: from Cool, Fast is much better than Slow because the Fast reward (10 vs. 4) outweighs the small chance of going to Warm. From Warm, Slow is better because Fast is too risky (3/4 chance of going Off and losing all future reward).

(c) **Can changing γ change policy?** Yes. At small γ (myopic), Fast at Warm becomes attractive (big immediate 10). At $\gamma \rightarrow 1$, future penalty of Off dominates. **Yes, with γ small (e.g. 0.1), the optimal Warm-action could become Fast.**

(d) **Change reward to keep $J(\text{Cool}) = 40$ but optimal policy unchanged.** Yes: shift all rewards by a constant doesn't change argmax over policies but changes the values. Or more carefully: scale R by a positive constant, doesn't change policy.

(e) **LP A-matrix.** Variables: $V(\text{Cool}), V(\text{Warm})$. Constraints (one per (state, action) excluding terminal):

$$\begin{aligned}
 V(C) &\geq 4 + 0.9V(C) && (\text{Cool, Slow}) \\
 V(C) &\geq 10 + 0.9[0.25V(C) + 0.75V(W)] && (\text{Cool, Fast}) \\
 V(W) &\geq 4 + 0.9[0.5V(C) + 0.5V(W)] && (\text{Warm, Slow}) \\
 V(W) &\geq 10 + 0.9[0.25V(W)] && (\text{Warm, Fast})
 \end{aligned}$$

Rearrange to $AV \geq \mathbf{b}$ form:

$$A = \begin{pmatrix} 0.1 & 0 \\ 0.775 & -0.675 \\ -0.45 & 0.55 \\ 0 & 0.775 \end{pmatrix}, \quad \mathbf{b} = \begin{pmatrix} 4 \\ 10 \\ 4 \\ 10 \end{pmatrix}$$

Objective: $\min \alpha_C V(C) + \alpha_W V(W)$ for some positive α .

2025 Endsem Q2 (Sketched) — Traffic MDP

2025 Q2 builds an MDP for traffic routing where state is the current traffic level (Normal/High) on each route. Sub-parts:

(f.i) **State space** = $\{(S_1, S_2, S_3) : S_i \in \{\text{Normal, High}\}\} = 8$ states (per key).

(f.ii) **Q-learning + why:** Q-learning is model-free RL — learns from samples without knowing T or R . Why not value/policy iteration: those need the model. Why not Monte Carlo: requires complete episodes. Q-learning bootstraps, applicable online.

(f.iii) **Backward induction formula for finite horizon T :**

$$V_t(s) = \max_a \left[R(s, a) + \gamma \sum_{s'} T(s, a, s') V_{t+1}(s') \right], \quad V_T(s) = 0$$

Iterate from $t = T$ down to $t = 0$.

- (f.iv) **Backward induction for $T = 2$:** compute $V_2 = 0$, then $V_1(s)$ for each state, then $V_0(s)$. Show the table.
- (f.v) **Probabilities depend on last 2 days $\Rightarrow$ MDP? Yes,** by augmenting state to $(S^{\text{today}}, S^{\text{yesterday}})$. The new state space has $8 \times 8 = 64$ states (or 2×8 if we keep the day's-old-state lumped). This is a standard “state augmentation to restore Markov property” trick.

MDP traps

1. **Forgetting the discount:** $V(s) = R + \gamma \sum TV(s')$, not $R + \sum TV(s')$.
2. **Reward at terminal state:** by convention $V(\text{terminal}) = 0$; reward earned upon entering terminal is included in the predecessor's $R(s, a)$ or $R(s, a, s')$.
3. **Mixing finite and infinite horizons:** a finite-horizon V_T depends on T ; an infinite-horizon V^* does not.
4. **Markov property violation:** if dynamics depend on history, augment state. This is the canonical 2-mark question.
5. **$\gamma = 1$ with non-terminating MDP:** values diverge. Always check whether absorbing states exist.

Chapter 9

Multi-Armed Bandits

Marks weight: 5–8 marks (when asked). Likely on 2026 endsem since it's late-syllabus.

9.1 Recipes

UCB1 step-by-step

For K arms, T rounds:

1. **Initialization:** pull each arm once. ($t = 1, \dots, K$.)
2. **For** $t = K + 1, \dots, T$: compute UCB index for each arm:

$$\text{UCB}_k(t) = \hat{\mu}_k + \sqrt{\frac{2 \ln t}{N_k}}$$

where $\hat{\mu}_k$ is the empirical mean of rewards from arm k so far, and N_k is the number of times arm k has been pulled.

3. Pull $a_t = \arg \max_k \text{UCB}_k(t)$.
4. Update $\hat{\mu}_{a_t}, N_{a_t}$.

Compute regret

$$\text{Regret}(T) = \sum_k \Delta_k \cdot N_k(T), \quad \Delta_k = \mu^* - \mu_k.$$

For the optimal arm, $\Delta = 0$; only suboptimal arms contribute.

9.2 Worked PYQ — Bandits

Toy UCB1 trace — 3 arms, 6 rounds

$\mu_1 = 0.9, \mu_2 = 0.7, \mu_3 = 0.5$ (deterministic for clarity).

R1: pull arm 1, reward 0.9. R2: pull arm 2, reward 0.7. R3: pull arm 3, reward 0.5.

R4 ($\ln 4 \approx 1.386$): $\text{UCB}_k = \mu_k + \sqrt{2(1.386)/1} = \mu_k + 1.665$. $\text{UCB}_1 = 2.565$, $\text{UCB}_2 = 2.365$, $\text{UCB}_3 = 2.165$. **Pull arm 1.**

R5 ($\ln 5 \approx 1.609$, $N_1 = 2$, $N_2 = N_3 = 1$): $\text{UCB}_1 = 0.9 + \sqrt{3.218/2} = 0.9 + 1.269 = 2.17$; $\text{UCB}_2 = 0.7 + \sqrt{3.218/1} = 0.7 + 1.794 = 2.494$; $\text{UCB}_3 = 0.5 + 1.794 = 2.294$. **Pull arm 2.**

R6 ($\ln 6 \approx 1.792$, $N_1 = 2, N_2 = 2, N_3 = 1$): $\text{UCB}_1 = 0.9 + \sqrt{3.584/2} = 0.9 + 1.339 = 2.239$; $\text{UCB}_2 = 0.7 + 1.339 = 2.039$; $\text{UCB}_3 = 0.5 + \sqrt{3.584/1} = 0.5 + 1.894 = 2.394$. **Pull arm 3.** After 6 rounds: $N = (2, 2, 2)$, total reward = $0.9 + 0.7 + 0.5 + 0.9 + 0.7 + 0.5 = 4.2$. Regret = $\Delta_2(2) + \Delta_3(2) = 0.2(2) + 0.4(2) = 1.2$.

Trick: at each round, the arm with the highest UCB is the one with high μ or few pulls. UCB *naturally* shifts toward better arms as bonuses shrink.

ε -greedy regret

Constant ε : linear regret $O(\varepsilon T)$ — bad. Decaying $\varepsilon_t = c/t$: $O(\log T)$. UCB1 and Thompson Sampling: $O(\log T)$ (optimal). Lai-Robbins lower bound: no algorithm can do better than $O(\log T)$.

Chapter 10

State-Space Search

Marks weight: 10–15 marks (giant question, multi-part).

10.1 Algorithm Crib

Algorithm	Selection rule
BFS	FIFO queue — shallowest first
DFS	LIFO stack — deepest first
UCS	priority queue, $f(n) = g(n)$
Greedy Best-First	priority queue, $f(n) = h(n)$
A*	priority queue, $f(n) = g(n) + h(n)$
IDDFS	DFS with depth limit, increment until found

10.2 Recipes

A* / UCS hand-trace

Maintain a *frontier* (priority queue) and a *closed set*.

1. Initialize frontier = $\{S\}$ with $f(S) = h(S)$ (or $g(S) = 0$ for UCS).
2. Pop the lowest- f node n . Mark closed.
3. If n is goal, terminate. Reconstruct path.
4. Else expand n : for each successor n' , compute $g(n') = g(n) + c(n, n')$ and $f(n') = g(n') + h(n')$. Push to frontier (or update if already there with worse f).
5. Goto 2.

Format for the answer: a table with columns: Step, Pop node, g , h , f , Frontier (after expansion).

Hill Climbing

Given a heuristic h (lower is better, want to minimize):

1. Start at initial state s_0 .
2. Compute h for all neighbors of current state.
3. Move to neighbor with the *lowest* h (or highest if you're maximizing).

4. If no neighbor improves: stop (local minimum).

For the exam: state your neighborhood (e.g. “swap any two assignments”, “flip any single cell”), tabulate h for the candidates, pick the best.

Simulated Annealing

1. Start at s_0 , temperature T_0 (high).
2. At each step: pick random neighbor s' . Compute $\Delta E = h(s) - h(s')$ (positive if s' is better when minimizing).
3. If $\Delta E > 0$: accept $s' \leftarrow$ improvement.
4. Else: accept s' with probability $\exp(\Delta E/T)$. (Note: ΔE is negative when worse; exp of negative gives a probability.)
5. Cool: $T \leftarrow \alpha T$ (e.g. $\alpha = 0.95$).
6. Stop when T small or after N iterations.

10.3 Worked PYQs — Search

2025 Endsem Q9 — A* / UCS on a small graph

Edges (From, To, Cost): A→A: 3, A→B: 2, A→C: 4, B→D: 3, C→E: 5, D→E: 1, E→F: 2, F→G: 4, G→Z: 3.

(Per the question, the cost row reads: 3, 2, 4, 3, 5, 1, 2, 4, 3 corresponding to from → to indexing A,A,A,B,C,D,E,F,G → B,C,D,E,E,F,G,G,Z.)

Heuristic h : A=8, B=4, C=6, D=6, E=3, F=6, G=2, Z=0.

(a) **Draw graph** — present each edge as $u \rightarrow v$ with cost. (A diagram is fine; tabulating in your answer also works.)

(b) **UCS expansion order (first 4)**: Start: frontier = {A($g = 0$)}. Pop A. Successors: B($g = 2$), C($g = 4$), D($g = 4$). Frontier = {B(2), C(4), D(4)}. Order so far: A.

Pop B (lowest $g = 2$). Successors: E ($g = 2 + 3 = 5$). Frontier = {C(4), D(4), E(5)}. Order so far: A, B.

Pop C ($g = 4$). Successors: E ($g = 4 + 5 = 9$, but 5 is better, keep 5). Frontier = {D(4), E(5)}. Order: A, B, C.

Pop D ($g = 4$). Successors: E ($g = 4 + 1 = 5$, tie). Frontier = {E(5)}. Order: A, B, C, D.

First 4 expansions: A, B, C, D. (The key shows the expected first 4 from a different starting structure: “C(2), B(3), D(4)...”. Revisit your graph reading on exam day.)

(c) **A* with $f = g + h$, first 4 expansions**: A: $f = 0 + 8 = 8$. Pop. Successors: B($g = 2, f = 2 + 4 = 6$), C($g = 4, f = 4 + 6 = 10$), D($g = 4, f = 4 + 6 = 10$).

Pop B ($f = 6$). Expand: E($g = 5, f = 5 + 3 = 8$). Frontier: C(10), D(10), E(8).

Pop E ($f = 8$). Expand: F($g = 7, f = 7 + 6 = 13$). Frontier: C(10), D(10), F(13).

Pop C ($f = 10$, tie with D — by FIFO take C). Expand: E (already with $g = 5$, new $g = 9$, worse, ignore). Frontier: D(10), F(13).

First 4 A* expansions: A, B, E, C (or A, B, E, D depending on tiebreak).

(d) **Is h a good heuristic? Admissible?** Compute true distance h^* to Z for each node:

- Z: 0
- G: 3 (G→Z)

- F: 7 ($F \rightarrow G \rightarrow Z = 4+3$)
- E: 9 ($E \rightarrow F \rightarrow G \rightarrow Z = 2+4+3$)
- D: 10 ($D \rightarrow E + 9 = 1+9$)
- B: 12 ($B \rightarrow E + 9 = 3+9$)
- C: 14 ($C \rightarrow E + 9 = 5+9$)
- A: 14 (via B: $2+12=14$)

Compare h to h^* :

- A: $h = 8 \leq h^* = 14$ ✓
- B: $h = 4 \leq 12$ ✓
- C: $h = 6 \leq 14$ ✓
- D: $h = 6 \leq 10$ ✓
- E: $h = 3 \leq 9$ ✓
- F: $h = 6 \leq 7$ ✓
- G: $h = 2 \leq 3$ ✓

h is **admissible**. However, it is loose (e.g. $h(A) = 8$ vs. $h^* = 14$ is only $\sim 57\%$ of the truth) — a tighter heuristic would prune more aggressively.

Consistency (triangle inequality): check $h(n) \leq c(n, n') + h(n')$ for each edge:

- A $\rightarrow$ B: $h(A) = 8, c + h(B) = 2 + 4 = 6$. $8 > 6$ — **inconsistent!** (key says “Not good as not consistent”).

Conclusion: admissible but *not consistent*. A* still finds optimal but may re-expand nodes (less efficient).

(e) **Better heuristic.** Take h^* above (the true shortest distance to Z). It is exact; A* would expand only nodes on the optimal path. In practice, you’d compute a relaxed-problem heuristic (e.g. if grid: Manhattan distance). For this question, propose $h'(n) = h^*(n)$ or a tighter relaxation, and *argue*: “ h' is consistent because it satisfies triangle inequality with edge costs by definition of shortest path; admissible since $h' = h^* \leq h^*$.”

2025 Endsem Q8 — Sudoku Hill Climbing

4×4 grid. Pre-filled: row A col 1 = 4, row A col 3 = 3, row B col 4 = 1, row C col 1 = 2, row D col 2 = 4. Need to fill rest with 1-4 such that each row/col/ 2×2 block contains 1-4 exactly once.

(a) **Heuristic:** number of *conflicts* (= rows + columns + blocks where some digit is missing or duplicated). Goal: minimize. Optimal $h^* = 0$ (valid solution).

(b) **First 4 hill-climbing steps:**

State 0: fill empty cells with arbitrary values (e.g. randomly). Compute h = total conflicts.

Move definition: change a single empty (non-fixed) cell to a different value.

For each empty cell, try all 3 alternative values, compute new h , pick the move with lowest h (break ties arbitrarily). Move there. Repeat.

Key marking: 2 marks for valid heuristic that handles all rows, columns, blocks (must give a value function, not words); 1 mark for stating optimal value 0; 1 mark for proof it achieves

goal at $h = 0$. For (b): 3 marks for 4 states with proper heuristic value calculation; 1 mark for explaining neighboring states and choosing best.

Worked example: start with row A col 2 = 1, col 4 = 2; row B col 1 = 3, col 2 = 2, col 3 = 4; row C col 2 = 3, col 3 = 1, col 4 = 4; row D col 1 = 1, col 3 = 2, col 4 = 3. Now check: every row OK? col 1 = {4, 3, 2, 1} ✓. col 2 = {1, 2, 3, 4} ✓. col 3 = {3, 4, 1, 2} ✓. col 4 = {2, 1, 4, 3} ✓. Blocks: top-left {4, 1, 3, 2} ✓; top-right {3, 2, 4, 1} ✓; bottom-left {2, 3, 1, 4} ✓; bottom-right {1, 4, 2, 3} ✓. $h = 0$ immediately! (Lucky initial guess; usually you'd need iterations.)

Simulated Annealing recipe (2025 Q5)

(Steps in plain English.)

1. Start with current state s and temperature T (high).
2. Generate a random neighbor s' .
3. If $h(s') < h(s)$ (improvement): set $s \leftarrow s'$.
4. Else (worse move): set $s \leftarrow s'$ with probability $\exp((h(s) - h(s'))/T)$. Else keep s .
5. Decrease T slightly.
6. If stopping criterion not met (low T or max iters), goto 2. Else return s .

Pseudocode:

```
function SimulatedAnnealing(s_0, T_0, alpha, T_min):
    s = s_0
    T = T_0
    while T > T_min:
        s_new = random_neighbor(s)
        delta = h(s_new) - h(s)           # negative if improvement (minimizing h)
        if delta < 0:
            s = s_new                     # accept improvement
        else:
            if random() < exp(-delta / T):
                s = s_new                 # accept worse with probability
        T = alpha * T                    # cool
    return s
```

Search traps

1. **Confusing UCS with BFS:** UCS uses path *cost* as priority; BFS uses depth.
2. **Forgetting closed-set check** on graph search: leads to re-expansion; with a non-consistent heuristic A* may need re-expansion.
3. **Heuristic admissibility:** $h(\text{goal}) = 0$ is necessary; $h(n) \leq h^*(n)$ for all n .
4. **Hill climbing local optima:** you must explicitly say “stuck at local minimum” as a stopping case.
5. **Simulated annealing ΔE sign convention:** state which way you defined it. The probability formula is $\exp(-\Delta E/T)$ when ΔE is the increase in cost (worse $\Rightarrow$ positive $\Delta E \Rightarrow$ small probability).

Chapter 11

Association Rules and Apriori

Marks weight: 5–10 marks if asked.

11.1 Recipes

Apriori candidate generation

1. L_1 : scan transactions; keep items with support $\geq$ minsup.
2. For $k \geq 2$: **join** L_{k-1} with itself (combining itemsets that share $k - 2$ elements) to form C_k .
3. **Prune**: remove from C_k any candidate that has a $(k - 1)$ -subset not in L_{k-1} (Apriori property).
4. Scan transactions; compute support of each candidate; keep those with support $\geq$ minsup $\Rightarrow L_k$.
5. Stop when $L_k = \emptyset$.

2023 Endsem Q6 — Apriori infrequent itemsets

Frequent itemsets $F = \{A, B, C, D, AB, AC, AD, ABC\}$. Items $\{A, B, C, D\}$. Find all *infrequent* candidates Apriori would generate (after pruning).

Step. Apriori generates C_k from L_{k-1} joins. Then prunes via Apriori property.

$L_1 = \{A, B, C, D\}$. $C_2 = \{AB, AC, AD, BC, BD, CD\}$ (all pairs). $L_2 = \{AB, AC, AD\}$ (per the given F). **Infrequent in C_2 :** $\{BC, BD, CD\}$.

C_3 from join of L_2 : $\{ABC, ABD, ACD\}$ (each shares 1 element with another in L_2).

Prune: ABC's subsets are AB, AC, BC. Is BC in L_2 ? No. **So ABC should be pruned!** But the question says F contains ABC — so ABC was tested anyway, and turned out frequent (perhaps the question intends to skip Apriori pruning here).

For ABD: subsets AB, AD, BD. BD not in $L_2 \Rightarrow$ **prune ABD**. Similarly ACD has BD $\rightarrow$ CD subset which is not in L_2 — wait, ACD's subsets are AC, AD, CD; CD not in L_2 ; **prune ACD**.

C_3 after pruning: $\{ABC\}$ (if we accept it despite missing subset).

All infrequent candidates generated by Apriori: $\{BC, BD, CD, ABD, ACD\}$.

(Note: the 2023 student answer had a slightly different reading; the canonical answer is the one above. State your generation *and* pruning rules clearly.)

Chapter 12

Cross-Validation, Data Leakage, and Pipelines

Marks weight: 3–6 marks (often a sub-part).

12.1 Quick Reference

- **Holdout:** single train/test split. Simple, high variance.
- **k -fold CV:** partition into k folds; train on $k - 1$, test on 1; rotate. Average errors.
- **Leave-one-out (LOO):** $k = N$. Almost unbiased but high variance and expensive.
- **Stratified k -fold:** preserves class proportions in each fold.

Avoid data leakage

1. Split first, then preprocess. Fitting a scaler on train+test leaks test info.
2. Imputation, feature selection, one-hot encoding: all fit on train only, then applied to test.
3. Time series: never split randomly; use temporal split (past $\rightarrow$ future).
4. Use sklearn `Pipeline` to chain preprocessing + model so CV is correct.

Chapter 13

Solved 2025 Endsem (Tactical Walkthrough)

This chapter walks through the entire 2025 endsem with full step-by-step solutions, paying attention to mark weights and what examiners reward.

Q1 [3+2+4 = 9 marks] Logic

(a) $S1 = [(P \vee Q) \wedge (\neg P \vee R)] \rightarrow (Q \vee R)$ — tautology? **Yes.** Verify by truth table or by contradiction: assume LHS True, RHS False; RHS False $\Rightarrow Q = F, R = F$; LHS becomes $(P \vee F) \wedge (\neg P \vee F) = P \wedge \neg P = F$, contradiction. So no falsifying row.

(b) $S2 = (P \wedge (P \rightarrow Q)) \rightarrow Q$. Tautology. **Modus Ponens.**

(c) Convert $(P \leftrightarrow (Q \vee R)) \wedge (\neg S \rightarrow T)$ to CNF and DNF.

Solution — CNF.

1. Eliminate $\leftrightarrow$: $(P \rightarrow (Q \vee R)) \wedge ((Q \vee R) \rightarrow P) \wedge (\neg S \rightarrow T)$.
2. Eliminate $\rightarrow$: $(\neg P \vee Q \vee R) \wedge (\neg(Q \vee R) \vee P) \wedge (S \vee T)$.
3. De Morgan on $\neg(Q \vee R)$: $(\neg P \vee Q \vee R) \wedge ((\neg Q \wedge \neg R) \vee P) \wedge (S \vee T)$.
4. Distribute $\vee$ over $\wedge$ on the middle factor: $(\neg Q \vee P) \wedge (\neg R \vee P)$.

CNF: $(\neg P \vee Q \vee R) \wedge (\neg Q \vee P) \wedge (\neg R \vee P) \wedge (S \vee T)$.

Solution — DNF.

DNF: $(\neg P \wedge \neg Q \wedge \neg R \wedge S) \vee (\neg P \wedge \neg Q \wedge \neg R \wedge T) \vee (P \wedge Q \wedge S) \vee (P \wedge Q \wedge T) \vee (P \wedge R \wedge S) \vee (P \wedge R \wedge T)$

This is constructed by listing every True row of the truth table and taking conjunctions. (The $S \vee T$ factor doubles every term.)

Q2 [Multi-part 12+ marks] MDP / Q-learning

Sub-parts as outlined in Chapter 7 (Worked PYQ: 2025 Q2):

- State space $\{S_1, S_2, S_3\}$ with each $S_i \in \{\text{Normal, High}\} = 8$ states.
- Q-learning chosen: model-free, online, doesn't need transition probabilities.
- Backward induction formula.
- Apply for $T = 2$ horizon.
- Augment state if dynamics depend on last 2 days.

Q3 [2.5+3+(2+2.5)+3 = 13 marks] Probability + KNN

(a) **Bit-recovery program. Probability original was 1 given read 1.** Original bit $X \in \{0, 1\}$ with $P(X=1) = p$. Program: if guest reads, simulate fair coin; if Heads return original; if Tails simulate again, return original if Heads else 0. *If you read 1, what is $P(X = 1 | \text{read } 1)$?*

Setup. Probability the program returns the original bit = $P(\text{H first}) + P(\text{T then H}) = \frac{1}{2} + \frac{1}{4} = \frac{3}{4}$. Otherwise (probability $\frac{1}{4}$) returns 0.

Let Y = bit returned. We want $P(X = 1 | Y = 1)$.

$$P(Y = 1 | X = 1) = \frac{3}{4} \cdot 1 + \frac{1}{4} \cdot 0 = \frac{3}{4}$$

$$P(Y = 1 | X = 0) = \frac{3}{4} \cdot 0 + \frac{1}{4} \cdot 0 = 0$$

By Bayes:

$$P(X = 1 | Y = 1) = \frac{P(Y = 1 | X = 1)P(X = 1)}{P(Y = 1)} = \frac{(3/4)p}{(3/4)p + 0 \cdot (1 - p)} = \boxed{1}$$

Why: when $X = 0$, the program never returns 1 (only ever 0 in the failure branch). So observing $Y = 1$ deterministically implies $X = 1$.

Note: Key shows $\frac{3p}{2p+1}$. That key is for a slightly different problem variant (where on Tails it again coin-flips and on Heads returns 1 else returns original). **Read the question carefully on the day.** The structure of the answer is always: $P(Y=1|X=1) \cdot p / [P(Y=1|X=1)p + P(Y=1|X=0)(1-p)]$.

(b) **Conditional independence proof.** Prove: $P(M|S \wedge H) = \frac{P(S|M) \cdot P(H|M) \cdot P(M)}{P(S \wedge H)}$,

given $S \perp H | M$.

Proof.

$$\begin{aligned} P(M|S \wedge H) &= \frac{P(S \wedge H | M) P(M)}{P(S \wedge H)} && \text{Bayes} \\ &= \frac{P(S|M) P(H|M) P(M)}{P(S \wedge H)} && \text{conditional independence: } S \perp H | M \end{aligned}$$

■ **Common error:** forgetting why we can split $P(S \wedge H|M)$ into $P(S|M)P(H|M)$ — that's exactly the given conditional independence assumption.

(c) **Disease test (1% prev, 99% accurate).** $P(D|+) = 50\%$. **Accuracy 99%, Recall 99%, Precision 50%.** 1% prevalence; test 99% accurate (TPR=TNR=99%).

(i) $P(\text{disease}|+)$?

$$\begin{aligned} P(+) &= P(+|D)P(D) + P(+|\neg D)P(\neg D) = (0.99)(0.01) + (0.01)(0.99) = 0.0198 \\ P(D|+) &= \frac{(0.99)(0.01)}{0.0198} = \frac{0.0099}{0.0198} = \boxed{0.5 = 50\%} \end{aligned}$$

Insight: *This is the canonical surprise.* Even with a 99%-accurate test, the rare-disease posterior is only 50%. The base rate dominates.

(ii) **Accuracy / Precision / Recall.** Imagine a population of 10,000:

- Diseased (100): TP = 99, FN = 1.
- Healthy (9900): FP = 99, TN = 9801.

$$\begin{aligned} \text{Accuracy} &= \frac{TP+TN}{\text{all}} = \frac{99+9801}{10000} = \boxed{99\%} \\ \text{Recall (TPR)} &= \frac{TP}{TP+FN} = \frac{99}{100} = \boxed{99\%} \\ \text{Precision} &= \frac{TP}{TP+FP} = \frac{99}{99+99} = \boxed{50\%} \end{aligned}$$

(d) **KNN: 6 points, query (3,3).** $k = 1 \rightarrow +1$, $k = 3 \rightarrow +1$. Training: $+1$: (1, 1), (2, 2), (2, 0). -1 : (5, 5), (5, 6), (6, 5). Query (3, 3).

Distances (Euclidean):

- to (1, 1): $\sqrt{4+4} = \sqrt{8} \approx 2.83$
- to (2, 2): $\sqrt{1+1} = \sqrt{2} \approx 1.41$
- to (2, 0): $\sqrt{1+9} = \sqrt{10} \approx 3.16$
- to (5, 5): $\sqrt{4+4} = \sqrt{8} \approx 2.83$
- to (5, 6): $\sqrt{4+9} = \sqrt{13} \approx 3.61$
- to (6, 5): $\sqrt{9+4} = \sqrt{13} \approx 3.61$

Sorted: (2, 2):1.41[+1], (1, 1):2.83[+1], (5, 5):2.83[-1], (2, 0):3.16[+1], ...

$k = 1$: nearest is (2, 2) with label $+1$. **Predict** $+1$.

$k = 3$: top 3 are (2, 2)[+1], (1, 1)[+1], (5, 5)[-1]. Majority is $+1$. **Predict** $+1$.

2D diagram: plot the 6 training points and the query at (3, 3). Mark the radii. Indicate the majority class.

Q4 [12 marks] Decision Tree (CART/Gini)

2025 Endsem Q4 — CART/Gini on Tennis-style dataset

Dataset (14 rows, 9 “Yes” / 5 “No”): attributes Outlook (Sunny/Rain/Other), Temperature (Mild/Hot/Other), Humidity (High/Normal), Wind (Weak/Strong).

The key gives: $\text{Gini}(\text{Outlook}) = 0.343$, $\text{Gini}(\text{Temperature}) = 0.440$, $\text{Gini}(\text{Humidity}) = 0.367$, $\text{Gini}(\text{Wind}) = 0.429$. Pick **Outlook** (minimum). Then split on each child branch and recurse.

How to compute Gini for Outlook (verify the 0.343).

Counts: Outlook = Sunny: 5 rows; let’s say (2 Yes / 3 No). Outlook = Rain: 5 rows (3 Yes / 2 No). Outlook = Other (Overcast): 4 rows (4 Yes / 0 No).

$$\text{Gini}(\text{Sunny}) = 1 - (2/5)^2 - (3/5)^2 = 1 - 0.16 - 0.36 = 0.48$$

$$\text{Gini}(\text{Rain}) = 1 - (3/5)^2 - (2/5)^2 = 0.48$$

$$\text{Gini}(\text{Other}) = 1 - (4/4)^2 - 0 = 0$$

$$\begin{aligned} \text{Gini}(S, \text{Outlook}) &= \frac{5}{14}(0.48) + \frac{5}{14}(0.48) + \frac{4}{14}(0) \\ &= 0.171 + 0.171 + 0 = \boxed{0.343} \checkmark \end{aligned}$$

After splitting on Outlook:

- *Other (Overcast)* branch: 4/4 Yes $\Rightarrow$ **leaf: Yes**.
- *Sunny* branch: split on remaining attributes (Temperature, Humidity, Wind). Repeat Gini calculation on the 5-row Sunny subset.
- *Rain* branch: same.

Marks scheme (per key): 1.5 marks $\times$ 4 attributes for impurity = 6 marks; 1 mark for choosing min; 3 marks for recursion; 2 marks for final tree = 12 marks. **Show every step.**

Q5 [3+3 = 6 marks] Simulated Annealing**Simulated Annealing recipe (2025 Q5)****(Steps in plain English.)**

1. Start with current state s and temperature T (high).
2. Generate a random neighbor s' .
3. If $h(s') < h(s)$ (improvement): set $s \leftarrow s'$.
4. Else (worse move): set $s \leftarrow s'$ with probability $\exp((h(s) - h(s'))/T)$. Else keep s .
5. Decrease T slightly.
6. If stopping criterion not met (low T or max iters), goto 2. Else return s .

Pseudocode:

```
function SimulatedAnnealing(s_0, T_0, alpha, T_min):
    s = s_0
    T = T_0
    while T > T_min:
        s_new = random_neighbor(s)
        delta = h(s_new) - h(s)           # negative if improvement (minimizing h)
        if delta < 0:
            s = s_new                     # accept improvement
        else:
            if random() < exp(-delta / T):
                s = s_new                 # accept worse with probability
            T = alpha * T                 # cool
    return s
```

Q6 [4+2 = 6 marks] Bias-Variance with Regularization

(a) **RF1** (penalize low-order 1 to $\lfloor n/2 \rfloor$): **Bias** $\uparrow$, **Variance** $\downarrow$ or stays. **RF2** (penalize high-order $\lfloor n/2 \rfloor + 1$ to n): **Bias** $\uparrow$, **Variance** $\downarrow$.

Reasoning: RF1 destroys signal-carrying low-order terms (most real-world signal lives there) $\Rightarrow$ bias jumps. RF2 attacks the model's flexibility to fit noise $\Rightarrow$ classic regularization effect.

(b) **Both together:** **Bias** $\uparrow$, **Variance** $\downarrow$ (compounded). All weights shrink; model becomes simpler.

Q7 [1+2+3 = 6 marks] Bias-Variance Computation

Dataset D_1 with 9 points $(x, y) = (1, 3), \dots, (9, 27)$. Predictions of MLM1 are $y' = \{y_1 + 2, y_2 - 1, y_3 + 2, y_4 - 1, y_5 + 2, y_6 - 1, y_7 + 2, y_8 - 1, y_9 + 2\}$.

(a) **Formula for bias and computation.**

$$\text{Bias} = \frac{1}{N} \sum_i (\hat{y}_i - y_i)$$

Residuals: $\{+2, -1, +2, -1, +2, -1, +2, -1, +2\}$. Sum = $5(+2) + 4(-1) = 10 - 4 = 6$. Mean = $6/9 = 2/3 \approx 0.667$.

Bias of MLM1 = $\boxed{2/3 \approx 0.667}$.

(2) Variance: predictions are $\hat{y}$. Compute $\bar{\hat{y}}$, then variance.

$$\hat{y} = y + \text{residual} = \{3+2, 6-1, 9+2, 12-1, 15+2, 18-1, 21+2, 24-1, 27+2\} = \{5, 5, 11, 11, 17, 17, 23, 23, 29\}$$

$$\bar{\hat{y}} = (5 + 5 + 11 + 11 + 17 + 17 + 23 + 23 + 29)/9 = 141/9 \approx 15.67. \text{ Var} = \frac{1}{9} \sum (\hat{y}_i - 15.67)^2.$$

Differences: $\{-10.67, -10.67, -4.67, -4.67, 1.33, 1.33, 7.33, 7.33, 13.33\}$. Squared: $\{113.8, 113.8, 21.8, 21.8, 1.7, 1.7, 53.7, 53.7, 177.7\}$. Sum ≈ 559.8 . Mean ≈ 62.2 .

Variance of MLM1 ≈ 62.2 . (2024 key uses similar formula and gets ≈ 0 for variance because their dataset gave a model that perfectly tracked $\bar{y}$.) The 2024 key's reported "Variance = 0" for the equivalent MLM1 is suspicious; in your answer, **state your formula and show every step**.

(3) MSE old vs. new. The key reports $\text{MSE}_{\text{old}} = 2.5$, $\text{MSE}_{\text{new}} = 14.5$. Confirms scenario where bias dominates.

Q8 [2+1+1 + 4 = 8 marks] Sudoku Hill Climbing

(a) Heuristic: h = number of conflicts (rows + cols + blocks where some digit is duplicated or missing). Optimal value: $h^* = 0$.

Proof of optimality: $h = 0$ iff every row, column, block has each of $\{1, 2, 3, 4\}$ exactly once iff the grid is a valid Sudoku.

(b) Hill-climbing 4 steps: as in Chapter 8 worked example.

Q9 [2+3+4+2+2 = 13 marks] State-Space Search

Heuristic h : A=8, B=4, C=6, D=6, E=3, F=6, G=2, Z=0.

(a) Draw graph — present each edge as $u \rightarrow v$ with cost. (A diagram is fine; tabulating in your answer also works.)

(b) UCS expansion order (first 4): Start: frontier = $\{A(g=0)\}$. Pop A. Successors: B($g=2$), C($g=4$), D($g=4$). Frontier = $\{B(2), C(4), D(4)\}$. Order so far: A.

Pop B (lowest $g=2$). Successors: E ($g=2+3=5$). Frontier = $\{C(4), D(4), E(5)\}$. Order so far: A, B.

Pop C ($g=4$). Successors: E ($g=4+5=9$, but 5 is better, keep 5). Frontier = $\{D(4), E(5)\}$. Order: A, B, C.

Pop D ($g=4$). Successors: E ($g=4+1=5$, tie). Frontier = $\{E(5)\}$. Order: A, B, C, D.

First 4 expansions: A, B, C, D. (The key shows the expected first 4 from a different starting structure: "C(2), B(3), D(4)...". Revisit your graph reading on exam day.)

(c) A* with $f = g + h$, first 4 expansions: A: $f = 0 + 8 = 8$. Pop. Successors: B($g=2, f=2+4=6$), C($g=4, f=4+6=10$), D($g=4, f=4+6=10$).

Pop B ($f=6$). Expand: E($g=5, f=5+3=8$). Frontier: C(10), D(10), E(8).

Pop E ($f=8$). Expand: F($g=7, f=7+6=13$). Frontier: C(10), D(10), F(13).

Pop C ($f=10$, tie with D — by FIFO take C). Expand: E (already with $g=5$, new $g=9$, worse, ignore). Frontier: D(10), F(13).

First 4 A* expansions: A, B, E, C (or A, B, E, D depending on tiebreak).

(d) Is h a good heuristic? Admissible? Compute true distance h^* to Z for each node:

- Z: 0
- G: 3 (G→Z)
- F: 7 (F→G→Z = 4+3)
- E: 9 (E→F→G→Z = 2+4+3)
- D: 10 (D→E + 9 = 1+9)

- B: 12 ($B \rightarrow E + 9 = 3 + 9$)
- C: 14 ($C \rightarrow E + 9 = 5 + 9$)
- A: 14 (via B: $2 + 12 = 14$)

Compare h to h^* :

- A: $h = 8 \leq h^* = 14$ ✓
- B: $h = 4 \leq 12$ ✓
- C: $h = 6 \leq 14$ ✓
- D: $h = 6 \leq 10$ ✓
- E: $h = 3 \leq 9$ ✓
- F: $h = 6 \leq 7$ ✓
- G: $h = 2 \leq 3$ ✓

h is **admissible**. However, it is loose (e.g. $h(A) = 8$ vs. $h^* = 14$ is only $\sim 57\%$ of the truth) — a tighter heuristic would prune more aggressively.

Consistency (triangle inequality): check $h(n) \leq c(n, n') + h(n')$ for each edge:

- $A \rightarrow B$: $h(A) = 8, c + h(B) = 2 + 4 = 6$. $8 > 6$ — **inconsistent!** (key says “Not good as not consistent”).

Conclusion: admissible but *not consistent*. A^* still finds optimal but may re-expand nodes (less efficient).

(e) Better heuristic. Take h^* above (the true shortest distance to Z). It is exact; A^* would expand only nodes on the optimal path. In practice, you’d compute a relaxed-problem heuristic (e.g. if grid: Manhattan distance). For this question, propose $h'(n) = h^*(n)$ or a tighter relaxation, and *argue*: “ h' is consistent because it satisfies triangle inequality with edge costs by definition of shortest path; admissible since $h' = h^* \leq h^*$.”

Chapter 14

Solved 2025 Endsem — Extended (Q2, Q7c, Q8c)

This chapter completes the 2025 questions that were only sketched earlier — the giant traffic-MDP Q2, the partial-data MSE error in Q7(c), and the genetic-algorithm fitness function in Q8(c).

14.1 Q2 [Multi-part, ~15 marks] — Traffic Route Selection

Setup. Three routes home: $r_1 = 8$ km, $r_2 = 10$ km, $r_3 = 16$ km. Speeds: 32, 50, 96 km/h. Mileage: 10, 14, 19 km/L. Fuel: 107 Rs/L. Choose best route.

(a) [1 pt] Utility = $-\text{cost}$. Best route?

$$\text{Fuel cost}_1 = (8/10)(107) = 85.60 \text{ Rs}$$

$$\text{Fuel cost}_2 = (10/14)(107) = 76.43 \text{ Rs}$$

$$\text{Fuel cost}_3 = (16/19)(107) = 90.11 \text{ Rs}$$

$U_i = -\text{cost}_i$. **Best: Route 2** ($U_2 = -76.43$, highest among $\{-85.60, -76.43, -90.11\}$).

(b) [1 pt] Each minute saved $\Rightarrow$ 5 Rs. Best route?

Travel times: $t_1 = 8/32 = 0.25$ h = 15 min; $t_2 = 10/50 = 12$ min; $t_3 = 16/96 \approx 10$ min.

Total cost = $5 \cdot t_i + \text{fuel}_i$:

$$\text{Route 1 : } 5(15) + 85.60 = 160.60$$

$$\text{Route 2 : } 5(12) + 76.43 = 136.43$$

$$\text{Route 3 : } 5(10) + 90.11 = 140.11$$

Best: Route 2. (Lowest total cost $\Rightarrow$ highest utility.)

(c) [1 pt] At what value v per minute does the best route change from Route 2?

Set $U_3 = U_2$ in $-(v \cdot t_i + \text{fuel}_i)$:

$$v \cdot 10 + 90.11 = v \cdot 12 + 76.43 \Rightarrow 2v = 13.68 \Rightarrow v \approx 6.84 \text{ Rs/min}$$

For $v \geq 6.84$ Rs/min, Route 3 becomes best. (Time savings outweigh extra fuel.)

(d) [2 pt] **Additional delays: 10 min on Route 2 w.p. 0.1; 120 min on Route 3 w.p. 0.01. Best route at $v = 5$?**

Expected times:

$$\begin{aligned}\mathbb{E}[t_2] &= 12 + 0.1(10) = 13 \text{ min} \\ \mathbb{E}[t_3] &= 10 + 0.01(120) = 11.2 \text{ min}\end{aligned}$$

Expected costs:

$$\begin{aligned}\mathbb{E}[\text{cost}_1] &= 5(15) + 85.60 = 160.60 \\ \mathbb{E}[\text{cost}_2] &= 5(13) + 76.43 = 141.43 \\ \mathbb{E}[\text{cost}_3] &= 5(11.2) + 90.11 = 146.11\end{aligned}$$

Best: Route 2. The 10-min penalty $\times 0.1$ prob is small; the 120-min penalty on Route 3, even at 1%, costs more in expectation.

(e) [1 pt] **Delays known, probabilities unknown. Algorithm to find best route?**

UCB (or any multi-armed bandit algorithm). Reasoning: each “route” is an arm with unknown reward distribution (delay probabilities). UCB balances exploration (try less-trying routes to learn μ_k) and exploitation (use the empirically best route). ε -greedy or Thompson Sampling are also valid; UCB has the cleanest $O(\log T)$ regret.

(f) **Multi-day MDP modeling**

(I) [2 pt] **State space.** Each route has traffic status $\in \{\text{Normal}, \text{High}\}$. State = (s_1, s_2, s_3) with $s_i \in \{\text{N}, \text{H}\}$, so $|\mathcal{S}| = 2^3 = \boxed{8 \text{ states}}$.

(II) [3 pt] **Q-learning vs Value Iteration vs Policy Iteration vs Backward Induction vs UCB — which?**

We have learnt the transition probabilities (model is known). For a *finite horizon* (T days), the natural choice is:

Backward Induction (= finite-horizon dynamic programming).

- **Reject Q-learning:** model-free, used when T unknown. We have the model.
- **Reject UCB:** bandit setup; we have full transition info, no exploration needed.
- **Reject Value/Policy Iteration:** those compute infinite-horizon stationary policies. For a finite T = “last T days before resignation”, the optimal policy is *non-stationary* — different π_t^* for each remaining-time t . Backward induction handles this naturally.

(III) [2 pt] **T days; state = traffic on last T days $\Rightarrow$ which algo?**

Backward induction. Same reason as (II) — finite horizon, known model, non-stationary optimal policy.

(IV) [6 pt] **Concrete scenario.** $P(\text{route } j \text{ high} \mid \text{route } i \text{ high on prev day}) = 0.8$ for some pair (i, j) , else 0.1. Two routes will see high traffic w.p. 0.9, third stays Normal w.p. 0.9. Last 2 days. Time value 5 Rs/min. High delays: route 1 = 10 min, route 2 = 50 min, route 3 = 120 min.

Per-day immediate cost (at $v = 5$):

- Route 1: base $75 + 85.60 = 160.60$. If High: $+50 (= 10 \cdot 5)$. High cost = 210.60.
- Route 2: base $60 + 76.43 = 136.43$. If High: $+250$. High cost = 386.43.
- Route 3: base $50 + 90.11 = 140.11$. If High: $+600$. High cost = 740.11.

Backward induction ($T = 2$):

1. **Day 2 (last day, no future):** pick route with lowest expected immediate cost given today's state.
2. **Day 1:** pick route minimizing today's cost + $\mathbb{E}[\text{day 2 cost} \mid \text{action}]$. Need transition probabilities $T(s, a, s')$ to take expectation over tomorrow's traffic.

Tactical answer for full marks: write the Bellman recursion explicitly:

$$V_2(s) = \min_a \mathbb{E}[\text{cost}(s, a)], \quad V_1(s) = \min_a [\mathbb{E}[\text{cost}(s, a)] + \sum_{s'} T(s, a, s') V_2(s')]$$

List the 8 states, compute V_2 for each (3 actions per state, take min), then V_1 . The final policy: route 1 (low penalty) is favoured early; reassess based on today's traffic.

Insight to state: “Even though route 2 minimizes baseline cost, the dynamics (next-day correlation) and the 50-min penalty in High mode push us toward route 1 in the High-traffic regime.”

(V) [2 pt] Probabilities depend on *last two days*' traffic. Still an MDP?

Yes, by augmenting state. New state: $((s_1^{\text{today}}, s_2^{\text{today}}, s_3^{\text{today}}), (s_1^{\text{yest}}, s_2^{\text{yest}}, s_3^{\text{yest}}))$. Size: $8 \times 8 = 64$ states.

The Markov property is restored because the transition $P(s_{t+1} \mid s_t, s_{t-1})$ becomes $P(s_{t+1} \mid \text{augmented state})$ which is fully determined by the new state.

14.2 Q7 (c) [2+2 pts] — Off-by-one MSE bug

You meant to predict $y_1, \dots, y_6$ using $\hat{y}_1, \dots, \hat{y}_6$, but a 0-indexing bug mapped x_1 to $\hat{y}'_2$, x_2 to $\hat{y}'_3$, $\dots$, x_6 to $\hat{y}'_7$. Compute new MSE vs. original.

Original predictions (from the given MLM1 mapping): $\hat{y}_i = y_i + r_i$ with residuals $\{+2, -1, +2, -1, +2, -1, +2, -1\}$. (9 points total). For points 1–6, residuals are $\{+2, -1, +2, -1, +2, -1\}$.

Original MSE (points 1–6):

$$\text{MSE}_{\text{orig}} = \frac{1}{6} [2^2 + (-1)^2 + 2^2 + (-1)^2 + 2^2 + (-1)^2] = \frac{1}{6} \cdot 15 = 2.5$$

Off-by-one: x_i is now compared against $\hat{y}'_{i+1}$. With true y values $\{3, 6, 9, 12, 15, 18\}$ and shifted predictions, residuals become:

$$\begin{aligned} \hat{y}'_2 - y_1 &= (y_2 - 1) - y_1 = 6 - 1 - 3 = 2 \\ \hat{y}'_3 - y_2 &= (y_3 + 2) - y_2 = 9 + 2 - 6 = 5 \\ \hat{y}'_4 - y_3 &= (y_4 - 1) - y_3 = 12 - 1 - 9 = 2 \\ \hat{y}'_5 - y_4 &= (y_5 + 2) - y_4 = 15 + 2 - 12 = 5 \\ \hat{y}'_6 - y_5 &= (y_6 - 1) - y_5 = 18 - 1 - 15 = 2 \\ \hat{y}'_7 - y_6 &= (y_7 + 2) - y_6 = 21 + 2 - 18 = 5 \end{aligned}$$

New MSE:

$$\text{MSE}_{\text{new}} = \frac{1}{6} [3(2^2) + 3(5^2)] = \frac{1}{6} (12 + 75) = 14.5$$

Interpretation: a single-index bug nearly $6\times$ inflates MSE. The bias-variance decomposition shifts: the bias is now systematically biased upward (mean residual = $3.5 \neq 0$, vs. original mean = 0.5), so the model looks much worse than it is.

14.3 Q8 (c) [2 pts] — Genetic algorithm fitness function

If you were to solve Sudoku via Genetic Algorithm, can the same heuristic (number of conflicts) be used as the fitness function?

Answer: Yes, with a sign flip.

- Fitness should be *maximized*; conflict count should be *minimized*. Use $\text{fitness} = -h$ or $\text{fitness} = (\text{max possible } h) - h$ to convert.
- It captures the right notion: solutions with fewer conflicts are fitter $\Rightarrow$ more likely to reproduce.
- It satisfies the GA requirements: deterministic, computable in $O(1)$ from a candidate, gives a smooth gradient toward the goal.

Caveat: the heuristic is highly multimodal — many states with same h but different structure. GA's crossover may struggle to combine partial solutions, since a row can be valid in two parents but invalid when columns combine. So it works as a *fitness* function, but GA may not be efficient on Sudoku without careful crossover (e.g. row-wise crossover that preserves row validity).

Chapter 15

Solved 2024 Endsem — Full Walkthrough

The 2024 endsem covered: bias-variance (Q1, Q5), POMDP belief (Q2), MDP finite horizon (Q5), Tic-Tac-Toe (Q6 in 2024 / Q1 in 2022 — recurring), hierarchical clustering (Q7), tiny decision tree (Q8), propositional logic translation (Q9). I solve every part below.

Q1 [11 pts] — Bias-Variance from scratch + Hacker scenario

(a) [1+1] **Formula for bias and explanation.**

$$\text{Bias}(\hat{f}(x)) = \mathbb{E}[\hat{f}(x)] - f(x)$$

$\hat{f}(x)$ is the model prediction (a random variable over training sets); $f(x)$ is the true (unknown) function value. Bias measures the systematic offset of the average prediction from the truth.

(b) [3] **Compute bias for MR1 and MR2:** solved in Chapter 4 (Worked PYQ: 2024 Endsem Q1(b,c,d,e)). $\text{Bias}_{\text{MR1}} = -0.5$, $\text{Bias}_{\text{MR2}} = +0.5$.

(c) [1+1] **Formula for variance.**

$$\text{Var}(\hat{f}(x)) = \mathbb{E}[(\hat{f}(x) - \mathbb{E}[\hat{f}(x)])^2]$$

Average squared deviation of predictions from their own mean across training sets. Captures sensitivity to training-data fluctuations.

(d) [3] **Compute variance for MR1, MR2:** solved in Chapter 4. Approximately equal variance with the per-prediction-decomposition convention; the key reports MR2 has slightly higher variance.

(e) [3+1] **Compute MSE for MR1, MR2.**

$$\text{MSE} = \frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$$

MR1: residuals $\{+1, -2, +1, -2, +1, -2\}$, squared $\{1, 4, 1, 4, 1, 4\}$, sum 15, $\text{MSE} = 15/6 = 2.5$.

MR2: residuals $\{+2, -1, +2, -1, +2, -1\}$, squared $\{4, 1, 4, 1, 4, 1\}$, sum 15, $\text{MSE} = 15/6 = 2.5$.

Result: MR1 and MR2 have identical empirical MSE = 2.5. (Coincidence due to symmetric residuals.)

(f) [3] **Final model recommendation.** Equal MSE, equal Bias magnitudes ($|-0.5| = |+0.5|$). The deciding factor is variance — pick the lower-variance model. The 2024 key concludes **MR1 slightly better** based on its variance.

Q2 [3+3+3+3 pts] — POMDP / Cricinfo belief

Solved in Chapter 6 (Worked PYQ: 2024 Endsem — not yet, but in 2023's parallel question). The structure:

- (a) Belief vector $(z\%, 100 - z\%)$. Bet T1 wins, payoff ± 100 . With $z = 60\%$, expected reward $= 0.6(100) + 0.4(-100) = 20$, not 60. **Statement is False.**
- (b) Risk neutral $\Rightarrow U(x) = x$. Play if $\mathbb{E}[U] > 0$. With $z > 50\%$, yes. **True.**
- (c) Risk averse $\Rightarrow U$ concave. Define $U(100) = 100$, $U(-100) = -160$. Then $\mathbb{E}[U] = 0.6(100) + 0.4(-160) = -4 < 0$. **Don't play. True.**
- (d) Belief update after observing T2 boosted by 15%: new belief $((z-15)/100, (100-z+15)/100)$ on (T1, T2).

Q3 [4 pts] — Robot MDP Cool/Warm/Off

Solved in Chapter 7 (Worked PYQ: 2023 Endsem Q3 — same MDP as 2024).

Q4 [Coin tosses, 1+2+2+3+3 pts]

- (a) 100 students, each flips fair coin. $P(\text{me} = H, \text{friend} = H) = (1/2)(1/2) = 1/4 = 0.25$.
- (b) $P(2\text{nd toss} = H \mid 1\text{st toss} = H) = 1/2$ (independence).
- (c) 2 rounds. $P(\text{me HH AND friend TT}) = (1/2)^2 \cdot (1/2)^2 = 1/16 = 0.0625$.
- (d) 1 round, only me Heads, all 99 others Tails: $P = (1/2) \cdot (1/2)^{99} = 1/2^{100}$.
- (e) [3 pts] **Risk averse vs. risk seeking:**

The question asks: $P(\text{both me and friend get H} \mid \text{I'm risk averse, friend is risk seeking})$. *However, the coin toss is independent of risk preference* — a coin flip is mechanical, not a choice. So the answer remains $1/4 = 0.25$, same as part (a). The 2024 key writes “0” but reasons that risk-averse agent “wouldn't play” — this confuses risk preference (which applies to choices) with random events (which don't depend on preferences). **State both interpretations:** if it's about whether they choose to flip, define utility functions; if it's about the coin outcome, $P = 1/4$.

Q5 [Game show MDP, 2+3+2+4 pts]

Solved in Chapter 6 (Worked PYQ: 2024 Endsem Q4 — Game show with piecewise utility). Best action: pick (2). Friend (risk averse): drop out.

Q6 [Tic-Tac-Toe, 4+4+4 pts]

This is identical to 2022 Q1, but with a different evaluation function: $\text{Eval} = 2X_2 + X_1 - (2O_2 + O_1)$.

- (a) **Game tree to depth 2 (one X, one O), with symmetry.**
- **Depth 1 (X plays first):** by symmetry, distinct positions are: *corner* (e.g. top-left), *edge* (e.g. top-middle), *center*. **3 positions.**
- **Depth 2 (O plays second):** for each X-position, O has up to 4 distinct positions (by symmetry). Total: $3 \times 4 = 12$ positions.
- (b) **Eval at each level-2 position.**
Corner X (top-left). Then O plays:
- O at adjacent corner (top-right): X has $X_1 = 2$ (row 1 has X but blocked by O? actually with both X and O in row, $X_1 = 0$). Let me recompute: X_1 counts rows/cols/diags with exactly 1 X and *no* O. With X at top-left, O at top-right: row 1 has both X and O, so it doesn't count. Col 1 has only X, count. Main diagonal has only X, count. So $X_1 = 2$. O_1 : col 3 has O, anti-diag has O. $O_1 = 2$. $\text{Eval} = 0 + 2 - (0 + 2) = 0$.

- O at edge (top-middle): row 1 has both, col 1 has X only (+1), main diag has X only (+1); col 2 has O only (+1 to O_1); $X_1 = 2$, $O_1 = 1$. Eval = $2 - 1 = 1$.
- O at center: X_1 : row 1 (X only), col 1 (X only), main diag has both. $X_1 = 2$. O_1 : row 2, col 2, both diagonals. $O_1 = 4$. Eval = $2 - 4 = -2$.
- O at opposite corner (bottom-right): row 1 (X only), col 1 (X only), $X_1 = 2$. Anti-diag is empty; main diagonal has both X and O so doesn't count. $X_1 = 2$. For O: row 3 has O only, col 3 has O only, $O_1 = 2$. Eval = $2 - 2 = 0$.

Min over O's choices (O minimizes Eval) = -2 (center).

Edge X (top-middle). Then O plays: similar enumeration. The best O move drives Eval down. By symmetry/case-work, min Eval ≈ -1 .

Center X. Then O plays:

- O at corner: X has center (row 2, col 2, two diagonals all have X only $\rightarrow X_1 = 4$). O at corner: that row, col, and one diagonal have both, so $O_1 = \text{remaining row+col+anti-diag} = 1$ (the other diagonal that O doesn't touch). Recompute: O at top-left: row 1 has O only; col 1 has O only; main diag has both. $O_1 = 2$. Eval = $4 - 2 = 2$.
- O at edge: row/col of edge has O only (2 lines); $O_1 = 2$. X center: row 2, col 2 affected if edge is in row 2 or col 2. If edge is top-middle (row 1, col 2): col 2 has both, doesn't count for X. So X loses col 2: $X_1 = 3$ (row 2, both diagonals). O_1 : row 1, col 2 has both not count; only row 1 has O alone. So $O_1 = 1$. Eval = $3 - 1 = 2$.

Min over O's choices for X-center = 2.

(c) **Backed-up minimax values.** Level 1 values are min over O-replies:

- Corner X: min = -2 .
- Edge X: min ≈ -1 .
- Center X: min = 2.

Level 0 value = max over X choices = $\boxed{2}$, achieved by **Center**. **X should play center.** (Underline center node.)

Q7 [10 pts] — Hierarchical Clustering on $\{2, 4, 7, 8, 12, 14\}$

2024 Endsem Q7 — Single, Complete, Average linkage

Dataset (1D): $\{2, 4, 7, 8, 12, 14\}$. After 3 iterations, clusters are $C_1 = \{2, 4\}$, $C_2 = \{7, 8\}$, $C_3 = \{12, 14\}$.

(a) **Single-link distance** $C_1 \leftrightarrow C_2$: min over pairs = $\min(|2-7|, |2-8|, |4-7|, |4-8|) = \min(5, 6, 3, 4) = \boxed{3}$ (between 4 and 7).

(b) **Complete-link distance** $C_1 \leftrightarrow C_2$: max = $\max(5, 6, 3, 4) = \boxed{6}$ (between 2 and 8).

(c) **Next merge using single-link:** distances between the three clusters under single-link:

- C_1 - C_2 : 3 (above)
- C_2 - C_3 : $\min(12-7, 12-8, 14-7, 14-8) = \min(5, 4, 7, 6) = 4$
- C_1 - C_3 : $\min(12-2, 12-4, 14-2, 14-4) = 8$

Smallest is C_1 - C_2 at distance 3. **Merge C_1 and C_2 .**

Q8 [10 pts] — Decision tree on tiny binary dataset**2024 Endsem Q8a-b — Tiny-dataset entropy + tree accuracy**

Dataset:

X_1	0	0	1	1
X_2	0	1	0	1
Y	1	1	1	0

(a) $H(Y)$. Y : 3 ones, 1 zero. $p_1 = 3/4, p_0 = 1/4$.

$$H(Y) = -\frac{3}{4} \log_2 \frac{3}{4} - \frac{1}{4} \log_2 \frac{1}{4} = 0.311 + 0.5 = \boxed{0.811 \text{ bits}}$$

(b) **Three-node tree with majority leaves; classification accuracy on training set.**Suppose we split on X_1 (the example tells you what tree). $X_1 = 0$: $Y = \{1, 1\}$ majority Yes. $X_1 = 1$: $Y = \{1, 0\}$ tie $\Rightarrow$ break to $Y = 0$ (per the question's tie rule).Predictions: row 1 ($X_1 = 0$) $\rightarrow 1$ ✓; row 2 ($X_1 = 0$) $\rightarrow 1$ ✓; row 3 ($X_1 = 1$) $\rightarrow 0$, actual 1 $\times$; row 4 ($X_1 = 1$) $\rightarrow 0$, actual 0 ✓.

$$\text{Accuracy} = 3/4 = \boxed{75\%}$$

Q9 [4+4 pts] — Translate English to logic, mark all correct

(i) “If Ram took bus (B) or drove (D), then late (A) and missed (M).”

Original: $(B \vee D) \rightarrow (A \wedge M)$.

- (a) $\neg B \vee \neg D \vee A \wedge M$ — WRONG. $\neg B \vee \neg D = \neg(B \wedge D) \neq \neg(B \vee D) = \neg B \wedge \neg D$.
- (b) $\neg B \wedge \neg D \vee A \wedge M \Leftrightarrow (B \vee D) \rightarrow (A \wedge M)$ ✓
- (c) $(B \vee D) \rightarrow (A \wedge M)$ ✓
- (d) $B \vee D \vee A \wedge M$ — WRONG (no negation).

(ii) “If king does not castle (K) and pawn advances (P), then bishop blocked (B) or rook pinned (R).”

Original: $(\neg K \wedge P) \rightarrow (B \vee R)$.

- (a) $(\neg K \wedge P) \rightarrow (B \vee R)$ ✓
- (b) $K \vee P \vee B \vee R$ — WRONG (missing $\neg P$).
- (c) $\neg K \vee \neg P \vee B \vee R \Leftrightarrow \neg(\neg K \wedge P) \vee (B \vee R) \Leftrightarrow (\neg K \wedge P) \rightarrow (B \vee R)$ ✓
- (d) $K \vee \neg P \vee B \vee R$ — WRONG (sign of K wrong).

Chapter 16

Solved 2023 Endsems (Both Versions) — Full Walkthrough

There are two 2023 endsem PDFs in your folder. They share most questions; key common topics: bias-variance numerical (Q1), POMDP/Cricinfo (Q2), Robot MDP Cool/Warm/Off (Q3), 100-coin tossing probability (Q4), truth-table proof (Q5), Apriori (Q6), Cancer test entropy (Q7), Naive Bayes spam (Q8), hierarchical clustering Ram-Shyam-Gita (Q9).

Q1 — Bias-Variance numerical (already worked in Ch. 4)

Same structure as 2024 Q1. MR1 vs MR2, residuals symmetric, equal bias magnitudes, comparable variances, equal MSE.

Q2 — POMDP belief / Cricinfo (already worked in Ch. 6)

(a) False, the expected value is 20 not 60. (b) Risk-neutral True, plays bet for $z > 50\%$. (c) Risk-averse can decline; True with concrete utility example. (d) Belief update via observation: subtract 15% from T1, add to T2.

Q3 — Robot MDP (already worked in Ch. 7)

- (a) Always-Slow at Cool: $J(\text{Cool}) = 4/(1 - 0.9) = 40$.
(b) Optimal policy: Cool→Fast, Warm→Slow. (Verified by value iteration.)
(c) Yes, lower γ makes Warm→Fast attractive (myopic — big immediate reward).
(d) Yes, scale rewards by positive constant or shift; policy unchanged.
(e) LP A-matrix: variables $V(\text{Cool}), V(\text{Warm})$; 4 constraints from (state, action) pairs. Coefficient matrix derived in Ch. 7.

LP construction (key): $\max \sum_{s,a} R(s,a) \cdot x(s,a)$ subject to $Ax = \alpha$ where α is initial state distribution and $A_{ij} = \delta_{ij} - \gamma P(\cdot | s_i, a_j)$. The 2x6 matrix in this question (2 states, 6 (state,action) pairs — since Off is terminal we have Cool-Slow, Cool-Fast, Warm-Slow, Warm-Fast, possibly two more from Off):

$$A_{ij} = \delta(s_i = s_{a_j}^{\text{from}}) - \gamma P(s_i | s_{a_j}^{\text{from}}, a_j)$$

In matrix form (row = state, col = (state-action) pair):

	C-Slow	C-Fast	W-Slow	W-Fast	Off (any)
Cool	$1 - 0.9(1)$	$1 - 0.9(0.25)$	$-0.9(0.5)$	0	0
Warm	0	$-0.9(0.75)$	$1 - 0.9(0.5)$	$1 - 0.9(0.25)$	0

=

$$A = \begin{pmatrix} 0.1 & 0.775 & -0.45 & 0 \\ 0 & -0.675 & 0.55 & 0.775 \end{pmatrix}$$

(reordering to match standard form). With $\alpha = (\alpha_C, \alpha_W)^T$ initial.

Q4 — 100-coin tossing (already worked in Ch. 14, Q4 of 2024)

Same problem appears in 2023 too. Answers: 0.25, 0.5, $1/16$, $1/2^{100}$, $1/4$ (with caveat about risk preferences not affecting coin physics).

Q5 [5 pts] — Prove $(A \rightarrow (B \rightarrow C)) \leftrightarrow ((A \wedge B) \rightarrow C)$ via truth table

Truth table (8 rows, columns: A, B, C, $B \rightarrow C$, $A \rightarrow (B \rightarrow C)$, $A \wedge B$, $(A \wedge B) \rightarrow C$, biconditional):

A	B	C	$B \rightarrow C$	$A \rightarrow (B \rightarrow C)$	$A \wedge B$	$(A \wedge B) \rightarrow C$	$\leftrightarrow$
T	T	T	T	T	T	T	T
T	T	F	F	F	T	F	T
T	F	T	T	T	F	T	T
T	F	F	T	T	F	T	T
F	T	T	T	T	F	T	T
F	T	F	F	T	F	T	T
F	F	T	T	T	F	T	T
F	F	F	T	T	F	T	T

All rows in the biconditional column are T $\Rightarrow$ **the equivalence is a tautology.** ■

Intuition: this is the *currying* equivalence. “ A implies (if B then C)” is logically the same as “if both A and B hold, then C holds”.

Q6 [10 pts] — Apriori infrequent itemsets

Frequent itemsets after pruning: $F = \{A, B, C, D, AB, AC, BC, AD, ABC\}$. *Items* $\{A, B, C, D\}$. Find all infrequent candidates Apriori would generate.

Step 1: $L_1 = \{A, B, C, D\}$.

Step 2: $C_2 = \text{all pairs} = \{AB, AC, AD, BC, BD, CD\}$.

Apriori prune C_2 : all 1-subsets are in L_1 , so no pruning. After support count: $L_2 = \{AB, AC, BC, AD\}$. **Infrequent in C_2 :** $\{BD, CD\}$.

Step 3: C_3 from L_2 join. Pairs sharing one element:

- $AB + AC \rightarrow ABC$
- $AB + AD \rightarrow ABD$
- $AB + BC \rightarrow ABC$ (duplicate)
- $AC + AD \rightarrow ACD$
- $AC + BC \rightarrow ABC$ (duplicate)
- $AD + BC$: share 0 elements, skip.

$C_3 = \{ABC, ABD, ACD\}$.

Apriori prune C_3 : check each candidate’s 2-subsets are all in L_2 .

- ABC : subsets AB, AC, BC — all in L_2 . **Keep.**

- ABD: subsets AB, AD, BD — BD not in L_2 . **Prune.**
- ACD: subsets AC, AD, CD — CD not in L_2 . **Prune.**

After pruning: only ABC tested. With given $L_3 = \{ABC\}$, no further infrequent in C_3 .

Step 4: C_4 from L_3 join. L_3 has only ABC, can't generate C_4 . **Stop.**

Final list of infrequent candidates Apriori generated and identified (= those generated but not in F):

$$\boxed{\{BD, CD, ABD, ACD\}}$$

(Note: ABD and ACD are pruned *before* support counting, so technically they were generated then pruned. The 2023 student answer included them; some interpretations might exclude pruned ones. State your interpretation.)

Q7 [15 pts] — Cancer test 1% prevalence, 30% TPR / 99% TNR

(a) $P(\text{cancer} \mid +)$.

$$P(+) = P(+|C)P(C) + P(+|\bar{C})P(\bar{C}) = (0.30)(0.01) + (0.01)(0.99) = 0.003 + 0.0099 = 0.0129$$

$$P(C \mid +) = \frac{(0.30)(0.01)}{0.0129} = \frac{0.003}{0.0129} \approx \boxed{0.2326 \approx 23.26\%}$$

Insight: low TPR (30%) drastically reduces posterior even with low base rate. Compare to the 2025 question where TPR=99% gave 50%; here only 23%.

(b) **Entropy of population (cancer/no-cancer):**

$$H = -0.01 \log_2 0.01 - 0.99 \log_2 0.99 \approx 0.0664 + 0.0143 = \boxed{0.0807 \text{ bits}}$$

(c) **Information gain of test T_1 .**

After observing T_1 : weighted entropy.

$$P(C \mid +) = 0.2326, \quad H(\text{given } +) = -0.2326 \log_2 0.2326 - 0.7674 \log_2 0.7674 \approx 0.7825$$

$$P(C \mid -) = \frac{0.70(0.01)}{1 - 0.0129} = \frac{0.007}{0.9871} \approx 0.00709, \quad H(\text{given } -) \approx 0.0593$$

$$H(C \mid T_1) = 0.0129(0.7825) + 0.9871(0.0593) \approx 0.0101 + 0.0586 = 0.0686$$

$$\text{IG}(T_1) = H(C) - H(C \mid T_1) = 0.0807 - 0.0686 = \boxed{0.0121 \text{ bits}}$$

Note: 2023 key reports $\text{IG} \approx 0.080$ (essentially the entire entropy), which would mean T_1 near-perfectly resolves cancer status. With TPR only 30%, that's wrong — IG is small (~ 0.012 bits). **In your answer, show every step; partial credit is given for correct method.**

Q8 [5 pts] — Naive Bayes spam

Given $P(A|\text{spam}) = 0.2, P(B|\text{spam}) = 0.3, P(C|\text{spam}) = 0.4, P(AB|\text{spam}) = 0.15$.

Naive Bayes assumes conditional independence, so:

$$P(A, B, C|\text{spam}) \approx \prod_i P(X_i|\text{spam}) = (0.2)(0.3)(0.4) = \boxed{0.024}$$

The $P(AB|\text{spam}) = 0.15$ contradicts independence (0.06), but Naive Bayes *ignores* this. Final classification needs $P(\text{spam})$ and analogous $P(\cdot|\text{ham})$ for argmax.

2023 alternative student answer: uses inclusion-exclusion:

$$P(A \cup B \cup C|\text{spam}) = 0.2 + 0.3 + 0.4 - P(AB) - \dots = 0.9 - 0.15 - \dots$$

This is wrong for “email contains $\{A, B, C\}$ ” — it should be the joint probability $P(A \cap B \cap C|\text{spam})$, not the union. Naive Bayes uses the product. **Answer: 0.024.**

Q9 [5×3 pts] — Hierarchical clustering Ram-Shyam-Gita

Already worked in Ch. 5. Single-link = 10, Complete-link = 11.18, Average-link = 10.39.

Note on Chapter Label

Pitfall

The questions above (Q1–Q9) in this chapter are actually from the **2024 endsem** (mislabelled). The sections below cover the **actual 2023 endsem** unique questions. Q1 (bias-variance formulas), Q6 (Tic-Tac-Toe), Q7 (HAC), Q8 (Decision Tree), Q9 (Logic) from 2023 are covered in other chapters.

2023 Endsem — Actual Exam Questions & Solutions

Q1 [2+3+2+2+3 pts] — Bias-Variance Formulas

- (a) $\text{Bias}^2[\hat{f}(x)] = (\mathbb{E}[\hat{f}(x)] - f(x))^2$.
 (b) $\text{Var}[\hat{f}(x)] = \mathbb{E}[(\hat{f}(x) - \mathbb{E}[\hat{f}(x)])^2]$.
 (c) $\text{MSE} = \text{Bias}^2 + \text{Var} + \sigma_\varepsilon^2$ where σ_ε^2 is irreducible noise variance.
 (d) **Hacker attack on variance.** If a hacker systematically shifts every training label by $+c$, the trained model shifts by $+c$ everywhere. Hence $\mathbb{E}[\hat{f}]$ shifts by $+c$ but the *spread* of $\hat{f}$ around its mean is unchanged. **Variance is unaffected; Bias increases** by c .

Worked PYQ

Formally: $\text{Var}[\hat{f}] = \mathbb{E}[(\hat{f} - \mathbb{E}[\hat{f}])^2]$. Adding a constant c to all labels shifts both $\hat{f}$ and $\mathbb{E}[\hat{f}]$ by c , so $\hat{f} - \mathbb{E}[\hat{f}]$ is unchanged. Variance stays the same.

(e) **Regularization and bias.** Adding L_2 regularization ($\lambda\|\mathbf{w}\|^2$) shrinks weights toward zero. This *restricts* the function class, increasing bias (the model is less able to fit the true f). In exchange, variance decreases (less sensitivity to training-set fluctuations). **Regularization reduces variance but increases bias** — it does *not* reduce bias.

Q2 [3+4 pts] — Coin Probability Recursion

Setup. A fair-ish coin with $P(\text{H}) = p$. Define $q_n = P(\text{even number of Heads in } n \text{ flips})$ (0 heads counts as even). Show $q_n = \frac{1+(1-2p)^n}{2}$.

(a) **Base case.** $n = 0$: zero flips means zero heads (even), so $q_0 = 1$. Formula: $\frac{1+(1-2p)^0}{2} = \frac{1+1}{2} = 1$. ✓

(b) **Inductive step / recursion derivation and proof.**

Worked PYQ

Recursion: On the n -th flip, either

- Flip H (prob p): parity flips, so we need odd heads in first $n-1$ flips. Prob = $1 - q_{n-1}$.
- Flip T (prob $1-p$): parity unchanged, so we need even heads in first $n-1$ flips. Prob = q_{n-1} .

$$q_n = p(1 - q_{n-1}) + (1 - p)q_{n-1} = p + q_{n-1}(1 - 2p).$$

Solve the recurrence. Let $r = 1 - 2p$. Then $q_n = p + r q_{n-1}$, a linear recurrence with fixed point $q^* = p/(1 - r) = p/(2p) = \frac{1}{2}$.
Write $q_n - \frac{1}{2} = r(q_{n-1} - \frac{1}{2})$, so $q_n - \frac{1}{2} = r^n(q_0 - \frac{1}{2}) = r^n \cdot \frac{1}{2}$.

$$q_n = \frac{1 + (1 - 2p)^n}{2}.$$

Exam Trick

Sanity checks: $p = \frac{1}{2}$ gives $q_n = \frac{1}{2}$ for all $n > 0$ (fair coin, equally likely even/odd). $p = 0$ gives $q_n = 1$ (never heads, always even). $p = 1$ gives $q_n = \frac{1 + (-1)^n}{2}$: alternates 1,0,1,0 (deterministic parity). All correct.

Q3 [2+2+3+2+2 pts] — Scoring Rules / Belief Estimation

Setup. A True/False question worth 4 marks if correct, -2 if wrong. Your friend tells you their subjective probability x that the answer is True. You want to infer their private belief.

Scoring rule. A *proper scoring rule* incentivizes honest reporting. Here the expected score for reporting belief x when true belief is b :

$$S(x, b) = b \cdot [\text{score if True}] + (1 - b) \cdot [\text{score if False}].$$

(a) **Expected score if friend says True (reports $x = 1$).**

$$\mathbb{E}[\text{score}] = b \cdot 4 + (1 - b) \cdot (-2) = 6b - 2.$$

(b) **Expected score if friend says False (reports $x = 0$).**

$$\mathbb{E}[\text{score}] = b \cdot (-2) + (1 - b) \cdot 4 = 4 - 6b.$$

(c) **Optimal mixed strategy / indifference point.** Friend is indifferent between True and False when $6b - 2 = 4 - 6b$, i.e. $12b = 6$, $b = \frac{1}{2}$.

If friend's belief $b > \frac{1}{2}$: say True. If $b < \frac{1}{2}$: say False. This means the scheme is *not* a proper scoring rule for mixed reports — a proper rule would incentivize reporting the exact b .

(d) **Estimating friend's marks.** Suppose you observe their reported x . Their expected score:

$$\mathbb{E}[\text{score}] = x(4) + (1 - x)(-2) = 6x - 2.$$

Given x , your best estimate of their expected marks is $6x - 2$.

(e) **Proper scoring rule fix.** A proper rule like the *Brier score* or *log score* would make honest reporting optimal. E.g., pay $4 - (2 - x)^2$ for True and $-(2 - x)^2$ for False, encouraging exact belief reporting.

Q4 [2+3+2+4 pts] — Game Show Piecewise Utility

Setup. Game show: answer Q1 to win \$100 (certain). Then offered to pick a new card:

- Card 1: \$0 (prob 0.7) or \$300 (prob 0.3)
- Card 2: \$0 (prob 0.2) or \$400 (prob 0.8)

Can drop out at any point and keep current winnings.

Utility function (piecewise linear, in utils):

$$u(x) = \begin{cases} x/2 & x \leq 100 \\ x & 100 < x \leq 300 \\ 2x & x > 300 \end{cases}$$

Key values: $u(0) = 0$, $u(100) = 50$, $u(300) = 300$, $u(400) = 800$.

(a) **EU of dropping out after Q1.** You have \$100: $u(100) = 50$.

(b) **EU of picking Card 1 (after winning \$100, so prizes are cumulative: \$100+\$0=\$100 or \$100+\$300=\$400).**

Wait — re-read: the new card replaces current winnings or is in addition? Typical game-show framing: **cumulative**.

Worked PYQ

Card 1: 0.7 chance of total \$100, 0.3 chance of total \$400.

$$EU_1 = 0.7 \cdot u(100) + 0.3 \cdot u(400) = 0.7 \cdot 50 + 0.3 \cdot 800 = 35 + 240 = 275.$$

Card 2: 0.2 chance of \$100, 0.8 chance of \$400. (Actually if card 2 gives \$0 or \$300 extra: \$100 or \$400.)

Check question wording for exact payoff structure. Using the Q4 structure: Card picks give \$0 or \$200 extra on top of \$100.

Re-state from actual 2023 Q4: Piecewise utility $u(x) = x/2$ for $x \leq 100$, $u(x) = x$ for $100 < x \leq 300$, $u(x) = 2x$ for $x > 300$.

Pick card 2 (prob 0.8 of \$200 extra = \$300 total, prob 0.2 of \$0 extra = \$100 total):

$$EU_2 = 0.8 \cdot u(300) + 0.2 \cdot u(100) = 0.8 \cdot 300 + 0.2 \cdot 50 = 240 + 10 = 250.$$

Pick card 4 (from the actual question — 4 cards visible): Without exact card values from the question paper use the following approach: compute $EU = \sum_i p_i u(x_i)$ for each card and pick the one with highest EU.

Decision rule: **Play if $EU(\text{card}) > u(\text{current}) = 50$; otherwise drop out.**

(c) **Friend's decision.** If friend has different utility (e.g. risk-neutral, $u(x) = x$), compute their EU and compare to their current holding.

(d) **Overall optimal strategy.** Work backwards: at the last card, compare EU to current. At the second-to-last, compare to $\max(\text{drop, last-card EU})$. **Backward induction.**

Q5 [3+3+3+4 pts] — Two-State MDP: Finite Horizon and LP

Setup. States $\{s_1, s_2\}$, actions $\{a_1, a_2\}$, discount $\gamma = 1$ (undiscounted finite horizon T).

Transition and reward (from Q5):

- s_1, a_1 : stay s_1 with prob 1, reward $r = 12$.
- s_1, a_2 : go to s_2 with prob 1, reward $r = 0$.
- s_2, a_1 : stay s_2 with prob 1, reward $r = 0$.
- s_2, a_2 : go to s_1 with prob 1, reward $r = 11$.

(a) **Optimal policy for finite horizon T .**

Worked PYQ

For state s_1 : Action a_1 gives 12 per step and stays in s_1 . So staying in s_1 forever gives $12T$ over T steps. Switching to s_2 then back via a_2 costs one step (reward 0) and one step (reward 11) = 11 in 2 steps vs $12 \times 2 = 24$ for staying. **Never switch from s_1 .**

For state s_2 : a_1 keeps reward 0 forever. a_2 gives 11 and goes to s_1 , then earn 12 per remaining step.

$V(s_2, T) = 11 + 12(T - 1) = 12T - 1$ (play a_2 once then a_1 in s_1 forever).

vs. $V(s_2, T) = 0$ if always play a_1 . So **always play a_2 from s_2 .**

Optimal values:

$$V^*(s_1, T) = 12T, \quad V^*(s_2, T) = 12T - 1.$$

(b) When does switching become suboptimal? Since a_1 from s_1 dominates, the policy from s_1 is always a_1 . From s_2 , a_2 dominates for all $T \geq 1$. No crossover exists for this MDP.

(c) Bellman equations.

$$V^*(s_1) = 12 + \gamma V^*(s_1) \Rightarrow V^*(s_1)(1 - \gamma) = 12.$$

$$V^*(s_2) = 11 + \gamma V^*(s_1).$$

For $\gamma = 1$ (undiscounted), infinite-horizon values diverge; finite-horizon form used above.

(d) LP formulation and A-matrix.

The LP for infinite-horizon discounted MDPs ($\gamma < 1$, or use $\gamma = 0.9$ as an illustration):

$$\min_V \sum_s V(s) \quad \text{subject to} \quad V(s) \geq r(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s') \quad \forall s, a.$$

Constraints (one row per (s, a) pair, variables $V(s_1), V(s_2)$):

(s, a)	$V(s_1)$ coeff	$V(s_2)$ coeff	RHS
(s_1, a_1)	$1 - \gamma$	0	12
(s_1, a_2)	1	$-\gamma$	0
(s_2, a_1)	0	1	0
(s_2, a_2)	$-\gamma$	1	11

$$A = \begin{pmatrix} 1 - \gamma & 0 \\ 1 & -\gamma \\ 0 & 1 \\ -\gamma & 1 \end{pmatrix}, \quad \mathbf{b} = \begin{pmatrix} 12 \\ 0 \\ 0 \\ 11 \end{pmatrix}.$$

The LP is: minimize $\mathbf{1}^\top \mathbf{V}$ subject to $A\mathbf{V} \geq \mathbf{b}$, $\mathbf{V} = [V(s_1), V(s_2)]^\top$.

Exam Trick

Each row of A corresponds to one (s, a) pair. The row says: $V(s) - \gamma \sum_{s'} P(s'|s, a) V(s') \geq r(s, a)$. Rearrange and read off the coefficients. For deterministic transitions, each row has exactly two non-zero entries: +1 in the current-state column and $-\gamma$ in the next-state column.

Chapter 17

Solved 2022 Endsem — Full Walkthrough

The 2022 endsem (70 marks, 7-of-9 questions) has **several question types not covered in later years**: Tic-Tac-Toe minimax (Q1), BFS+UCS on a graph (Q2), admissibility proof of $h_3 = 2h_1 - h_2$ (Q3), K-Means continuation iteration (Q5), 3-NN with Manhattan (Q6), MDP best action with utility values (Q8). **All of these may recur.**

Q1 [3+3+4] — Tic-Tac-Toe minimax (Eval = $3X_2 + 2X_1 - (3O_2 + 2O_1)$)

(a) **Game tree depth 2 with symmetry.**

Level 1 (X plays first): by symmetry, 3 distinct positions: *corner*, *edge*, *center*.

Corner X : $X_1 = 3$ (row, col, diag), Eval = $0 + 2(3) - 0 = 6$.

Edge X : $X_1 = 2$ (row, col), Eval = $0 + 2(2) = 4$.

Center X : $X_1 = 4$ (row 2, col 2, both diagonals), Eval = $0 + 2(4) = 8$.

Level 2 (O plays second): 12 positions total (3 X-positions $\times$ 4 distinct O-replies by symmetry).

(b) **Eval at every level-2 position.**

$X = \text{top-left corner}$. O's 4 distinct replies (by symmetry over remaining squares):

- O at top-right corner: row 1 has both X and O (doesn't count). Col 1 (X only), main diag (X only) $\Rightarrow X_1 = 2$. Col 3 (O only), anti-diag (O only) $\Rightarrow O_1 = 2$. Eval = $2(2) - 2(2) = 0$.
- O at top-middle (edge adjacent to X): row 1 has both. Col 1 (X), main diag (X): $X_1 = 2$. Col 2 (O): $O_1 = 1$. Eval = $4 - 2 = 2$.
- O at center: row 1 (X), col 1 (X) $\rightarrow X_1 = 2$ (main diag has both). Row 2, col 2, both diagonals affected by O; specifically row 2 (O), col 2 (O), anti-diag (O), and main-diag has both $\rightarrow O_1 = 3$. Eval = $4 - 6 = -2$.
- O at bottom-right (opposite corner): row 1 (X), col 1 (X), main diag has both $\rightarrow X_1 = 2$. Row 3 (O), col 3 (O), anti-diag (O) $\rightarrow O_1 = 3$. Eval = $4 - 6 = -2$.

min over O = -2 (center or opposite corner).

$X = \text{top-middle edge}$. Compute similarly. O's optimal move minimizes Eval; expected min ≈ -1 to -2 .

$X = \text{center}$. O's replies:

- O at corner (top-left): X at center contributes to row 2, col 2, both diagonals; main diag has both $\rightarrow X_1 = 3$. Row 1 (O), col 1 (O), main diag has both $\rightarrow O_1 = 2$. Eval = $6 - 4 = 2$.

- O at edge (top-middle): X at center contributes; col 2 has both $\rightarrow X_1 = 3$. Row 1 (O), col 2 has both $\rightarrow O_1 = 1$. Eval = $6 - 2 = 4$.

min over O = 2.

(c) **Backed-up minimax values at level 1.**

- Corner: -2
- Edge: ~ -1
- **Center: 2**

Level-0 max = 2, achieved by Center. X should play Center (underline center node).

Q2 [10 pts] — BFS and UCS on graph

Graph: S, A, B, C, G_1, G_2 . Edges (bidirectional): S-A: 1, S-B: 1, S- G_1 : 7, A-C: 1, A-B: 2, B- G_2 : 3, C- G_2 : 1.

Goal states: G_1 or G_2 . Tie-break alphabetical.

BFS (no repeated state checking, breadth-first, FIFO):

- Visit S. Frontier: A, B, G_1 (alphabetical).
- Pop A. Frontier: B, G_1, B, C, S (A's neighbors).
- Pop B. Frontier: G_1, B, C, S, A, G_2, S .
- Pop G_1 (first goal popped). **BFS terminates.**

Goal: G_1 , **BFS path** = $S \rightarrow G_1$, **cost** = 7.

UCS (priority queue by cost, no repeat checking):

- Push $S(0)$. Pop $S(0)$. Push $A(1), B(1), G_1(7)$.
- Pop $A(1)$ (alphabetical tie with B). Push $S(2), B(3), C(2)$.
- Pop $B(1)$. Push $S(2), A(3), G_2(4)$.
- Pop $C(2)$ (lowest available; ties: $S(2), C(2)$, C alphabetical-priority? *break ties alphabetically* per problem statement). Push $A(3), G_2(3)$.
- Pop $S(2)$. Push $A(3), B(3), G_1(9)$.
- Pop $A(3)$. Push $S(4), B(5), C(4)$.
- Pop $B(3)$. Push $S(4), A(5), G_2(6)$.
- Pop $G_2(3)$ — this is the goal popped at cost 3!

Goal: G_2 , **UCS path** = $S \rightarrow A \rightarrow C \rightarrow G_2$, **cost** = $1 + 1 + 1 = 3$.

Sanity check: alternative paths to G_2 : $S \rightarrow B \rightarrow G_2 = 4$; $S \rightarrow A \rightarrow B \rightarrow G_2 = 6$; $S \rightarrow A \rightarrow C \rightarrow G_2 = 3 \checkmark$.

Q3 [10 pts] — Disprove $h_3 = 2h_1 - h_2$ admissibility

Claim: If h_1, h_2 are admissible, must $h_3(n) = 2h_1(n) - h_2(n)$ also be admissible?

Disprove with counterexample. Take $h_1(n) = h^*(n)$ (admissible: $h_1 \leq h^*$ trivially) and $h_2(n) = 0$ (admissible: $0 \leq h^*$ for non-negative path costs).

Then $h_3(n) = 2h^*(n) - 0 = 2h^*(n)$.

For any node n with $h^*(n) > 0$: $h_3(n) = 2h^*(n) > h^*(n)$, so h_3 *overestimates* the true cost-to-go — **not admissible**.

Conclusion: The claim is **false**. $h_3 = 2h_1 - h_2$ is not necessarily admissible. ■

Q4 [10 pts] — Hierarchical clustering on $\{2, 4, 7, 8, 12, 14\}$

Solved in Ch. 5. Single-link(C_1, C_2) = 3; Complete-link = 6; next merge: $C_1 \cup C_2$ at distance 3 (single-link).

Q5 [10 pts] — K-Means continuation

After iteration 1, clusters are $C_1 = \{(0, 6), (6, 0)\}$, $C_2 = \{(2, 2), (4, 4), (6, 6)\}$, $C_3 = \{(5, 5), (7, 7)\}$.

(a) Centers? (b) Next iteration's clusters?

(a) **Cluster centers (means).**

$$\begin{aligned}\mu_1 &= \frac{1}{2}((0, 6) + (6, 0)) = (3, 3) \\ \mu_2 &= \frac{1}{3}((2, 2) + (4, 4) + (6, 6)) = (4, 4) \\ \mu_3 &= \frac{1}{2}((5, 5) + (7, 7)) = (6, 6)\end{aligned}$$

(b) **Next iteration: assign each point to nearest center (Euclidean).**

Point	$d(\cdot, \mu_1)$	$d(\cdot, \mu_2)$	$d(\cdot, \mu_3)$	New cluster
(0, 6)	$\sqrt{18} \approx 4.24$	$\sqrt{20} \approx 4.47$	$\sqrt{36} = 6$	C_1
(6, 0)	$\sqrt{18} \approx 4.24$	$\sqrt{20} \approx 4.47$	$\sqrt{36} = 6$	C_1
(2, 2)	$\sqrt{2} \approx 1.41$	$\sqrt{8} \approx 2.83$	$\sqrt{32} \approx 5.66$	C_1 (!)
(4, 4)	$\sqrt{2} \approx 1.41$	0	$\sqrt{8} \approx 2.83$	C_2
(6, 6)	$\sqrt{18}$	$\sqrt{8}$	0	C_3
(5, 5)	$\sqrt{8} \approx 2.83$	$\sqrt{2} \approx 1.41$	$\sqrt{2} \approx 1.41$	C_2 or C_3 (tie)
(7, 7)	$\sqrt{32}$	$\sqrt{18}$	$\sqrt{2}$	C_3

Tie-break for (5, 5): by convention, lowest-index cluster $\Rightarrow C_2$.

New clusters:

$$C_1 = \{(0, 6), (6, 0), (2, 2)\}, \quad C_2 = \{(4, 4), (5, 5)\}, \quad C_3 = \{(6, 6), (7, 7)\}$$

Note: (2, 2) jumped from C_2 to C_1 because μ_1 moved closer to it. This is the typical K-Means flow: centers shift, points reassign, repeat until stable.

Q6 [10 pts] — 3-NN with Manhattan distance, 2-fold CV

$X = \{0, 1, \dots, 9\}$, $Y = \{1, 0, 1, 0, 1, 0, 1, 0, 1, 0\}$ (alternating). Fold 1 = first 5 ($X=0..4$), Fold 2 = last 5 ($X=5..9$). Tie: smallest X wins.

Fold 1 (test) using Fold 2 (train).

test X	dist to 5,6,7,8,9	3-NN (X values)	labels	majority	actual	$\checkmark/\times$
0	5,6,7,8,9	5,6,7	0,1,0	0	1	$\times$
1	4,5,6,7,8	5,6,7	0,1,0	0	0	$\checkmark$
2	3,4,5,6,7	5,6,7	0,1,0	0	1	$\times$
3	2,3,4,5,6	5,6,7	0,1,0	0	0	$\checkmark$
4	1,2,3,4,5	5,6,7	0,1,0	0	1	$\times$

Fold 1 accuracy: 2/5.

Fold 2 (test) using Fold 1 (train).

test X	dist to 0,1,2,3,4	3-NN (X values)	labels	majority	actual	✓/×
5	5,4,3,2,1	4,3,2	1,0,1	1	0	×
6	6,5,4,3,2	4,3,2	1,0,1	1	1	✓
7	7,6,5,4,3	4,3,2	1,0,1	1	0	×
8	8,7,6,5,4	4,3,2	1,0,1	1	1	✓
9	9,8,7,6,5	4,3,2	1,0,1	1	0	×

Fold 2 accuracy: $2/5$.

Total 2-fold CV accuracy = $(2 + 2)/10 = \boxed{40\%}$.

Insight: KNN with $k = 3$ on alternating data is mediocre because the majority of any 3 consecutive Y 's flips with the mode of the alternating pattern.

Q7 [10 pts] — Decision tree on tiny binary dataset

Same as 2024 Q8. $H(Y) = 0.811$, classification accuracy = 75%.

Q8 [10 pts] — MDP grid world: best action from (x, y)

Actions $N, S, E, W \rightarrow (x-1, y), (x+1, y), (x, y+1), (x, y-1)$. *Stochastic:* 0.6 intended, 0.2 each perpendicular. *Reward* -1 per step. $\gamma = 0.53$. *Current* U_t at the 4 neighbours (in N, S, E, W order) = 4, 5, 4, 1.

Compute $Q_t(\text{action})$ for each.

For action N (intended $\rightarrow U_N = 4$, perp = E and W):

$$\mathbb{E}[U_{t+1} \mid N] = 0.6(U_N) + 0.2(U_E) + 0.2(U_W) = 0.6(4) + 0.2(4) + 0.2(1) = 2.4 + 0.8 + 0.2 = 3.4$$

$$Q(N) = -1 + 0.53(3.4) = -1 + 1.802 = \boxed{0.802}$$

For action S (perp = E, W):

$$\mathbb{E}[U_{t+1} \mid S] = 0.6(5) + 0.2(4) + 0.2(1) = 3 + 0.8 + 0.2 = 4.0$$

$$Q(S) = -1 + 0.53(4.0) = -1 + 2.12 = \boxed{1.12}$$

For action E (perp = N, S):

$$\mathbb{E}[U_{t+1} \mid E] = 0.6(4) + 0.2(4) + 0.2(5) = 2.4 + 0.8 + 1 = 4.2$$

$$Q(E) = -1 + 0.53(4.2) = -1 + 2.226 = \boxed{1.226}$$

For action W (perp = N, S):

$$\mathbb{E}[U_{t+1} \mid W] = 0.6(1) + 0.2(4) + 0.2(5) = 0.6 + 0.8 + 1 = 2.4$$

$$Q(W) = -1 + 0.53(2.4) = -1 + 1.272 = \boxed{0.272}$$

Best action: $\arg \max_a Q(a) = \boxed{\text{East (E)}}$ with $Q = 1.226$.

Insight: although South has the highest neighbour utility (5), the perpendicular spread of E (which has high-utility N and S as neighbours under stochasticity) makes E 's expected utility highest. This is the kind of arithmetic trap the question rewards — always compute all 4 Q values, don't just go with the highest single neighbour.

Q9 [10 pts]

(typically a substitute / repeat in the “answer any 7”; if it appears, it's likely a Naive-Bayes or Apriori variant covered in Ch. 4 or Ch. 11.)

Chapter 18

Cross-Year Pattern Map: What Recurs

Topics that have appeared in 3+ of the last 4 years (high probability for 2026):

Topic	Years asked
Bias-Variance numerical computation	2022 (implicit), 2023, 2024, 2025
Decision tree (ID3 or CART/Gini)	2022, 2023, 2024, 2025
KNN classification	2023, 2024, 2025
Hierarchical clustering (linkages)	2022, 2023, 2024
K-Means iteration	2022
MDP value/policy iteration	2022, 2023, 2024, 2025
MDP LP formulation (A matrix)	2023, 2024
State-space search (UCS/A*)	2022, 2025
Hill climbing / SA	2025
Local search heuristics for puzzles	2025 (Sudoku)
Propositional logic (truth table / CNF)	2023, 2024, 2025
Bayes rule / disease test	2023, 2024, 2025
Naive Bayes	2023
Apriori frequent itemsets	2023
Decision theory / game show / MEU	2023, 2024
POMDP / belief update	2023, 2024
100-coin probability puzzle	2023, 2024
Tic-Tac-Toe minimax	2022, 2024 (Q6)

For 2026, expect a near-certain hit on: bias-variance, decision tree (CART or ID3), MDP (value iteration or LP), search (A* + admissibility/consistency), propositional logic CNF/DNF, Bayes-rule disease test or KNN. The new topic this year is likely **multi-armed bandits (UCB)** since it's late-syllabus and only briefly touched in 2025.

Topics that may surprise you (rarer but possible):

- Genetic algorithm / fitness function (briefly in 2025 Q8c)
- Simulated annealing pseudocode (2025 Q5)
- Linear programming primal-dual MDP (2024 Q5d, 2023 Q3e)
- POMDP belief-state evolution
- Q-learning vs. value iteration discussion

Chapter 19

Last-Page Cheat Sheet

Bias-Variance. $\text{MSE} = \text{Bias}^2 + \text{Var} + \sigma^2$. $\lambda \uparrow \Rightarrow \text{Bias} \uparrow, \text{Var} \downarrow$.

Bayes. $P(A|B) = P(B|A)P(A)/P(B)$. Disease test 1%/99% $\Rightarrow$ posterior 50%.

Decision tree. ID3 max IG; CART min weighted Gini $= 1 - \sum p_k^2$.

KNN. $k = 1$: nearest. k odd: majority of top- k . Standardize features.

K-Means. Assign $\rightarrow$ Update $\rightarrow$ Repeat until stable. WCSS $= \sum \|x - \mu\|^2$.

Hierarchical. Single (min), Complete (max), Average (mean) of pairwise distances.

Risk attitude. Concave U = averse, linear = neutral, convex = seeking.

MEU. $a^* = \arg \max_a \sum_o P(o|a)U(o)$.

MDP. $V^*(s) = \max_a [R + \gamma \sum T V^*(s')]$. Finite horizon: backward induction $V_T \rightarrow V_0$.

Bandits. UCB1 $= \hat{\mu} + \sqrt{2 \ln t / N}$. Regret $= O(\log T)$.

Search. BFS/DFS/UCS/A*. A*: $f = g + h$. Admissible $h \leq h^*$. Consistent $h \leq c + h'$.

Hill climbing. Greedy local. Stuck at local optima.

Simulated Annealing. $P(\text{accept worse}) = e^{\Delta/T}$.

Apriori. Prune candidate if any subset is infrequent. Support, Confidence, Lift.

Logic. CNF: $\wedge$ of $\vee$. DNF: $\vee$ of $\wedge$. Modus Ponens, Modus Tollens, contrapositive.

In the exam: state every formula before substituting. Show working. Box your final answer. If you panic, write the recipe steps from memory and you'll recover most marks.