

SMAI – Complete Study Guide

Statistical Methods in Artificial Intelligence

IIIT Hyderabad • Jan–Feb 2026

Auto-generated from Lecture Notes

Last compiled: March 1, 2026

Contents

1	Introduction & Course Overview	3
1.1	Types of Machine Learning	3
1.2	Data Representation	3
1.3	Classification vs. Regression	4
1.4	Nearest Neighbour (NN) Classifier	4
2	Linear Regression	5
2.1	The Model	5
2.2	Loss Function: Mean Squared Error (MSE)	5
2.3	Closed-Form Solution (Normal Equation)	5
2.4	Gradient Descent	6
3	Regularization	7
3.1	Overfitting vs. Underfitting	7
3.2	L2 Regularization (Ridge Regression)	7
3.3	L1 Regularization (Lasso Regression)	7
3.4	L1 vs. L2 Comparison	8
4	Linear Algebra Foundations	9
4.1	Matrix Rank	9
4.2	Eigenvalues and Eigenvectors	9
4.3	Singular Value Decomposition (SVD)	10
4.4	Low-Rank Approximation	10
5	Principal Component Analysis (PCA)	11
5.1	Step-by-Step PCA Algorithm	11
5.2	The Covariance Matrix	11
5.3	Why Eigenvectors of Covariance Matrix?	11
5.4	Reconstruction from PCA	12
5.5	Variance Explained	12
5.6	PCA via SVD	12
5.7	Important Properties of PCA	12
5.8	Applications of PCA & SVD	12

6	Bayesian Decision Theory	14
6.1	Bayes' Theorem	14
6.2	Bayesian Decision Rule	14
6.3	Normal (Gaussian) Distribution	14
6.4	Bayesian Decision Boundary	15
7	Maximum Likelihood Estimation & Gaussian Mixture Models	16
7.1	Maximum Likelihood Estimation (MLE)	16
7.2	Gaussian Mixture Model (GMM)	16
7.3	Expectation-Maximization (EM) Algorithm	16
7.4	K-Means Clustering	17
8	Classification Algorithms	18
8.1	Naive Bayes Classifier	18
8.2	Logistic Regression	18
8.3	Softmax Regression (Multi-Class)	19
8.4	Loss Function Comparison	19
9	Support Vector Machines (SVM)	20
9.1	Geometric Intuition	20
9.2	Hard-Margin SVM	20
9.3	Soft-Margin SVM	20
9.4	The Kernel Trick	20
10	Practical Machine Learning & Model Evaluation	22
10.1	Data Splitting	22
10.2	Cross-Validation	22
10.3	Performance Metrics	22
10.4	Data Normalization	23
10.5	Hyperparameter Tuning	23
11	Neural Networks	24
11.1	Single-Layer Perceptron	24
11.2	Perceptron Learning Rule	24
11.3	Delta Rule (for Regression)	25
11.4	Activation Functions	25
11.5	Multi-Layer Perceptron (MLP)	25
11.6	Biological vs. Artificial Neurons	25
12	Formula Reference Sheet	26
13	Practice Exam Questions	28
14	Quick-Reference: Topic-by-Lecture Map	30
15	Glossary of Key Terms	31

1 Introduction & Course Overview

What is this course about?

This course covers **Statistical Methods in AI (SMAI)**, split roughly 50:50 between:

1. **Classical / Basic Methods:** Linear regression, classification, regularization, PCA, SVD, Bayesian methods, SVMs.
2. **Neural Networks & Deep Learning:** Perceptrons, multi-layer networks, backpropagation, modern architectures.

1.1 Types of Machine Learning

Supervised Learning

The algorithm learns from **labelled data** — each training example comes with a correct answer (label).

Given $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$, learn $f: \mathbf{x} \mapsto y$.

Examples: spam detection, image classification, predicting house prices.

Unsupervised Learning

The algorithm learns from **unlabelled data** — no correct answers are provided. The goal is to discover hidden patterns, groupings, or structure.

Given $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$, find patterns.

Examples: clustering (K-means), dimensionality reduction (PCA), anomaly detection.

Self-Supervised Learning

A variant where the algorithm creates its own labels from the unlabelled data and uses them to learn useful representations. Very popular in modern NLP and computer vision.

1.2 Data Representation

Vectors as Data

Every data point (sample) is represented as a **feature vector**:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \in \mathbb{R}^d$$

where d is the **dimensionality** (number of features).

Example: Representing a student by features like height, weight, GPA, age gives $d = 4$.

Handcrafted vs. Learned Representations

- **Handcrafted:** Features designed by a human expert using domain knowledge (e.g., color histogram for images).

- **Learned:** Features automatically extracted from data by algorithms, especially neural networks (e.g., word embeddings, CNN features).

Modern deep learning mostly uses *learned* representations.

1.3 Classification vs. Regression

Aspect	Classification	Regression
Output type	Discrete / Categorical ($y \in \{0, 1, \dots, K\}$)	Continuous ($y \in \mathbb{R}$)
Goal	Predict a <i>class label</i>	Predict a <i>real number</i>
Example	Spam vs. Not Spam	Predict house price
Loss function	Cross-entropy, 0–1 loss	Mean Squared Error (MSE)

1.4 Nearest Neighbour (NN) Classifier

Nearest Neighbour Algorithm – The Simplest Classifier

Given a new test point $\mathbf{x}^*$:

1. Compute the distance between $\mathbf{x}^*$ and *every* training point.
2. Find the training point $\mathbf{x}_j$ that is closest (smallest distance).
3. Predict the label y_j of that nearest training point.

Distance metric: Typically Euclidean distance $\|\mathbf{x}^* - \mathbf{x}_j\|_2$.

Limitations of Nearest Neighbour

- **Slow at test time:** Must compare with *all* n training points — $O(nd)$ per prediction.
- **Curse of dimensionality:** Performance degrades in very high dimensions.
- **Stores entire dataset:** Memory-intensive.

2 Linear Regression

What is Linear Regression?

Linear regression is a **supervised learning** algorithm that predicts a **continuous** output by fitting a linear relationship between input features and the target variable.

2.1 The Model

Linear Regression Equation

$$\hat{y} = \mathbf{w}^\top \mathbf{x} + b = w_1x_1 + w_2x_2 + \dots + w_dx_d + b$$

- $\mathbf{w} = [w_1, w_2, \dots, w_d]^\top$ is the **weight vector** (parameters to learn).
- b is the **bias** (intercept).
- $\mathbf{x} = [x_1, x_2, \dots, x_d]^\top$ is the **input feature vector**.
- $\hat{y}$ is the **predicted output**.

Compact Notation

We often absorb the bias into the weight vector by appending a 1 to the feature vector:

$$\tilde{\mathbf{x}} = \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}, \quad \tilde{\mathbf{w}} = \begin{bmatrix} \mathbf{w} \\ b \end{bmatrix}, \quad \hat{y} = \tilde{\mathbf{w}}^\top \tilde{\mathbf{x}}.$$

2.2 Loss Function: Mean Squared Error (MSE)

Mean Squared Error

$$\mathcal{L}(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n \left(y_i - \mathbf{w}^\top \mathbf{x}_i \right)^2$$

This measures the average squared difference between predicted and actual values. Our goal is to **minimize** this loss.

2.3 Closed-Form Solution (Normal Equation)

Normal Equation

Stack all data into a matrix $\mathbf{X} \in \mathbb{R}^{n \times d}$ and vector $\mathbf{y} \in \mathbb{R}^n$. The optimal weights are:

$$\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

- This requires $\mathbf{X}^\top \mathbf{X}$ to be **invertible**.
- Computational cost: $O(d^3 + nd^2)$ — expensive when d is large.

2.4 Gradient Descent

Gradient Descent — Iterative Optimization

When the closed-form solution is expensive, we use **gradient descent**:

1. Initialize $\mathbf{w}$ randomly or as zeros.
2. Compute the gradient: $\nabla_{\mathbf{w}}\mathcal{L} = -\frac{2}{n}\mathbf{X}^\top(\mathbf{y} - \mathbf{X}\mathbf{w})$.
3. Update: $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}}\mathcal{L}$, where η is the **learning rate**.
4. Repeat until convergence.

Variant	Description	Pros / Cons
Batch GD	Uses <i>all</i> n samples per update	Stable but slow for large n
Stochastic GD (SGD)	Uses <i>one</i> random sample per update	Fast but noisy
Mini-batch GD	Uses a small batch of B samples	Best trade-off in practice

Closed-Form vs. Gradient Descent

	Closed-Form	Gradient Descent
Speed	One-shot computation	Iterative
Memory	Needs to store $\mathbf{X}^\top\mathbf{X}$	Can work on mini-batches
Scalability	Poor when d or n large	Scales well
Parallelizable	Yes (matrix operations)	Yes (each sample independent)

3 Regularization

Why Regularization?

Regularization is a technique to **prevent overfitting** by adding a **penalty term** to the loss function. This discourages the model from learning overly complex patterns that fit noise in the training data.

$$\mathcal{L}_{\text{regularized}} = \underbrace{\mathcal{L}_{\text{data}}}_{\text{fit the data}} + \underbrace{\lambda \cdot R(\mathbf{w})}_{\text{keep weights small}}$$

where $\lambda > 0$ is the **regularization strength** (hyperparameter).

3.1 Overfitting vs. Underfitting

The Central Problem of Machine Learning

- **Overfitting (High Variance)**: Model is too complex. Fits training data perfectly but fails on new data. Training error $\ll$ test error.
- **Underfitting (High Bias)**: Model is too simple. Cannot capture the underlying patterns. Both training error and test error are high.

Goal: Find the sweet spot between bias and variance — the **bias–variance trade-off**.

3.2 L2 Regularization (Ridge Regression)

Ridge Regression

Add the squared magnitude of weights as penalty:

$$\mathcal{L}_{\text{Ridge}} = \frac{1}{n} \sum_{i=1}^n (y_i - \mathbf{w}^\top \mathbf{x}_i)^2 + \lambda \|\mathbf{w}\|_2^2$$

where $\|\mathbf{w}\|_2^2 = w_1^2 + w_2^2 + \dots + w_d^2$.

Closed-form solution: $\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}$

Effect: Shrinks all weights towards zero, but does *not* set them exactly to zero.

Geometric interpretation: Constraint region is a **circle** (or hypersphere).

3.3 L1 Regularization (Lasso Regression)

Lasso Regression

Add the absolute value of weights as penalty:

$$\mathcal{L}_{\text{Lasso}} = \frac{1}{n} \sum_{i=1}^n (y_i - \mathbf{w}^\top \mathbf{x}_i)^2 + \lambda \|\mathbf{w}\|_1$$

where $\|\mathbf{w}\|_1 = |w_1| + |w_2| + \dots + |w_d|$.

Effect: Can set some weights *exactly to zero* $\Rightarrow$ **feature selection** (sparse model).

Geometric interpretation: Constraint region is a **diamond** (in 2D).

3.4 L1 vs. L2 Comparison

Property	L1 (Lasso)	L2 (Ridge)
Penalty term	$\lambda \sum w_j $	$\lambda \sum w_j^2$
Sparsity	Yes — some weights become 0	No — weights are small but non-zero
Feature selection	Built-in	Not directly
Constraint shape	Diamond	Circle
Differentiable?	No (at $w_j = 0$)	Yes, everywhere
Best for	Many irrelevant features	All features contribute somewhat

4 Linear Algebra Foundations

4.1 Matrix Rank

Rank of a Matrix

The **rank** of a matrix $\mathbf{A}$ is the maximum number of **linearly independent** rows (or columns).

- If $\mathbf{A} \in \mathbb{R}^{m \times n}$, then $\text{rank}(\mathbf{A}) \leq \min(m, n)$.
- Rank = number of **non-zero singular values**.
- Rank = number of **non-zero eigenvalues** (for square matrices).

Linear Independence

Vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_k$ are **linearly independent** if no one vector can be written as a linear combination of the others. Formally:

$$c_1 \mathbf{v}_1 + c_2 \mathbf{v}_2 + \dots + c_k \mathbf{v}_k = \mathbf{0} \implies c_1 = c_2 = \dots = c_k = 0.$$

4.2 Eigenvalues and Eigenvectors

Eigenvalue Equation

For a square matrix $\mathbf{A} \in \mathbb{R}^{d \times d}$, a scalar λ and non-zero vector $\mathbf{v}$ satisfying

$$\boxed{\mathbf{A}\mathbf{v} = \lambda\mathbf{v}}$$

are called an **eigenvalue** (λ) and **eigenvector** ($\mathbf{v}$) of $\mathbf{A}$.

What do eigenvalues/vectors mean?

- An **eigenvector** gives a *direction* that is unchanged by the transformation $\mathbf{A}$ (only scaled).
- The corresponding **eigenvalue** tells you *how much* the vector is stretched or shrunk along that direction.
- If $\lambda > 1$: stretching. If $0 < \lambda < 1$: shrinking. If $\lambda < 0$: flipping.

Eigendecomposition

If $\mathbf{A}$ is a symmetric matrix, it can be decomposed as:

$$\mathbf{A} = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^\top$$

where:

- $\mathbf{Q}$ is an orthogonal matrix whose columns are eigenvectors.
- $\mathbf{\Lambda}$ is a diagonal matrix of eigenvalues: $\mathbf{\Lambda} = \text{diag}(\lambda_1, \lambda_2, \dots, \lambda_d)$.

4.3 Singular Value Decomposition (SVD)

SVD – The Most Important Decomposition

Any matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ can be decomposed as:

$$\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^\top$$

where:

- $\mathbf{U} \in \mathbb{R}^{m \times m}$: orthogonal matrix (left singular vectors).
- $\mathbf{\Sigma} \in \mathbb{R}^{m \times n}$: diagonal matrix of **singular values** $\sigma_1 \geq \sigma_2 \geq \dots \geq 0$.
- $\mathbf{V} \in \mathbb{R}^{n \times n}$: orthogonal matrix (right singular vectors).

Key Properties of SVD

- SVD works for **any** matrix (not just square or symmetric).
- The number of non-zero singular values = rank of $\mathbf{A}$.
- Singular values are always ≥ 0 .
- Relationship to eigenvalues: $\sigma_i = \sqrt{\lambda_i(\mathbf{A}^\top \mathbf{A})}$.
- Can be used to compute pseudo-inverses, low-rank approximations, and more.

4.4 Low-Rank Approximation

Low-Rank Approximation via SVD

Given $\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^\top$, the best rank- k approximation (in Frobenius norm) is:

$$\mathbf{A}_k = \sum_{i=1}^k \sigma_i \mathbf{u}_i \mathbf{v}_i^\top$$

This is the **Eckart–Young theorem**. We keep only the top k singular values and their vectors, discarding the rest as “noise.”

Applications: Image compression, denoising, recommendation systems, latent semantic analysis.

Choosing k — How many components to keep?

Plot the singular values $\sigma_1 \geq \sigma_2 \geq \dots$. Look for a **sharp drop** (“elbow”) in the spectrum. Components after the elbow mostly capture noise.

5 Principal Component Analysis (PCA)

What is PCA?

PCA is the most popular **unsupervised dimensionality reduction** technique. It finds new axes (called **principal components**) along which the data has **maximum variance**.

Goal: Reduce dimensionality from d to d' (where $d' \ll d$) while retaining as much information as possible.

5.1 Step-by-Step PCA Algorithm

PCA Algorithm

1. **Center the data:** Subtract the mean: $\bar{\mathbf{x}} = \frac{1}{n} \sum_{i=1}^n \mathbf{x}_i$, then $\tilde{\mathbf{x}}_i = \mathbf{x}_i - \bar{\mathbf{x}}$.

2. **Compute the covariance matrix:**

$$\Sigma = \frac{1}{n} \sum_{i=1}^n \tilde{\mathbf{x}}_i \tilde{\mathbf{x}}_i^\top = \frac{1}{n} \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \in \mathbb{R}^{d \times d}$$

3. **Compute eigenvalues and eigenvectors of Σ :** $\Sigma \mathbf{v}_j = \lambda_j \mathbf{v}_j$.

4. **Sort** eigenvectors by decreasing eigenvalue: $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$.

5. **Select top d' eigenvectors:** Form $\mathbf{U} = [\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_{d'}] \in \mathbb{R}^{d \times d'}$.

6. **Project:** $\mathbf{x}'_i = \mathbf{U}^\top \tilde{\mathbf{x}}_i \in \mathbb{R}^{d'}$.

5.2 The Covariance Matrix

Covariance Matrix – Intuition

The covariance matrix Σ captures how features **vary together**:

$$\Sigma_{jk} = \frac{1}{n} \sum_{i=1}^n (x_{ij} - \bar{x}_j)(x_{ik} - \bar{x}_k)$$

- **Diagonal entries** Σ_{jj} = variance of feature j .
- **Off-diagonal entries** Σ_{jk} = covariance between features j and k .
- Σ is always **symmetric** and **positive semi-definite**.

5.3 Why Eigenvectors of Covariance Matrix?

PCA = Max Variance Directions

The first principal component $\mathbf{v}_1$ is the direction that maximizes the variance of the projected data:

$$\mathbf{v}_1 = \arg \max_{\|\mathbf{v}\|=1} \text{Var}(\mathbf{v}^\top \tilde{\mathbf{X}}) = \arg \max_{\|\mathbf{v}\|=1} \mathbf{v}^\top \Sigma \mathbf{v}$$

The solution is the **eigenvector** of Σ corresponding to the **largest eigenvalue**.

Each subsequent principal component is the direction of maximum variance *orthogonal to all previous ones*.

5.4 Reconstruction from PCA

Reconstruction

After projecting $\mathbf{x}'_i = \mathbf{U}^\top \tilde{\mathbf{x}}_i$, we can reconstruct approximately:

$$\hat{\mathbf{x}}_i = \mathbf{U} \mathbf{x}'_i + \bar{\mathbf{x}}$$

This reconstruction is **lossy** if $d' < d$ because we discarded some eigenvectors (those with small eigenvalues).

5.5 Variance Explained

Proportion of Variance Explained

The fraction of total variance captured by the top d' components is:

$$\text{Variance Explained} = \frac{\sum_{j=1}^{d'} \lambda_j}{\sum_{j=1}^d \lambda_j}$$

Common rule of thumb: Keep enough components to explain **90–95%** of the variance.

5.6 PCA via SVD

Connection: PCA and SVD

Given centered data matrix $\tilde{\mathbf{X}} \in \mathbb{R}^{n \times d}$ with SVD: $\tilde{\mathbf{X}} = \mathbf{U}_X \Sigma_X \mathbf{V}_X^\top$:

- The right singular vectors $\mathbf{V}_X$ are the **principal components** (eigenvectors of $\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}}$).
- The singular values relate to eigenvalues: $\lambda_j = \frac{\sigma_j^2}{n}$.
- SVD is often more **numerically stable** and efficient than explicitly computing the covariance matrix.

5.7 Important Properties of PCA

PCA Gotchas

- PCA is **sensitive to scale**. Always **standardize** (zero mean, unit variance) features before PCA.
- PCA is a **linear** technique — it can only capture linear relationships.
- PCA assumes the directions of maximum variance are the most “interesting.”
- For non-linear dimensionality reduction, consider **t-SNE** or **UMAP**.

5.8 Applications of PCA & SVD

Real-World Applications

- **Image Compression**: Represent images using top- k principal components $\Rightarrow$ fewer bytes.
- **Face Recognition (Eigenfaces)**: PCA on face images finds the main “face directions.”

- **Denoising:** Low-rank approximation via SVD removes noise from data.
- **Latent Semantic Indexing (LSI):** Apply SVD to term-document matrix to find latent topics.
- **Word Embeddings:** Dense vector representations of words capturing semantic meaning.
- **Recommendation Systems:** Matrix factorization (SVD-based) for collaborative filtering (e.g., Netflix).
- **Data Visualization:** Project high-dimensional data to 2D / 3D for plotting.
- **Preprocessing for ML:** Reduce features before feeding into classifiers to reduce overfitting.

6 Bayesian Decision Theory

The Bayesian Approach

Bayesian decision theory provides a **principled framework** for making optimal decisions under **uncertainty**. The key idea: use **probability** to represent uncertainty and **Bayes' theorem** to update beliefs when new evidence arrives.

6.1 Bayes' Theorem

Bayes' Theorem – The Foundation

$$P(\omega_j | \mathbf{x}) = \frac{P(\mathbf{x} | \omega_j) P(\omega_j)}{P(\mathbf{x})}$$

where:

- $P(\omega_j | \mathbf{x})$ = **Posterior**: probability of class ω_j given observation $\mathbf{x}$.
- $P(\mathbf{x} | \omega_j)$ = **Likelihood**: probability of observing $\mathbf{x}$ if it belongs to class ω_j .
- $P(\omega_j)$ = **Prior**: probability of class ω_j before seeing $\mathbf{x}$.
- $P(\mathbf{x})$ = **Evidence**: total probability of observing $\mathbf{x}$ (normalization constant).

When to Remember

$$\text{Posterior} = \frac{\text{Likelihood} \times \text{Prior}}{\text{Evidence}}$$

6.2 Bayesian Decision Rule

Optimal Decision

Classify $\mathbf{x}$ as class ω_j that maximizes the posterior:

$$\hat{y} = \arg \max_j P(\omega_j | \mathbf{x})$$

Since $P(\mathbf{x})$ is the same for all classes, this is equivalent to:

$$\hat{y} = \arg \max_j P(\mathbf{x} | \omega_j) P(\omega_j)$$

This decision rule **minimizes the probability of error**.

6.3 Normal (Gaussian) Distribution

Univariate Normal Distribution

A random variable X follows a normal distribution $\mathcal{N}(\mu, \sigma^2)$ if its PDF is:

$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

- μ = **mean** (center of the bell curve).
- σ^2 = **variance** (spread of the bell curve).

- σ = standard deviation.

Multivariate Normal Distribution

For a d -dimensional random vector $\mathbf{x}$:

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

- $\boldsymbol{\mu} \in \mathbb{R}^d$ = **mean vector**.
- $\Sigma \in \mathbb{R}^{d \times d}$ = **covariance matrix** (symmetric, positive definite).
- $|\Sigma|$ = determinant of Σ .

6.4 Bayesian Decision Boundary

Decision Boundary Between Two Classes

For two Gaussian classes ω_1, ω_2 with parameters $(\boldsymbol{\mu}_1, \Sigma_1)$ and $(\boldsymbol{\mu}_2, \Sigma_2)$:

Case 1: Same covariance ($\Sigma_1 = \Sigma_2 = \Sigma$):

The decision boundary is **linear** (a hyperplane).

Case 2: Different covariances ($\Sigma_1 \neq \Sigma_2$):

The decision boundary is **quadratic** (a curve / ellipse).

7 Maximum Likelihood Estimation & Gaussian Mixture Models

7.1 Maximum Likelihood Estimation (MLE)

What is MLE?

Given observed data $\{x_1, x_2, \dots, x_n\}$ and a parametric model (e.g., Gaussian with unknown μ, σ^2), MLE finds the parameter values that make the observed data **most probable**.

MLE for Normal Distribution

Likelihood function (assuming i.i.d. samples):

$$L(\mu, \sigma^2) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right)$$

Log-likelihood (easier to optimize):

$$\ell(\mu, \sigma^2) = -\frac{n}{2} \ln(2\pi) - \frac{n}{2} \ln(\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$$

MLE solutions (set derivatives to zero):

$$\hat{\mu}_{\text{MLE}} = \frac{1}{n} \sum_{i=1}^n x_i \quad \hat{\sigma}_{\text{MLE}}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

7.2 Gaussian Mixture Model (GMM)

GMM – Modeling Complex Distributions

A GMM models the data as coming from a **mixture of K Gaussian distributions**:

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k \mathcal{N}(\mathbf{x} \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$$

- π_k = **mixing coefficient** (weight) of component k . Must satisfy $\sum_k \pi_k = 1$, $\pi_k \geq 0$.
- $\boldsymbol{\mu}_k$ = mean of component k .
- $\boldsymbol{\Sigma}_k$ = covariance of component k .

Applications: clustering, density estimation, anomaly detection.

7.3 Expectation-Maximization (EM) Algorithm

EM Algorithm for GMM

Goal: Find parameters $\{\pi_k, \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k\}$ that maximize the likelihood of the data.

Iterate until convergence:

E-Step (Expectation): Compute the “responsibility” of component k for data point i :

$$\gamma_{ik} = \frac{\pi_k \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(\mathbf{x}_i \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)}$$

M-Step (Maximization): Update parameters:

$$N_k = \sum_{i=1}^n \gamma_{ik}, \quad \boldsymbol{\mu}_k = \frac{1}{N_k} \sum_{i=1}^n \gamma_{ik} \mathbf{x}_i$$
$$\boldsymbol{\Sigma}_k = \frac{1}{N_k} \sum_{i=1}^n \gamma_{ik} (\mathbf{x}_i - \boldsymbol{\mu}_k)(\mathbf{x}_i - \boldsymbol{\mu}_k)^\top, \quad \pi_k = \frac{N_k}{n}$$

EM Caveats

- EM converges to a **local maximum**, not necessarily the global one.
- **Initialization matters**: Different starting points can give different solutions. Run multiple times.
- Must choose K (number of components) — use BIC, AIC, or cross-validation.

7.4 K-Means Clustering

K-Means — A Simpler Clustering Algorithm

K-means is a **hard clustering** algorithm (each point belongs to exactly one cluster). It can be seen as a special case of GMM with equal, spherical covariances.

Algorithm:

1. **Initialize** K cluster centroids $\boldsymbol{\mu}_1, \dots, \boldsymbol{\mu}_K$ (randomly).
2. **Assign** each point to the nearest centroid: $c_i = \arg \min_k \|\mathbf{x}_i - \boldsymbol{\mu}_k\|^2$.
3. **Update** centroids: $\boldsymbol{\mu}_k = \frac{1}{|C_k|} \sum_{i \in C_k} \mathbf{x}_i$.
4. **Repeat** steps 2–3 until assignments stop changing.

8 Classification Algorithms

8.1 Naive Bayes Classifier

Naive Bayes

A probabilistic classifier based on Bayes' theorem with the **naive** assumption that all features are **conditionally independent** given the class:

$$P(\mathbf{x} \mid \omega_j) = \prod_{i=1}^d P(x_i \mid \omega_j)$$

Decision rule:

$$\hat{y} = \arg \max_j P(\omega_j) \prod_{i=1}^d P(x_i \mid \omega_j)$$

Pros: Simple, fast, works well with high-dimensional data (e.g., text).

Cons: Independence assumption is often violated in practice.

Applications: Spam filtering, text classification, sentiment analysis.

8.2 Logistic Regression

Logistic Regression – Despite the Name, It's for Classification!

Logistic regression models the **probability** that $\mathbf{x}$ belongs to class 1:

$$P(y = 1 \mid \mathbf{x}) = \sigma(\mathbf{w}^\top \mathbf{x} + b) = \frac{1}{1 + e^{-(\mathbf{w}^\top \mathbf{x} + b)}}$$

where $\sigma(z) = \frac{1}{1+e^{-z}}$ is the **sigmoid (logistic) function**.

Sigmoid Function Properties

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- Output is always in $(0, 1)$ — can be interpreted as a probability.
- $\sigma(0) = 0.5$.
- $\sigma(-z) = 1 - \sigma(z)$.
- Derivative: $\sigma'(z) = \sigma(z)(1 - \sigma(z))$.

Cross-Entropy Loss for Logistic Regression

$$\mathcal{L}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n [y_i \ln \hat{p}_i + (1 - y_i) \ln(1 - \hat{p}_i)]$$

where $\hat{p}_i = \sigma(\mathbf{w}^\top \mathbf{x}_i + b)$. Minimize via gradient descent.

8.3 Softmax Regression (Multi-Class)

Softmax Function

For K classes, the probability of class k is:

$$P(y = k \mid \mathbf{x}) = \frac{e^{\mathbf{w}_k^\top \mathbf{x}}}{\sum_{j=1}^K e^{\mathbf{w}_j^\top \mathbf{x}}}$$

This generalizes the sigmoid to multiple classes. The loss is **categorical cross-entropy**.

8.4 Loss Function Comparison

Loss Function	Description	Used For
0-1 Loss	$\ell = \mathbb{I}[\hat{y} \neq y]$	Ideal but non-differentiable
Hinge Loss	$\ell = \max(0, 1 - y \cdot f(x))$	SVM
Logistic Loss	$\ell = \ln(1 + e^{-y \cdot f(x)})$	Logistic Regression
Cross-Entropy	$-\sum y_k \ln \hat{p}_k$	Multi-class classification
MSE	$\frac{1}{n} \sum (y_i - \hat{y}_i)^2$	Regression

9 Support Vector Machines (SVM)

What is an SVM?

A Support Vector Machine finds the **hyperplane** that **maximally separates** two classes. Instead of just finding *any* separating hyperplane, SVM finds the one with the **largest margin**.

9.1 Geometric Intuition

Margin

The **margin** is the distance between the decision boundary (hyperplane) and the closest data points from either class. These closest points are called **support vectors**.

SVM objective: Maximize the margin. A larger margin means better generalization.

9.2 Hard-Margin SVM

Hard-Margin SVM Formulation

Assume data is linearly separable. Find $\mathbf{w}$ and b that:

$$\min_{\mathbf{w}, b} \frac{1}{2} \|\mathbf{w}\|^2 \quad \text{subject to} \quad y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \quad \forall i$$

- The margin width is $\frac{2}{\|\mathbf{w}\|}$.
- Minimizing $\|\mathbf{w}\|^2 \Leftrightarrow$ maximizing the margin.
- Constraints ensure all points are correctly classified with margin ≥ 1 .

9.3 Soft-Margin SVM

Soft-Margin SVM – Handling Non-Separable Data

Allow some misclassifications by introducing **slack variables** $\xi_i \geq 0$:

$$\min_{\mathbf{w}, b, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \quad \text{s.t.} \quad y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0$$

- $C > 0$ = regularization parameter controlling the trade-off between margin and errors.
- Large C : fewer errors allowed (could overfit).
- Small C : wider margin but more errors allowed (could underfit).

9.4 The Kernel Trick

Kernel Trick – Making SVM Non-Linear

When data is **not linearly separable**, we map it to a **higher-dimensional space** where it *is* separable. The kernel trick lets us do this **without explicitly computing the mapping**.

A **kernel function** $K(\mathbf{x}_i, \mathbf{x}_j) = \langle \phi(\mathbf{x}_i), \phi(\mathbf{x}_j) \rangle$ computes the inner product in the high-

dimensional space directly.

Kernel	Formula
Linear	$K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^\top \mathbf{x}_j$
Polynomial (degree p)	$K(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i^\top \mathbf{x}_j + c)^p$
RBF / Gaussian	$K(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\frac{\ \mathbf{x}_i - \mathbf{x}_j\ ^2}{2\sigma^2}\right)$

The XOR Problem

XOR is **not linearly separable** in 2D. But mapping to a higher dimension using a polynomial kernel makes it separable. This is a classic example of why kernels are powerful.

SVM Advantages & Disadvantages

Advantages:

- Works well in high-dimensional spaces.
- Effective when number of features $>$ number of samples.
- Memory efficient (only stores support vectors).
- Convex optimization — guaranteed global optimum (for given kernel).

Disadvantages:

- Doesn't directly output probabilities.
- Slow for very large datasets.
- Sensitive to feature scaling.
- Kernel choice requires experimentation.

10 Practical Machine Learning & Model Evaluation

The Most Important Lesson

Theory alone is not enough! You need hands-on practice. There is no single “best” algorithm — it depends on the data, the problem, and the computational resources.

10.1 Data Splitting

Train / Validation / Test Split

1. **Training Set** (e.g., 60–70%): Used to **train** the model (fit parameters).
2. **Validation Set** (e.g., 15–20%): Used to **tune hyperparameters** and **select the best model**.
3. **Test Set** (e.g., 15–20%): Used **only once at the very end** to report final performance.

Golden Rule: Never let the model see the test set during training or tuning!

10.2 Cross-Validation

k -Fold Cross-Validation

1. Split training data into k equal folds.
2. For each fold i : train on the other $k - 1$ folds, validate on fold i .
3. Average the k validation scores.

Typical: $k = 5$ or $k = 10$.

Benefit: Uses all data for both training and validation; more reliable estimate of performance.

10.3 Performance Metrics

Metric	Description
Accuracy	$\frac{\text{Correct predictions}}{\text{Total predictions}}$ — can be misleading with imbalanced classes
Precision	$\frac{TP}{TP+FP}$ — of those predicted positive, how many are actually positive?
Recall (Sensitivity)	$\frac{TP}{TP+FN}$ — of the actual positives, how many did we catch?
F1 Score	$2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$ — harmonic mean of precision and recall

Confusion Matrix

A table showing actual vs. predicted classes:

	Predicted +	Predicted −
Actual +	True Positive (TP)	False Negative (FN)
Actual −	False Positive (FP)	True Negative (TN)

10.4 Data Normalization

Why Normalize?

Features can have very different scales (e.g., age: 0–100, salary: 0–1,000,000). Algorithms like gradient descent, SVM, and PCA are sensitive to scale.

Common methods:

- **Min-Max Scaling:** $x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}} \in [0, 1]$.
- **Standardization (Z-score):** $x' = \frac{x - \mu}{\sigma}$ (zero mean, unit variance).

10.5 Hyperparameter Tuning

Common Hyperparameters

Algorithm	Key Hyperparameters
Linear Regression	Regularization λ
Logistic Regression	Regularization λ , learning rate η
SVM	C , kernel type, σ (for RBF)
K-Means	K (number of clusters)
PCA	d' (number of components)
Neural Network	Learning rate, number of layers, number of neurons, batch size

11 Neural Networks

What is a Neural Network?

A neural network is a **computational model** loosely inspired by biological neural networks. It is a **function** that maps inputs to outputs through a series of **layers** of interconnected “neurons.”

Key insight: Neural networks are **universal function approximators** — they can learn almost any mapping from inputs to outputs given enough data and capacity.

11.1 Single-Layer Perceptron

The Perceptron – The Simplest Neural Network

A single neuron computes:

$$\hat{y} = \text{activation}\left(\sum_{j=1}^d w_j x_j + b\right) = \text{activation}(\mathbf{w}^\top \mathbf{x} + b)$$

- **Inputs:** $x_1, x_2, \dots, x_d$.
- **Weights:** $w_1, w_2, \dots, w_d$ (learned from data).
- **Bias:** b .
- **Activation function:** Introduces non-linearity (e.g., step function, sigmoid, ReLU).
- **Output:** $\hat{y}$.

Limitation of Single-Layer Perceptron

A single-layer perceptron can only learn **linearly separable** functions. It **cannot** learn XOR!

Solution: Use **multi-layer** networks with non-linear activations.

11.2 Perceptron Learning Rule

Perceptron Learning Rule (for Classification)

For a misclassified point $(\mathbf{x}_i, y_i)$:

$$\mathbf{w} \leftarrow \mathbf{w} + \eta y_i \mathbf{x}_i$$

- $\eta > 0$ is the learning rate.
- $y_i \in \{-1, +1\}$ is the true label.
- Updates only when a mistake is made.
- **Guaranteed to converge** if data is linearly separable.

11.3 Delta Rule (for Regression)

Delta Rule

Minimize the MSE by gradient descent:

$$\mathbf{w} \leftarrow \mathbf{w} + \eta (y_i - \hat{y}_i) \mathbf{x}_i$$

- Updates proportional to the **error** ($y_i - \hat{y}_i$).
- Works for **continuous outputs** (regression).
- Essentially gradient descent on MSE.

11.4 Activation Functions

Function	Formula	Notes
Step	$f(z) = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases}$	Used in original perceptron; not differentiable
Sigmoid	$\sigma(z) = \frac{1}{1+e^{-z}}$	Output in $(0, 1)$; used for probabilities
Tanh	$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Output in $(-1, 1)$; zero-centered
ReLU	$f(z) = \max(0, z)$	Most popular in deep learning; simple and fast

11.5 Multi-Layer Perceptron (MLP)

MLP Architecture

An MLP consists of:

- **Input layer:** Receives the feature vector $\mathbf{x}$.
- **Hidden layer(s):** One or more layers of neurons that learn intermediate representations.
- **Output layer:** Produces the final prediction.

Each layer computes: $\mathbf{h} = \text{activation}(\mathbf{W}\mathbf{x} + \mathbf{b})$.

With at least **one hidden layer** and non-linear activations, an MLP can approximate any continuous function (**Universal Approximation Theorem**).

11.6 Biological vs. Artificial Neurons

Important Distinction

Aspect	Biological Neuron	Artificial Neuron
Input	Dendrites	Feature vector $\mathbf{x}$
Processing	Cell body (soma)	Weighted sum + activation
Output	Axon	$\hat{y}$
Learning	Synaptic plasticity	Weight updates via gradient descent
Speed	Slow ($\sim$ ms)	Fast (nanoseconds)

Artificial neural networks are **engineering tools**, not biological models.

12 Formula Reference Sheet

All Key Formulas at a Glance

Linear Regression:

$$\hat{y} = \mathbf{w}^\top \mathbf{x} + b, \quad \mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

MSE Loss:

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Ridge Regression:

$$\mathcal{L} = \frac{1}{n} \sum (y_i - \hat{y}_i)^2 + \lambda \|\mathbf{w}\|_2^2, \quad \mathbf{w}^* = (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y}$$

Lasso Regression:

$$\mathcal{L} = \frac{1}{n} \sum (y_i - \hat{y}_i)^2 + \lambda \|\mathbf{w}\|_1$$

Gradient Descent Update:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} \mathcal{L}$$

Eigenvalue Equation:

$$\mathbf{A}\mathbf{v} = \lambda \mathbf{v}$$

SVD:

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$$

PCA Projection:

$$\mathbf{x}' = \mathbf{U}_{d'}^\top (\mathbf{x} - \bar{\mathbf{x}})$$

Covariance Matrix:

$$\mathbf{\Sigma} = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^\top$$

Bayes' Theorem:

$$P(\omega | \mathbf{x}) = \frac{P(\mathbf{x} | \omega) P(\omega)}{P(\mathbf{x})}$$

Gaussian PDF (Multivariate):

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2} |\mathbf{\Sigma}|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \mathbf{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$

MLE for Gaussian:

$$\hat{\mu} = \frac{1}{n} \sum x_i, \quad \hat{\sigma}^2 = \frac{1}{n} \sum (x_i - \hat{\mu})^2$$

Sigmoid:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Softmax:

$$P(y = k | \mathbf{x}) = \frac{e^{\mathbf{w}_k^\top \mathbf{x}}}{\sum_j e^{\mathbf{w}_j^\top \mathbf{x}}}$$

SVM (Soft Margin):

$$\min \frac{1}{2} \|\mathbf{w}\|^2 + C \sum \xi_i \quad \text{s.t.} \quad y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i$$

RBF Kernel:

$$K(\mathbf{x}, \mathbf{x}') = \exp\left(-\frac{\|\mathbf{x} - \mathbf{x}'\|^2}{2\sigma^2}\right)$$

Perceptron Update:

$$\mathbf{w} \leftarrow \mathbf{w} + \eta y_i \mathbf{x}_i \quad (\text{if misclassified})$$

Delta Rule:

$$\mathbf{w} \leftarrow \mathbf{w} + \eta (y_i - \hat{y}_i) \mathbf{x}_i$$

13 Practice Exam Questions

Q1: Conceptual – Supervised vs. Unsupervised Learning

Question: Explain the difference between supervised and unsupervised learning. Give two examples of each.

Answer: Supervised learning uses labelled data to learn a mapping from inputs to outputs (e.g., spam detection, house price prediction). Unsupervised learning finds structure in unlabelled data (e.g., K-means clustering, PCA for dimensionality reduction).

Q2: Computation – Normal Equation

Question: Given 3 data points $(1, 2)$, $(2, 3)$, $(3, 5)$, compute the linear regression weights using the normal equation $\mathbf{w}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$.

Hint: Set up $\mathbf{X} = \begin{bmatrix} 1 & 1 \\ 2 & 1 \\ 3 & 1 \end{bmatrix}$, $\mathbf{y} = \begin{bmatrix} 2 \\ 3 \\ 5 \end{bmatrix}$ and compute step by step.

Q3: L1 vs. L2 Regularization

Question: Compare L1 and L2 regularization. When would you prefer one over the other?

Answer: L1 (Lasso) adds $\lambda \sum |w_j|$ and produces **sparse** models (some weights become exactly zero), making it good for feature selection. L2 (Ridge) adds $\lambda \sum w_j^2$ and shrinks all weights but does not set any to zero, better when all features are relevant. Geometrically, L1 has a diamond constraint region and L2 has a circular one.

Q4: PCA Steps

Question: Describe the steps of PCA. What does each eigenvalue tell you?

Answer: (1) Center the data. (2) Compute covariance matrix. (3) Find eigenvalues/eigenvectors. (4) Sort by decreasing eigenvalue. (5) Project onto top d' eigenvectors. Each eigenvalue represents the **amount of variance** captured by the corresponding principal component.

Q5: SVD Application

Question: How is SVD used in recommendation systems?

Answer: The user-item rating matrix is decomposed as $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$. The top- k singular values/vectors capture latent factors (e.g., genre preferences). Missing entries can be predicted from the low-rank approximation, enabling collaborative filtering.

Q6: Bayes' Theorem Application

Question: A disease affects 1% of the population. A test has 95% sensitivity and 5% false positive rate. If a person tests positive, what is the probability they have the disease?

Hint: Use Bayes' theorem with $P(D) = 0.01$, $P(+|D) = 0.95$, $P(+|\neg D) = 0.05$.

$$P(D|+) = \frac{P(+|D) P(D)}{P(+|D) P(D) + P(+|\neg D) P(\neg D)} = \frac{0.95 \times 0.01}{0.95 \times 0.01 + 0.05 \times 0.99} \approx 0.161$$

Q7: SVM Concepts

Question: What are support vectors? What happens when you increase C in a soft-margin SVM?

Answer: Support vectors are the data points closest to (or on the wrong side of) the decision boundary. They “support” (define) the hyperplane. Increasing C penalizes misclassifications more, leading to a **narrower margin** (less tolerance for errors), which can cause overfitting.

Q8: Kernel Trick

Question: Explain the kernel trick. Why is it useful?

Answer: The kernel trick computes dot products in a high-dimensional feature space **without explicitly** computing the mapping $\phi(\mathbf{x})$. This allows SVMs to learn non-linear decision boundaries efficiently. For example, the RBF kernel $K(\mathbf{x}, \mathbf{x}') = \exp(-\|\mathbf{x} - \mathbf{x}'\|^2 / 2\sigma^2)$ maps to an infinite-dimensional space.

Q9: MLE

Question: Derive the maximum likelihood estimate for the mean of a Gaussian distribution.

Answer: The log-likelihood is $\ell(\mu) = -\frac{n}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum (x_i - \mu)^2$. Differentiating w.r.t. μ and setting to zero: $\frac{\partial \ell}{\partial \mu} = \frac{1}{\sigma^2} \sum (x_i - \mu) = 0 \Rightarrow \hat{\mu} = \frac{1}{n} \sum x_i$, which is the sample mean.

Q10: EM Algorithm

Question: Describe the EM algorithm for GMMs. Why might it converge to different solutions?

Answer: E-Step: compute responsibilities γ_{ik} (soft assignments of points to clusters). M-Step: update means, covariances, and mixing coefficients using weighted data. Repeat until convergence. EM converges to a **local maximum**, and different random initializations lead to different local maxima, hence different solutions.

Q11: Neural Networks

Question: Why can't a single-layer perceptron learn XOR? How do we fix this?

Answer: XOR is **not linearly separable** — no single hyperplane can separate the two classes. A single-layer perceptron can only learn linear decision boundaries. **Fix:** Use a **multi-layer perceptron** (MLP) with at least one hidden layer and non-linear activations, or use a kernel method that maps the data to a higher-dimensional space where it becomes linearly separable.

Q12: Bias–Variance Trade-off

Question: Explain the bias–variance trade-off. How does regularization help?

Answer: **Bias** = error from overly simplistic model assumptions (underfitting). **Variance** = error from sensitivity to training data fluctuations (overfitting). Total error $\approx \text{Bias}^2 + \text{Variance} + \text{Irreducible noise}$. Regularization **increases bias slightly** but **reduces variance significantly**, often improving overall generalization.

14 Quick-Reference: Topic-by-Lecture Map

Date	Lecture #	Topics Covered
5 Jan	L01	Course overview, supervised/unsupervised learning, nearest neighbour, linear classifier
8 Jan	L02	Data representation, handcrafted vs. learned features, classification vs. regression
12 Jan	L03	Linear regression, loss functions (MSE), gradient descent, closed-form solution, bias-variance
19 Jan	L04	Optimization, regularization (L1 & L2), Lasso vs. Ridge, geometric interpretation
22 Jan	L05	Matrix rank, eigenvalues/eigenvectors, SVD, low-rank approximation
27 Jan	L06	PCA (intro), covariance matrix, eigenvector projection, data compression
29 Jan	L07	Course review: supervised learning, train/val/test, overfitting, performance metrics
2 Feb	L08	PCA (deep dive), SVD applications, LSI, word embeddings, matrix factorization
5 Feb	L09	Bayesian decision theory, normal distributions, Bayes' theorem, decision boundaries
9 Feb	L10	MLE, Gaussian Mixture Models, EM algorithm, K-means clustering
12 Feb	L11	K-means, Naive Bayes, logistic regression, softmax, loss functions
16 Feb	L12	SVM (hard & soft margin), kernel trick, polynomial/RBF kernels, convex optimization
19 Feb	L13	Practical ML: cross-validation, regularization, overfitting, hyperparameter tuning
23 Feb	L14	Neural networks: perceptron, delta rule, activation functions, MLP, biological vs. artificial

15 Glossary of Key Terms

Term	Definition
Bias	Systematic error from incorrect model assumptions (underfitting).
Classification	Predicting a discrete class label.
Clustering	Grouping similar data points (unsupervised).
Covariance Matrix	Matrix summarizing variance/covariance between all pairs of features.
Cross-Entropy	Loss function for classification tasks.
Cross-Validation	Technique for evaluating models using multiple train/test splits.
Dimensionality Reduction	Reducing the number of features/dimensions.
Eigenvalue	Scalar λ such that $\mathbf{A}\mathbf{v} = \lambda\mathbf{v}$.
Eigenvector	Vector $\mathbf{v}$ unchanged in direction by transformation $\mathbf{A}$.
EM Algorithm	Iterative method for MLE with latent variables (E-step + M-step).
Feature	A measurable property of a data point.
Gradient Descent	Iterative optimization: move in direction of steepest descent.
GMM	Gaussian Mixture Model: data modeled as mixture of Gaussians.
Hyperparameter	Parameter set before training (not learned from data).
Kernel	Function computing inner products in a transformed feature space.
Lasso (L1)	Regularization that promotes sparsity.
Likelihood	Probability of data given model parameters.
MLE	Maximum Likelihood Estimation.
MSE	Mean Squared Error.
Naive Bayes	Classifier assuming feature independence given class.
Normal Equation	Closed-form solution for linear regression.
Overfitting	Model memorizes training data; poor generalization.
PCA	Principal Component Analysis: unsupervised dimensionality reduction.
Perceptron	Simplest neural network model.
Posterior	Updated probability after observing evidence.
Prior	Probability before observing evidence.
Regression	Predicting a continuous value.
Regularization	Penalty on model complexity to reduce overfitting.
ReLU	Rectified Linear Unit: $\max(0, z)$.
Ridge (L2)	Regularization that shrinks all weights.
Sigmoid	Logistic function mapping $\mathbb{R} \rightarrow (0, 1)$.
Singular Value	Non-negative value in $\mathbf{\Sigma}$ of SVD.
Softmax	Generalization of sigmoid for multi-class probabilities.
Support Vector	Data point on or inside the SVM margin boundary.
SVD	Singular Value Decomposition: $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^\top$.
SVM	Support Vector Machine: max-margin classifier.
Underfitting	Model too simple to capture data patterns.
Variance	Error from sensitivity to training data fluctuations (overfitting).